



DSE 3151

OPERATING SYSTEMS

V Semester BTech DSE

Dr. Poornima P Kundapur
Associate Professor

✉ poornima.girish@manipal.edu
☎ 99808 53396



DSE 3153 OPERATING SYSTEMS

Operating System Structure and Operations, Process Management, **Memory Management**, Storage Management, Operating System Services, User Operating System Interfaces, Types of **System Calls**, System Programs, Operating System Structure, System Boot, Overview, Process Scheduling, Operations on Processes, **Inter-process Communication**, Multithreaded Models, Thread Libraries, **Scheduling Algorithms**, Thread Scheduling, Linux scheduling, Critical Section Problem, Peterson's Solution, **Synchronization Hardware**, Semaphores, Logical Versus Physical Address Space, Segmentation, Contiguous **Memory Allocation**, Paging, Structure of Page Table, Segmentation, Demand Paging, Copy-On-Write, **Page Replacement**, Allocation of Frames, Thrashing, Disk Scheduling, Swap-Space Management, System Model, Deadlock: Deadlock prevention, Avoidance, Detection, Recovery, **File Concept**, Protection.






Multicore systems


Mobile computing

Virtualization

Reference Books



<https://www.minix3.org/>



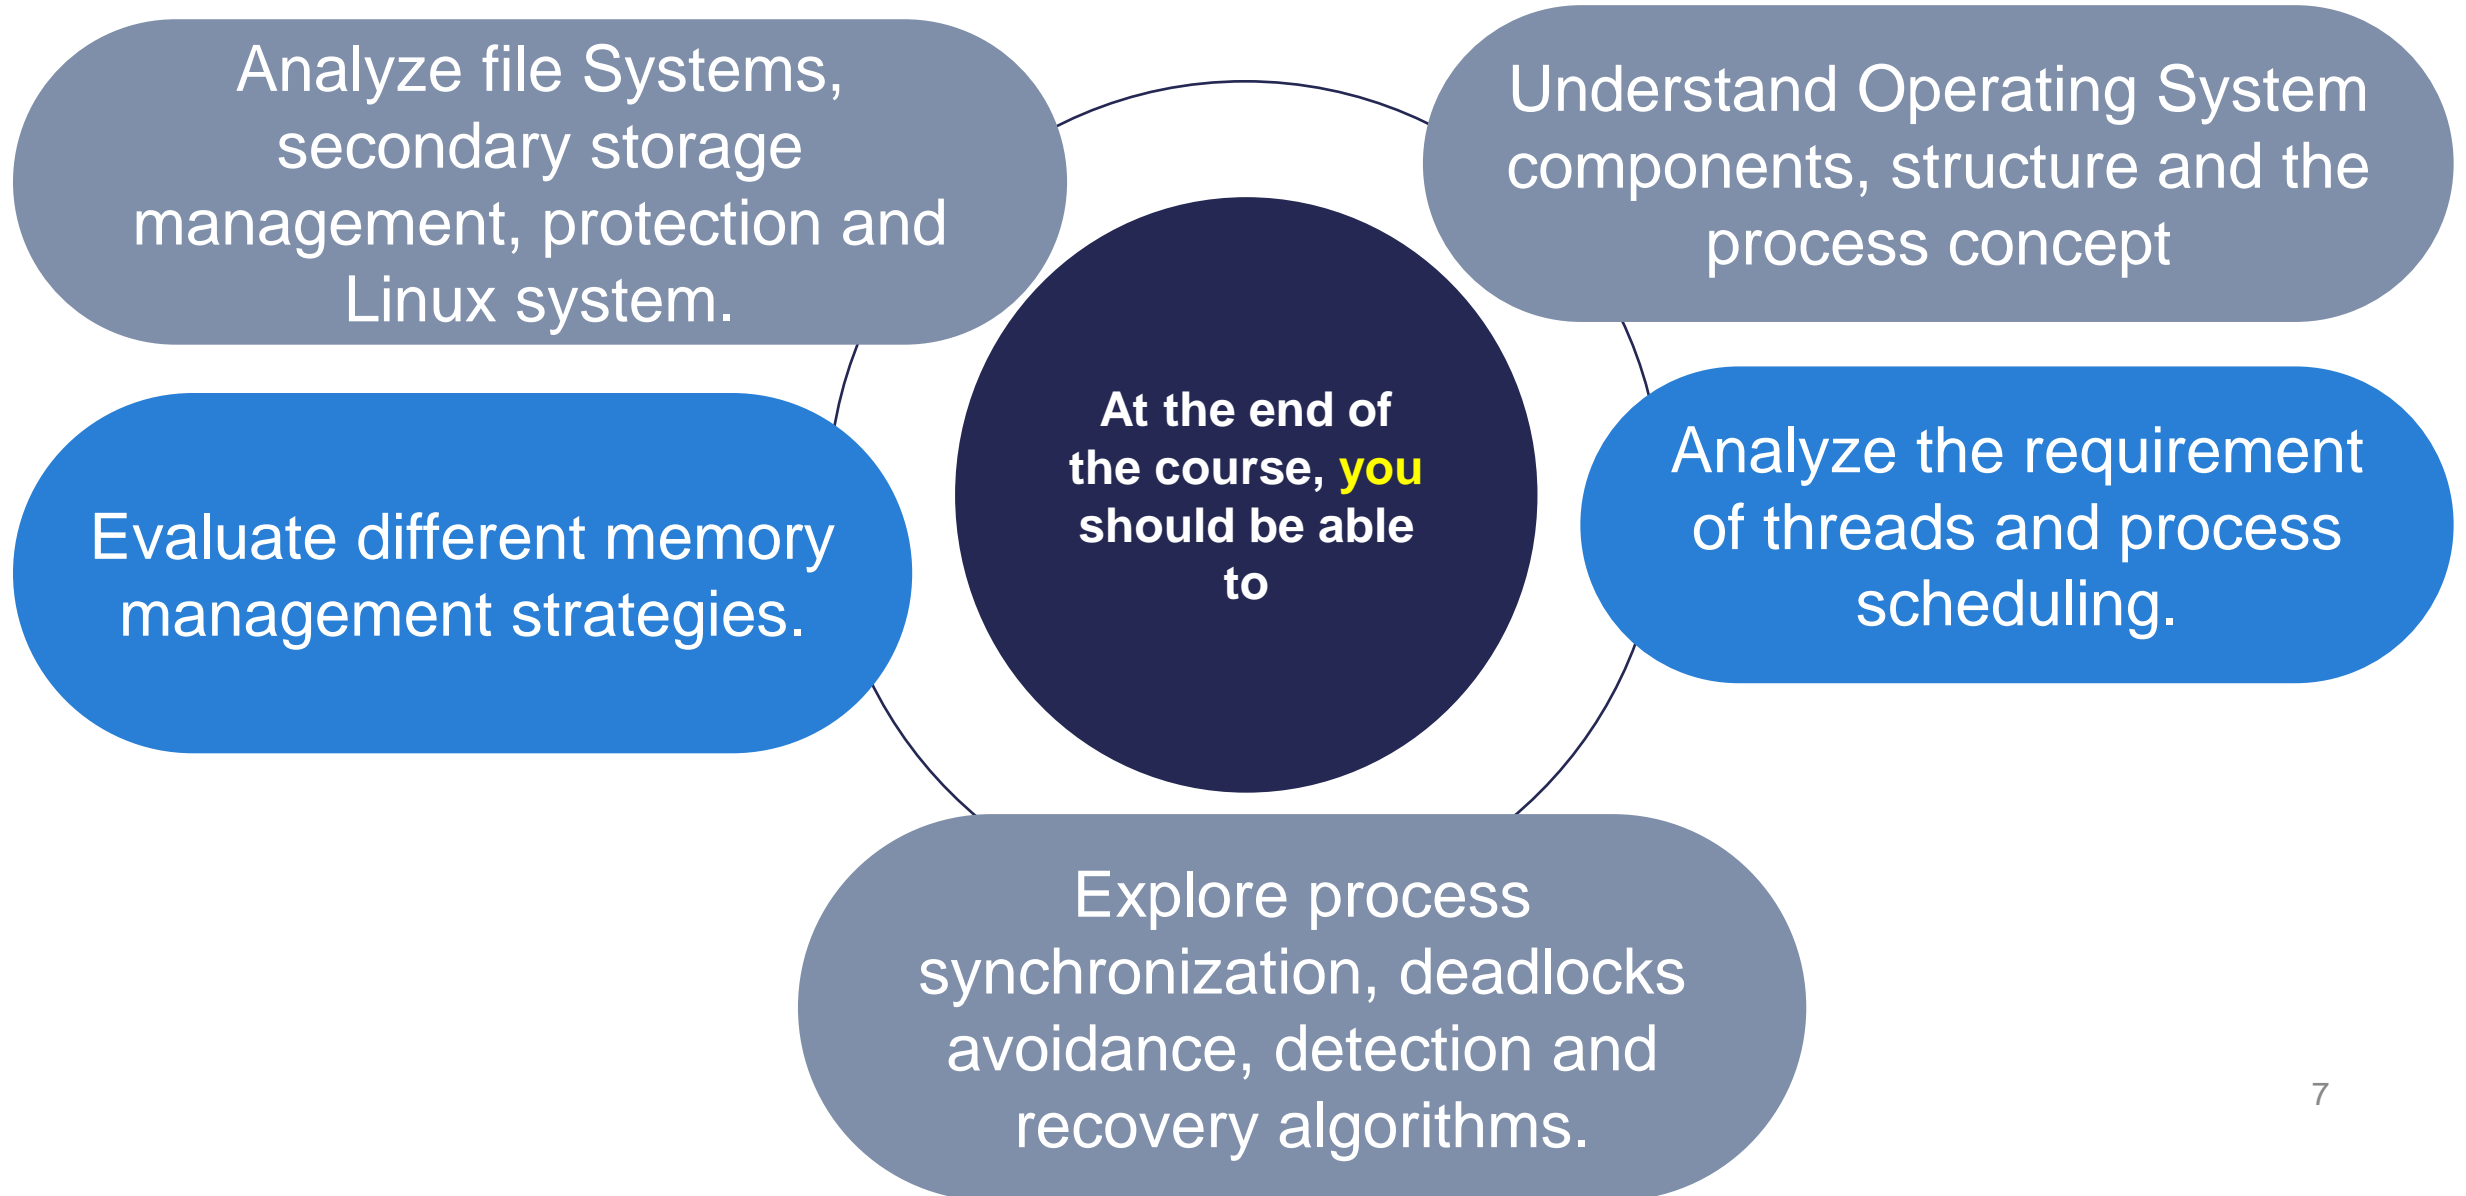
<https://www.youtube.com/watch?v=MG29rUtvNXg>

Class Engagement Rules.

- **Be on time** to class. No Late entry.
- **Late entry:** No attendance but you can sit through the class.
- Carry a notebook for all classes, helps in your revision.
- 75% isn't what you signed up for, so let's try to work above that.
- Questions are allowed, there is no silly question.
- Problem solving is an essential part of learning, engage in it.
- Let the learning happen.



Course Outcomes





Chapter 1: Introduction

Objectives of this chapter.


- To describe the basic organization of computer systems
- To provide a grand tour of the major components of operating systems
- To give an overview of the many types of computing environments
- To explore several open-source operating systems



Computer System Structure



Hardware



Operating system




Application programs



Users



Four Components of a Computer System



What Operating Systems Do?

- Depends on the point of view:
 - **Users:** Convenience, ease of use and good performance and don't care about resource utilization
 - **Shared computer** (mainframe or minicomputer) must keep all users happy
 - **Users of dedicate systems** (workstations) have dedicated resources but frequently use shared resources from servers
 - **Handheld computers:** Resource poor, optimized for usability and battery life
 - **Embedded computers:** (devices and automobiles) have little or no user interface.



Definition of an Operating System.

- An Operating System (OS) is a software that acts as an **interface between** computer hardware components and the user.
- Operating system **goals:**
 - Execute user programs and make solving user problems easier
 - Make the computer system convenient to use
 - Use the computer hardware in an efficient manner



Definition of an Operating System.

- OS is a **resource allocator**
 - Manages all resources
 - Decides between conflicting requests for efficient and fair resource use
- OS is a **control program**
 - Controls execution of programs to prevent errors and improper use of the computer



Definition of an Operating System.

- No universally accepted definition
- “**Everything a vendor ships when you order an operating system**” is a **good approximation**
- But varies wildly
 - “The one program running at all times on the computer” is **the kernel**.
 - Everything else is either
 - a system program (ships with the operating system) , or
 - an application program.



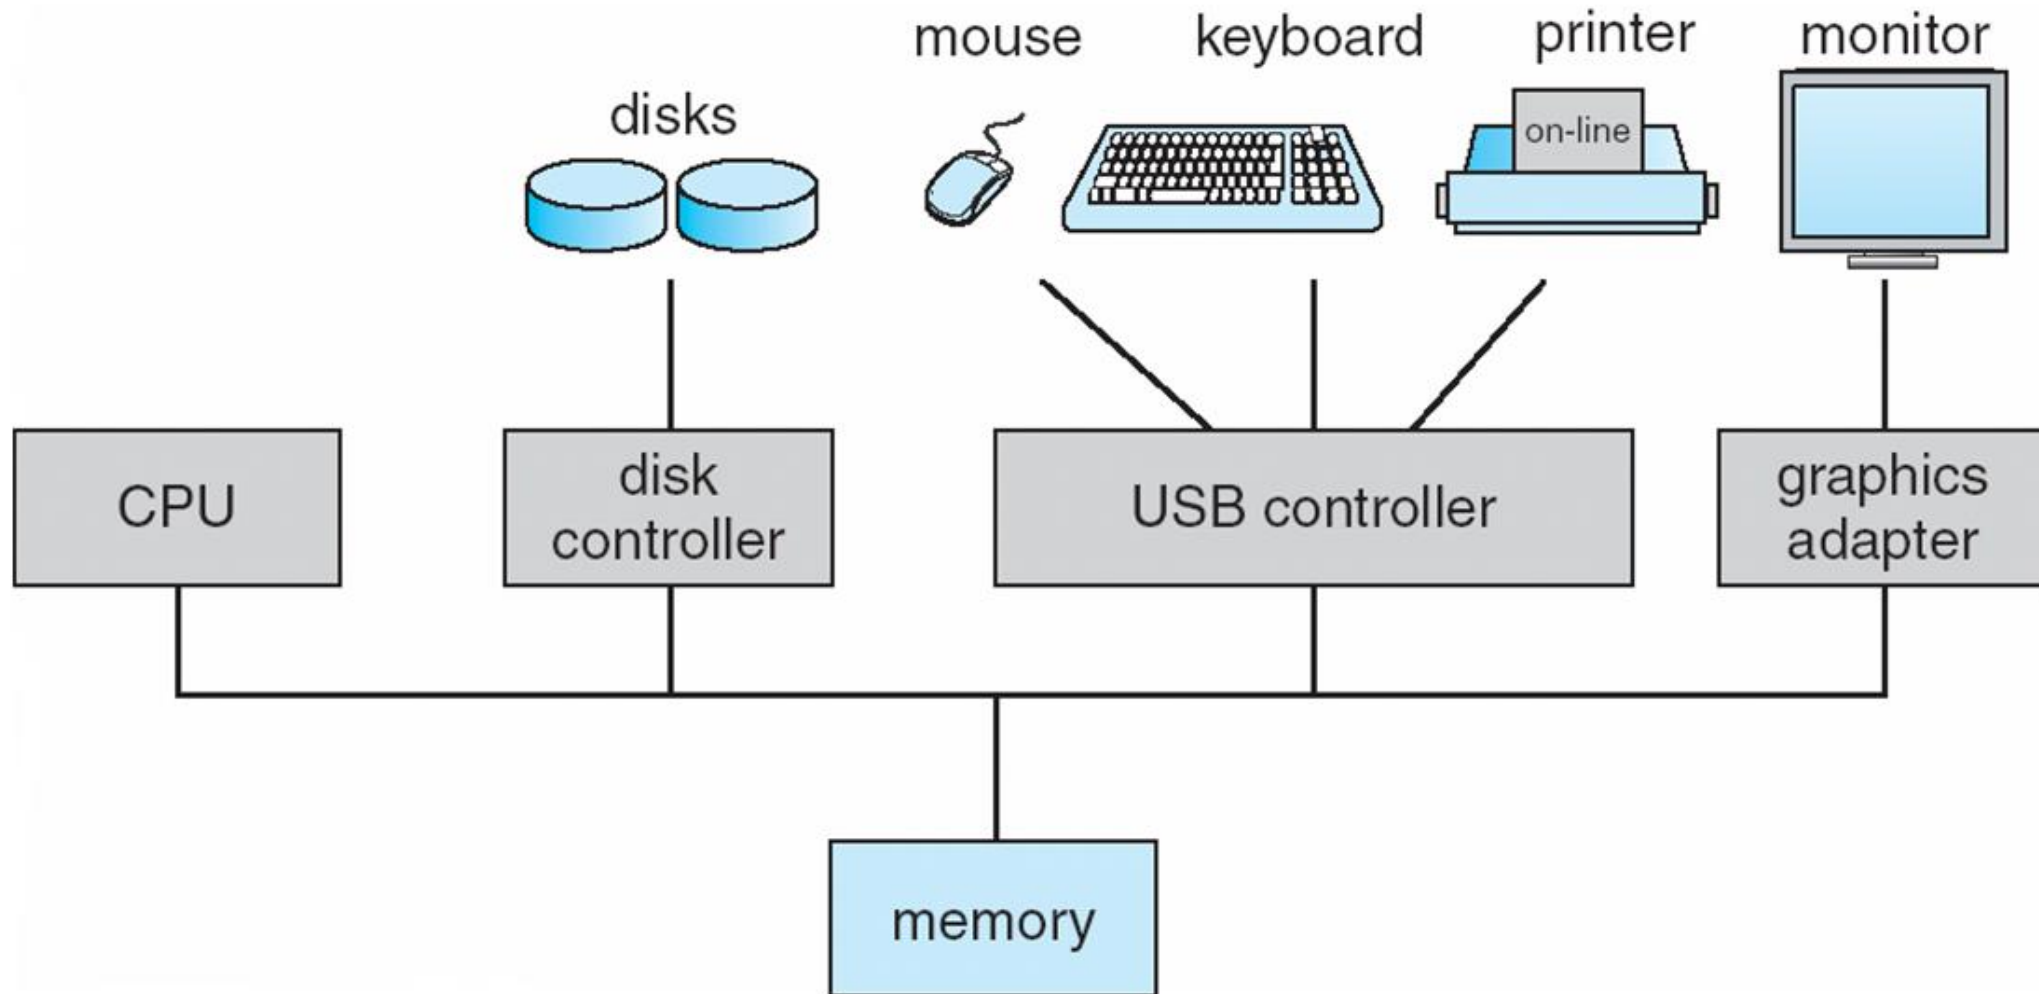
Definition of a “reliable” Operating System.

- An operating system is said to be reliable when a typical user has never experienced even a single failure in his or her lifetime and does not know anybody who has ever experienced a failure.

Andrew S. Tanenbaum: MINIX 3



Computer System organization.



Storage definitions

- Bit: 0 or 1 (basic unit)
- Nibble: 4 bit aggregation
- Byte: 8 bit octet
- Kilobyte: 1024 bytes (KB)
- Megabyte: 1024^2 bytes (MB)
- Gigabyte: 1024^3 bytes (GB)
- Terabyte: 1024^4 bytes (GB)
- Petabyte: 1024^5 bytes (GB)
- **Network measurements all in bits**



After the break.



Operating System Structure

Multiprogramming

Timesharing



Multiprogramming (Batch system)

Job pool



- Needed for efficiency
- Single program ***cannot keep CPU or I/O devices*** busy at all times
- Single users have multiple programs running
- Increases CPU utilization
 - Organizes **jobs (code & data)** so CPU always has one to execute




Multiprogramming (Jobs)

- How does it work?
- OS keeps several jobs in **memory** simultaneously
- *Main memory (RAM)* is **too small** to accommodate all jobs, the jobs are kept initially on the disk in the **job pool**.
- The set of jobs in memory can be a **subset of the jobs** kept in the job pool.
- One job selected and run via job scheduling
- When it has to wait (for I/O for example), OS switches to another job



Multiprogramming (Jobs)

Job pool
consists of **all**
processes residing on
disk awaiting allocation
of main memory.



Timesharing (multitasking).

*Logical
extension*

- CPU executes **multiple jobs** by ***switching*** among them, but the switches occur so **frequently** that the **users can interact with each *program*** while it is running
- *Creating* **interactive** computing
- Response time should be **< 1 second**



Timesharing (multitasking).

- **Allows many users to share the computer simultaneously**
- **Each action** or command tends to be **short**, only a **little CPU** time is *needed for each user*.
- As the **system switches rapidly** from **one user** to the **next**, each **user is given the impression** that the entire computer system is **dedicated** to his use, even though it is being shared among many users.



Timesharing (Process).

- Each user has **at least one separate program** in memory.
- A program loaded into memory and executing is called a **process**

CPU scheduling

Multiprogramming



Timesharing (Job scheduling)

If **several jobs are ready** to be brought into memory, and if there is **not enough room** for all of them, then the **system must choose** among them.

Making this decision involves **job scheduling**
(Chapter 6)



Timesharing (Memory Management)

When the operating system **selects a job** from the job pool,

Loads that job into **memory** for **execution**

Several programs in memory at the *same time* requires some form of **memory management** (Chapters 8 and 9)



Timesharing (CPU Scheduling)

If **several jobs** are **ready to run** at the **same time**, system *must choose which job* will run first.

Making this decision is called **CPU scheduling** (*Chapter 6*)



Time sharing (Virtual Memory)

Reasonable
response time

- Processes are swapped in and out of main memory to the disk.
- **Virtual memory:** *a technique that allows the execution of a process that is not completely in memory (Chapter 9).*
- Why virtual-memory scheme?
 - Enables users to run programs that are larger than actual physical memory.
 - Abstracts main memory into a large, uniform array of storage, separating logical memory as viewed by the user from physical memory.
 - Frees programmers from concern over memory-storage limitations.



Time sharing (File System Management)

- **File system** resides on a collection of disks: Disk management (Chapter 10).
- To ensure orderly execution, the system must provide mechanisms for **job synchronization and communication** (Chapter 5)
- Ensure that jobs do not get stuck in a **deadlock**, forever waiting for one another (Chapter 7).



Modern Operating-System Operations.

- **Interrupt driven** (hardware and software)
 - Hardware interrupt by one of the devices
 - Software interrupt (**exception or trap**):
 - Software error (e.g., division by zero)
 - Request for operating system service
 - Other process problems include infinite loop, processes modifying each other or the operating system



Modern Operating-System Operations.

More **subtle errors can occur** in a **multiprogramming system**, where one erroneous program might modify another program, the data of another program, or even the operating system itself

A **properly designed operating system** must ensure that an **incorrect (or malicious) program** cannot cause other programs to execute incorrectly.



Dual-Mode and Multimode Operation.

In order to ensure the proper execution of the operating system, we must be able to **distinguish** between the execution of **operating-system code** and **user-defined code**



Modes of Operation.

User mode

Kernel mode

Supervisor mode, system mode, or privileged mode





Modes of Operation.

A bit, called the **mode bit**, indicates the current mode: **kernel (0)** or **user (1)**.

Mode bit distinguishes between a task that is executed on behalf of the operating system and one that is executed on behalf of the user.

When the **computer system** is **executing on behalf of a user** application, the system is in **user mode**.

When a **user application** requests a **service** from the operating system (via a **system call**), the system must transition from user to **kernel mode** to fulfil the request.



Dual-Mode and Multimode Operation.

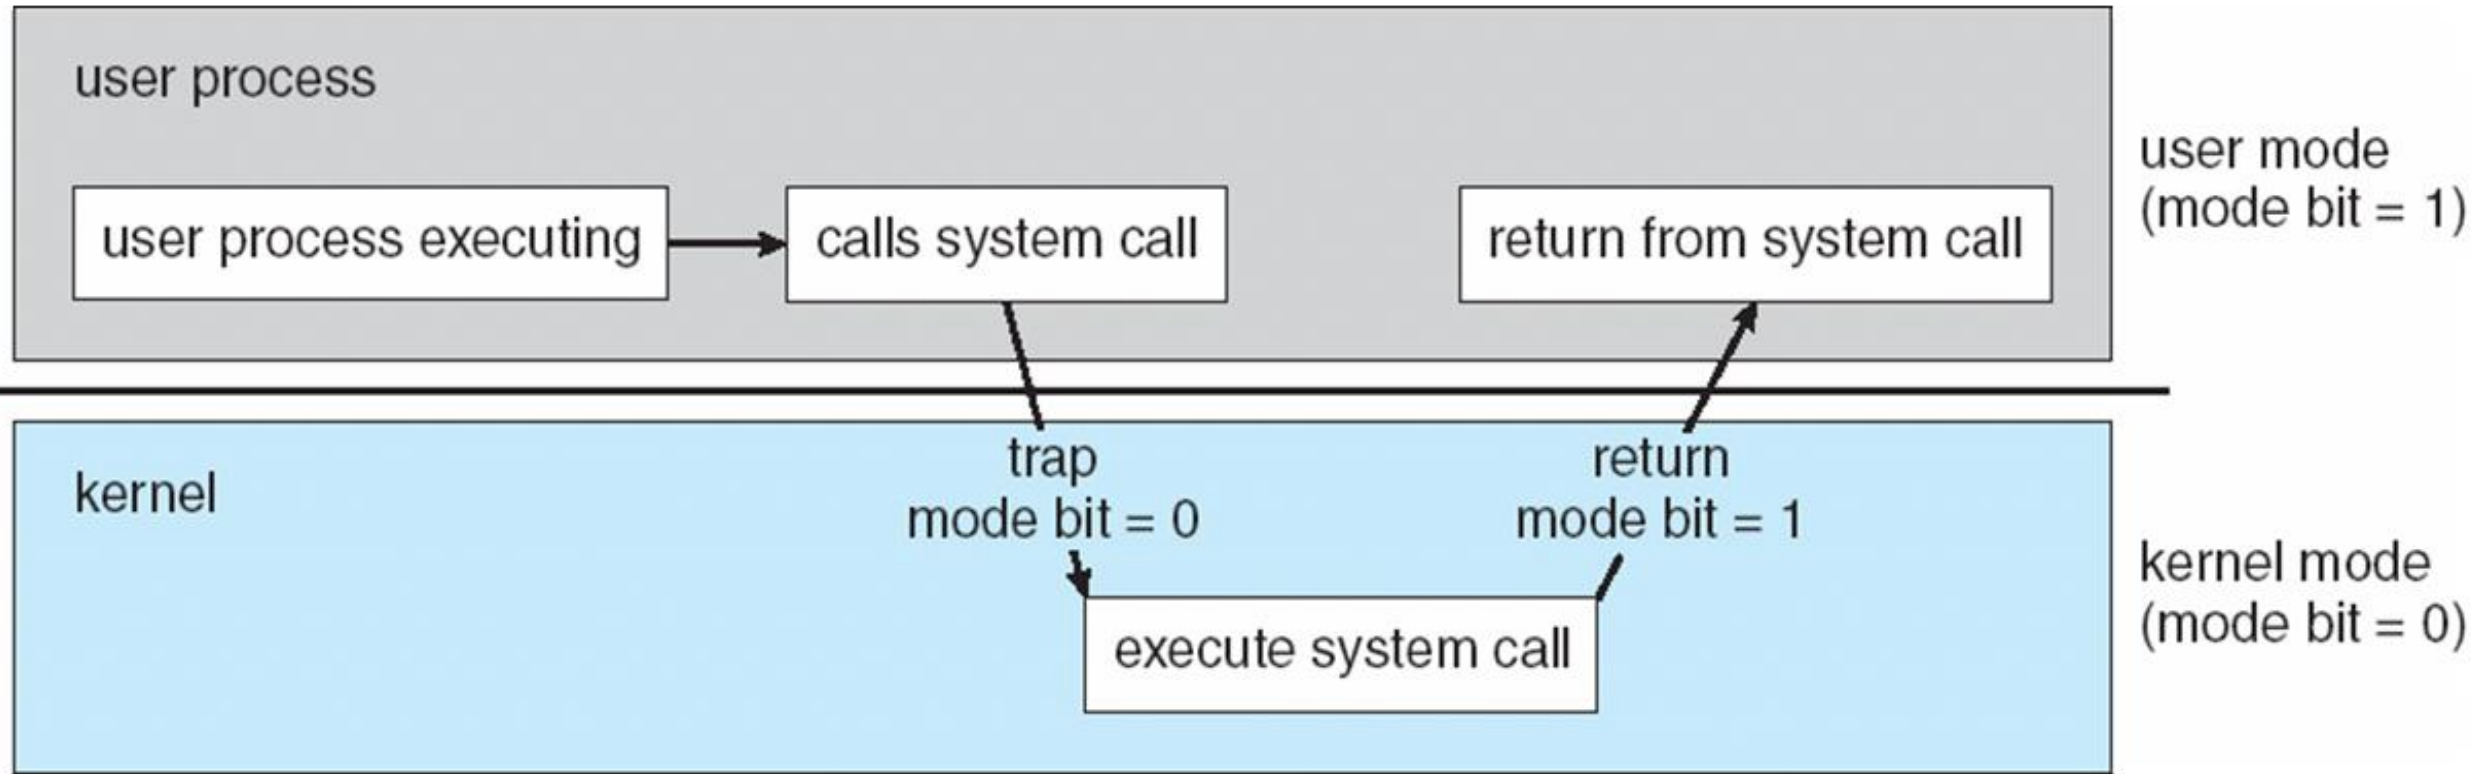


Figure 1.10 Transition from user to kernel mode



Dual-Mode and Multimode Operation

- System boot time: **Hardware** starts in **kernel mode**.
- OS then loaded and starts **user applications** in **user mode**.
- Trap or interrupt occurs:
 - Hardware switches from user mode to kernel mode (state of mode bit – 0)
- Whenever the OS gains control of the computer, it is in **kernel mode**.
- System always switches to user mode (setting mode bit to 1) before passing control to a user program.



Dual-Mode Operation? How does it help?

- **Means for protecting the OS** from errant users—and errant users from one another.
- How?
 - By designating some of the machine instructions that may cause harm as **privileged instructions**.
- Hardware allows **privileged instructions** to be executed only **in kernel mode**.
- If an attempt is made to execute a privileged instruction in user mode, the hardware does not execute (*illegal & traps it*)



So, what is a privileged instruction?

The instruction to switch to kernel mode is an example of a **privileged instruction**.

I/O control

Interrupt management

Timer management

Virtual Memory Manager



AFTER LAST CLASS.





Timer.

- **Does not allow** a user program to get stuck in an infinite loop or to fail to call system services and never return control to the operating system.
- Can be **set to interrupt** the computer after a specified period.

Fixed

(1/60 second)

Variable

(1 millisecond to 1 second)
Implemented by a fixed-rate clock and a counter





Who/What sets the counter?

Timer.

- **OS sets the counter.** Every time the clock ticks, the counter is decremented. When the counter reaches 0, an interrupt occurs.
- For instance, a ***10-bit counter*** with a 1-millisecond clock allows interrupts at intervals from 1 millisecond to 1,024 milliseconds, in steps of 1 millisecond.



Process Management (1)

- A program does nothing *unless its instructions* are executed by a CPU.
- **A program in execution**, as mentioned earlier many times, **is a process**.
 - *Compiler is a process.*
 - *A word-processing program* on a PC is a process.
 - *A system task*, such as sending output to a printer, can also be a process (or at least part of one).



Process Management (2)

- A **process** needs certain resources (CPU time, memory, files, and I/O devices) to accomplish its task.
 - Given to the process when it is **created or allocated** to it while it is running.
- A program by itself is not a process.
- A program is a passive entity, like the contents of a file stored on disk.
- A **process is an active entity**. A single-threaded process has one program counter specifying the next instruction to execute (Chapter 4)



Process Management (3)

- A **process** is the unit of work in a system.
- A **system consists of a collection of processes:**

*Operating system
processes*

(those that execute system
code)

User processes

(those that execute
user code)

- All **these processes execute concurrently** by multiplexing on a single CPU, for example.



Process Management (4)

- OS is responsible for the **following activities** in connection with process management:
 - **Scheduling processes and threads** on the CPUs
 - **Creating and deleting** both *user and system processes*
 - **Suspending and resuming** processes
 - Providing mechanisms for **process synchronization**
 - Providing mechanisms for **process communication**



MEMORY MANAGEMENT (1)

- **Main memory** is central to the operation of a modern computer system.
- To **execute a program** all (or part) of the instructions must be in ***memory***
- All (or part) of the **data** that is needed by the ***program*** must be in memory.
- Memory management determines what is in memory and when.
- Optimizing CPU utilization and computer response to users



MEMORY MANAGEMENT (1)

- OS is responsible for the **following activities** in connection with **memory management**:
 - **Keeping track** of which parts of **memory** are **currently being used** and **who** is using them
 - **Deciding which processes** (or parts of processes) and data to **move into and out of memory**
 - **Allocating** and **deallocating memory space** as needed



STORAGE MANAGEMENT.

- OS provides **uniform, logical view** of information storage
 - Abstracts physical properties to logical storage unit: **file**
 - Each **medium** is **controlled by device** (i.e., disk drive, tape drive)
 - Varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)



FILE SYSTEM MANAGEMENT.

- **Files** usually organized into **directories/folder**
- **Access control** on most systems to determine who can access what
- **OS activities** include
 - **Creating and deleting** files and directories
 - **Primitives to manipulate** files and directories
 - **Mapping files** onto secondary storage
 - **Backup files** onto stable (non-volatile) storage media



MASS STORAGE MANAGEMENT(1)

- **Disks** used to **store data** does not fit in main memory or data that must be kept for a “long” periods of time.
- Proper management is of central importance
- Entire **speed of computer operation** hinges on disk subsystem and its algorithms
- **OS activities**
 - Free-space management
 - Storage allocation
 - Disk scheduling



MASS STORAGE MANAGEMENT(2)

- Some **storage** need **not** be **fast**
 - **Tertiary storage** (*not crucial to system performance*) includes:
 - optical storage, magnetic tape
 - Still must be managed – by OS or applications
 - Varies between:
 - **WORM**: write-once, read-many-times, and
 - **RW**: read-write



Performance of Various Levels of Storage.

Level	1	2	3	4	5
Name	registers	cache	main memory	solid state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 - 25	80 - 250	25,000 - 50,000	5,000,000
Bandwidth (MB/sec)	20,000 - 100,000	5,000 - 10,000	1,000 - 5,000	500	20 - 150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape




Caching.

- Information is normally kept in some storage system (*main memory*).
- As it is used, it is copied into a faster storage system (*the cache*) on a temporary basis.
- When a particular piece of **information** is needed, **cache** is checked. **If found**, *information is directly used from the cache.*
- **If not**, *information from the source is used and a copy is placed in the cache* under assumption it will be needed again.



Migration of data “A” from Disk to Register.



- **Multitasking environments** **uses most recent value**, no matter where it is stored in the storage hierarchy
- Multiprocessor environment must provide **cache coherency** in hardware such that all CPUs have the most recent value in their cache
- Distributed environment situation even more complex, several copies of a datum can exist.



Migration of data “A” from Disk to Register.




Figure 1.6 Symmetric multiprocessing architecture.

- A **copy of A** may exist simultaneously in several caches.
- Since the various ***CPUs can all execute in parallel***, an **update to value of A** in one cache is ***immediately reflected in all other caches*** where A resides.
- This situation is called **cache coherency**, and it is usually a hardware issue.



I/O Subsystems.

- One purpose of OS is to **hide peculiarities** of hardware devices from the user
- **I/O subsystem** responsible for:
 - **Memory management of I/O** including **buffering** (storing data temporarily while it is being transferred), **caching**, **spooling** (overlapping of output of one job with input of other jobs)
 - General device-driver interface
 - Drivers for specific hardware devices



PROTECTION AND SECURITY

Protection

Any mechanism for controlling access of processes or users to resources defined by the OS

- Memory-addressing hardware ensures that a process can execute only within its own address space.
- The timer ensures that no process can gain control of the CPU without eventually relinquishing control.

SECURITY

Defense of the system against internal and external attacks

Huge range, including denial-of-service, worms, viruses, identity theft, theft of service

PROTECTION AND SECURITY

- Systems generally first distinguish among users, to determine who can do what
 - User identities (**user IDs**, security IDs) include name and associated number, one per user
 - **User ID** then associated with **all files, processes of that user** to determine access control
 - Group identifier (**group ID**) allows set of users to be defined and controls managed, then also associated with each process, file
 - **Privilege escalation** allows user to change to effective ID with more rights



End of Chapter 1

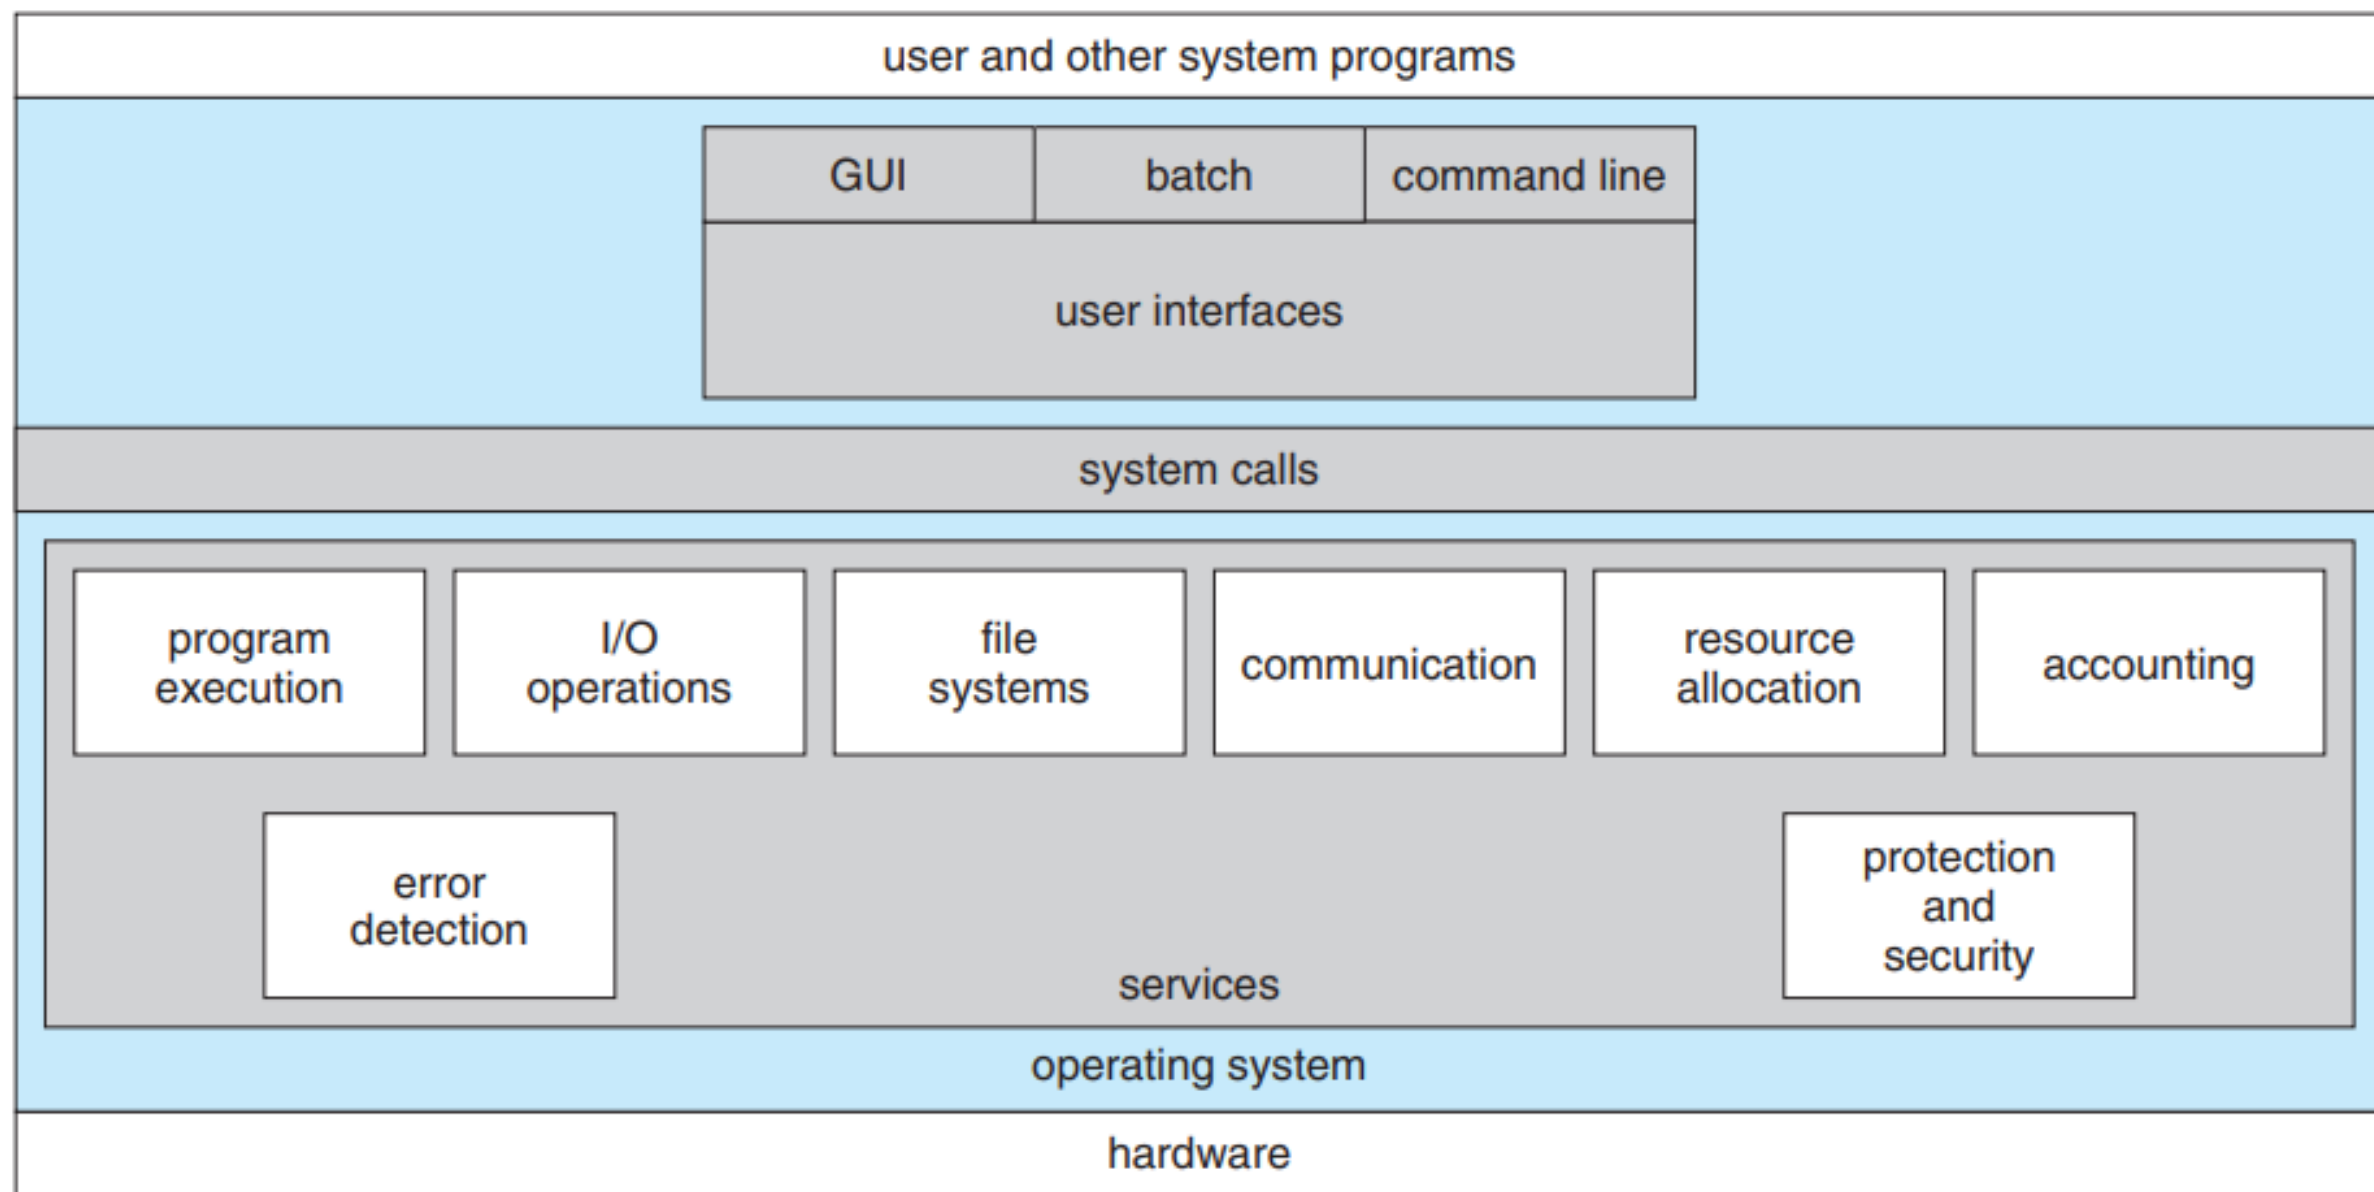


Figure 2.1 A view of operating system services.