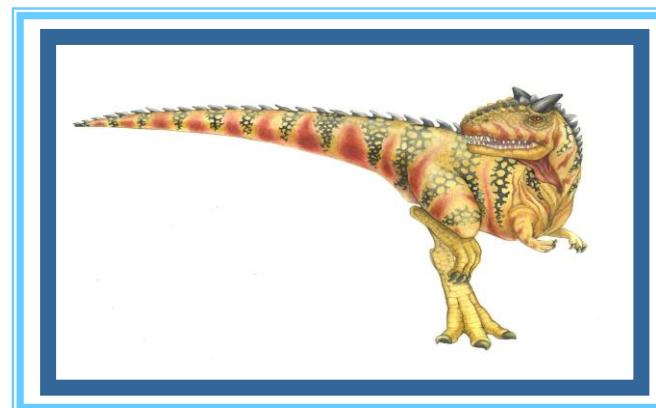


Chapter 5: Process Synchronization





Chapter 5: Process Synchronization

- Background
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Mutex Locks
- Semaphores
- Classic Problems of Synchronization
- Monitors
- Synchronization Examples
- Alternative Approaches





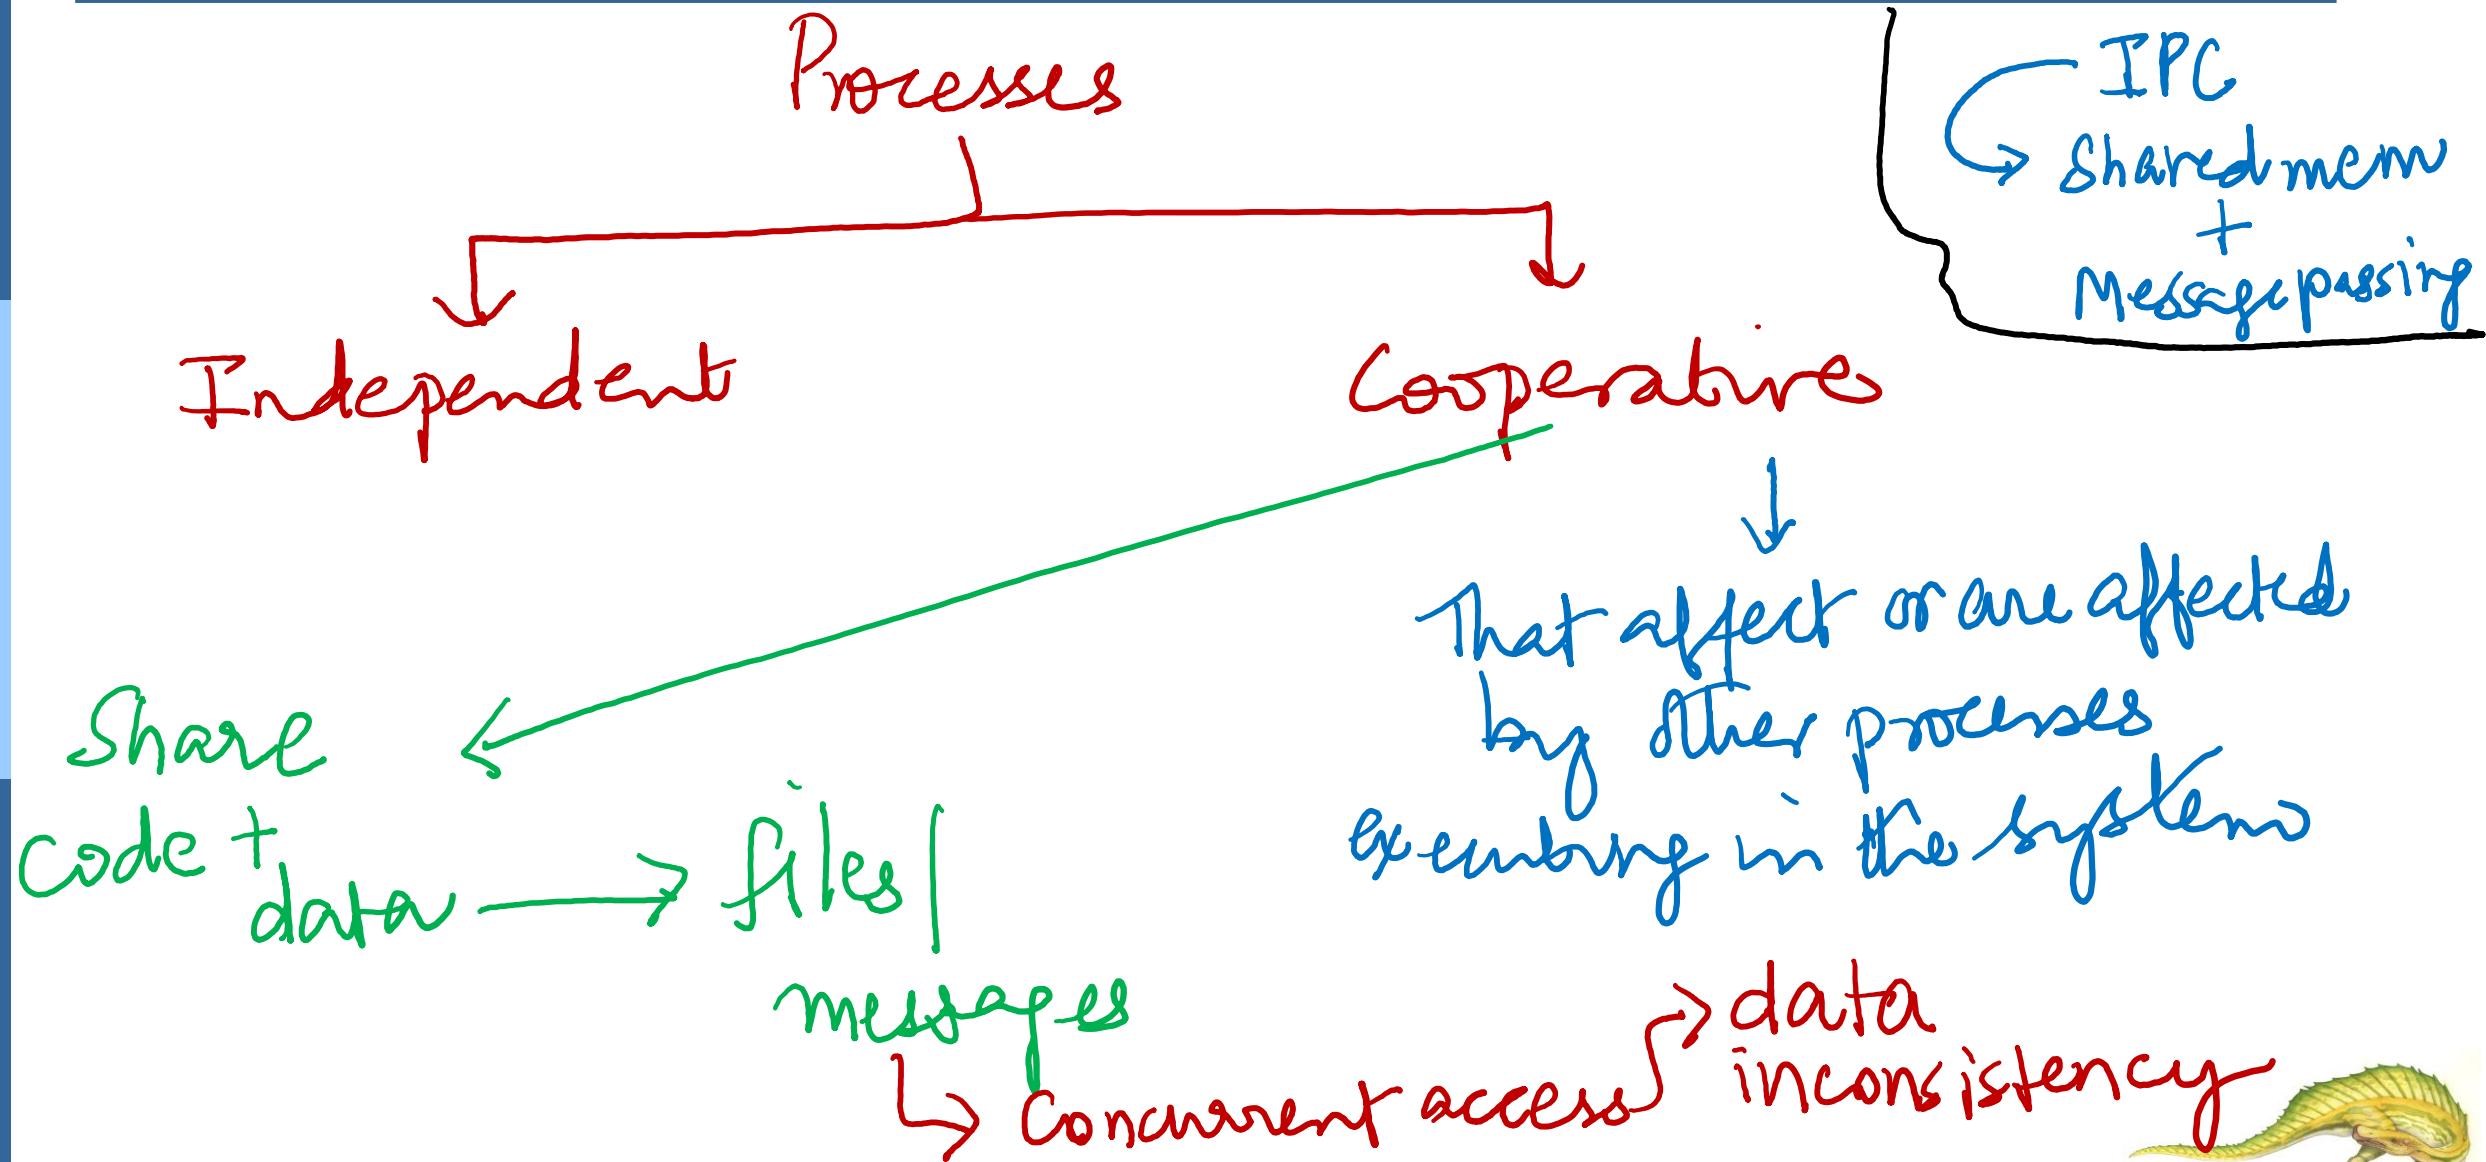
Objectives

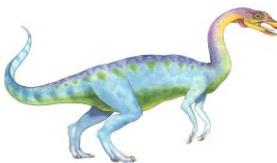
- To present the concept of process synchronization.
- To introduce the **critical-section problem**, whose solutions can be used to ensure the consistency of shared data
- To present both **software and hardware solutions** of the critical-section problem
- To examine several classical process-synchronization problems
- To explore several tools that are used to solve process synchronization problems





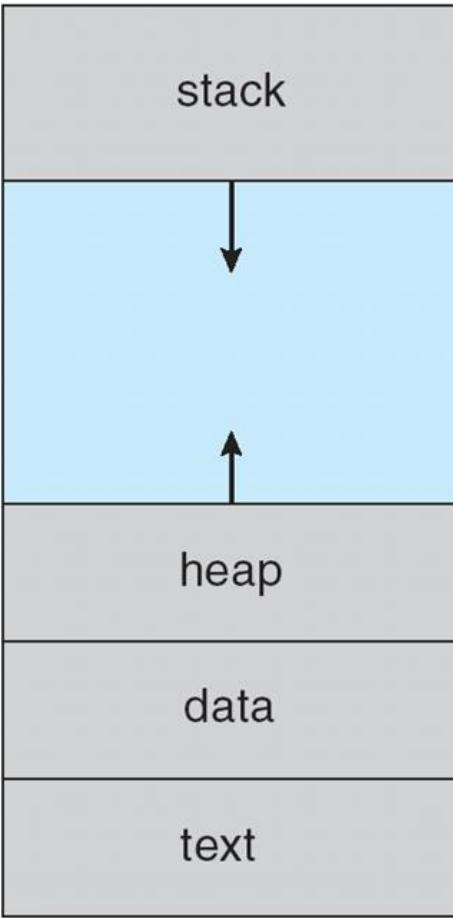
Recap & Looking ahead



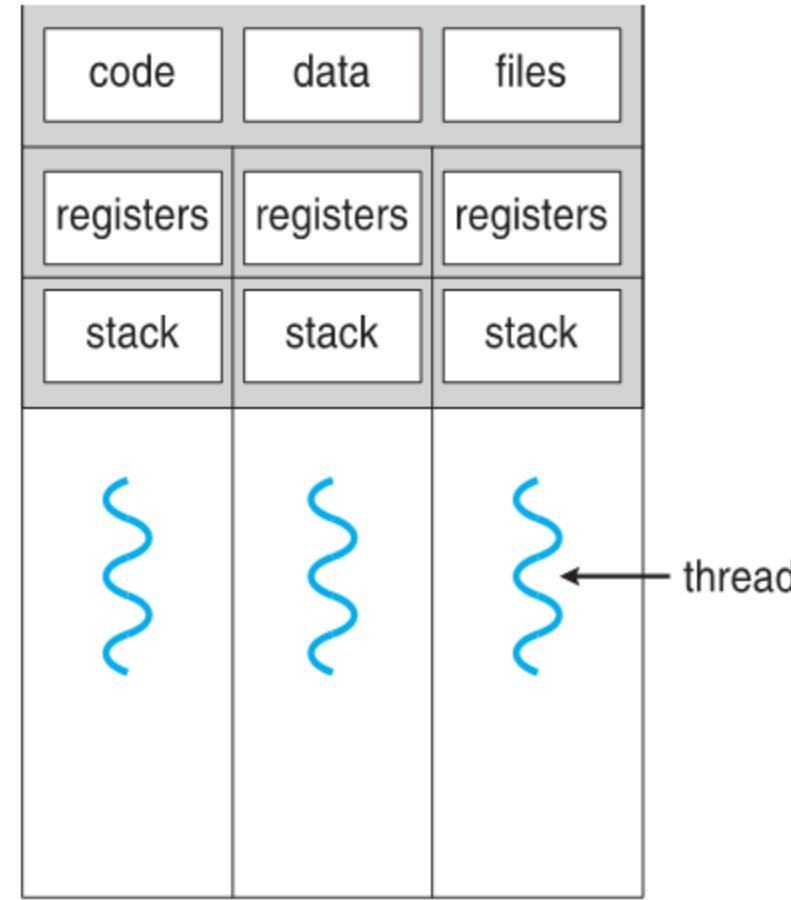


Recap & Looking ahead

max

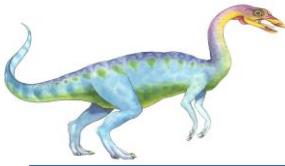


0



multithreaded process

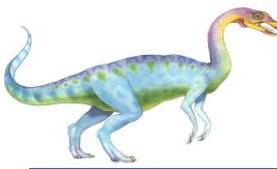




Background

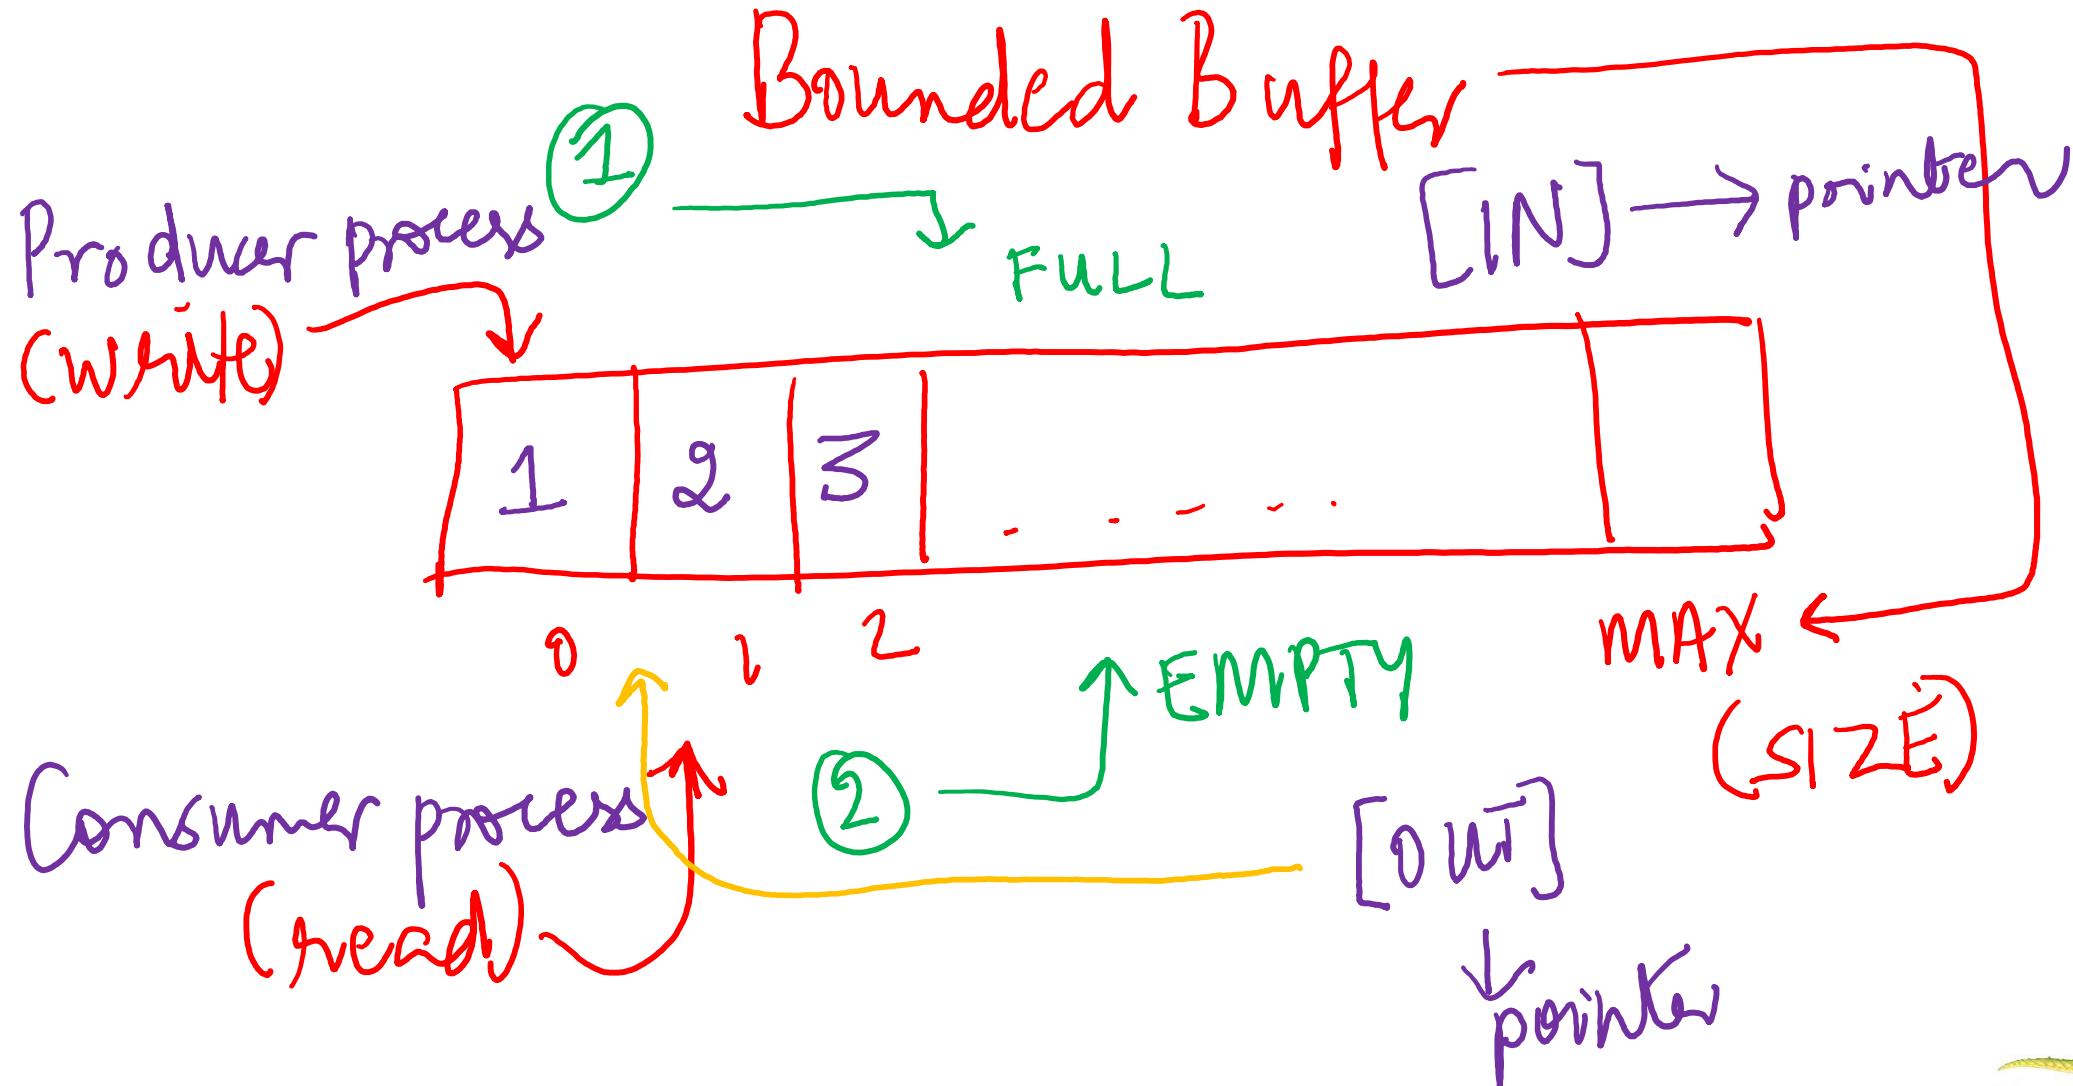
- Processes can **execute concurrently**
 - May *be interrupted at any time*, partially completing execution
- **Concurrent access to shared data may result in data inconsistency**
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Illustration of the problem:
- Suppose that we wanted to provide a solution to the consumer-producer problem that fills all the buffers. We can do so by having an integer counter that keeps track of the number of full buffers. Initially, counter is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.

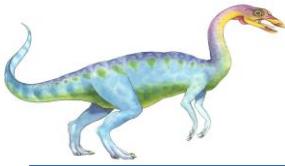




Recap on Producer Consumer Problem

COUNTER

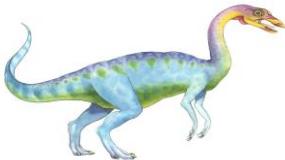




Producer

```
while (true) {  
    /* produce an item in next produced */  
  
    1 while (counter == BUFFER_SIZE);  
        /* do nothing */  
  
    2 buffer[in] = next_produced;  
    3 in = (in + 1) % BUFFER_SIZE;  
    4 counter++;  
}
```





Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```

Handwritten annotations in red:

- A vertical line with a bracket underlines the entire code block from the first brace to the final closing brace.
- The word "empty" is written in red above the line "next_consumed = buffer[out];".
- A red bracket groups the lines "out = (out + 1) % BUFFER_SIZE;" and "counter--;".





Race Condition

- `counter++` could be implemented as

```
register1 = counter  
register1 = register1 + 1  
counter = register1
```

5 | 5

- `counter--` could be implemented as

```
register2 = counter  
register2 = register2 - 1  
counter = register2
```

5 | 4

- Consider this execution interleaving with “count = 5” initially:

S0: producer execute `register1 = counter`
S1: producer execute `register1 = register1 + 1`
S2: consumer execute `register2 = counter`
S3: consumer execute `register2 = register2 - 1`
S4: producer execute `counter = register1`
S5: consumer execute `counter = register2`

{register1 = 5}
{register1 = 6}
{register2 = 5}
{register2 = 4}
{counter = 6 }
{counter = 4}





Critical Section Problem

Twin interleaved flag

- Consider system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has **critical section** segment of code
 - Process may be changing common variables, updating table, writing file, etc ||
 - When one process in critical section, no other may be in its critical section
- **Critical section problem** is to design protocol to solve this
- Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**





Critical Section

- General structure of process P_i

```
do {
```

```
    entry section
```

```
    critical section
```

```
    exit section
```

```
    remainder section
```

```
} while (true);
```

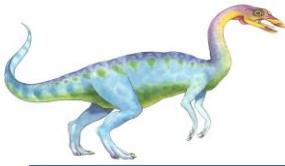




Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
 - Assume that each process executes at a nonzero speed
 - No assumption concerning **relative speed** of the n processes





Algorithm for Process P_i

```
do {  
    while (turn == j);  
        critical section  
        turn = j;  
    remainder section  
} while (true);
```





Algorithm for Process P_i

```
do {  
    turn != i  
    while (turn == j); //N  
        critical section  
        turn = j; //before return,  
        //set turn=j  
    remainder section  
} while (true);
```

(P_i)

turn is other process

j

(P_j)

turn = j

P_i → enters CS

while (turn == i);

CS

turn = i;

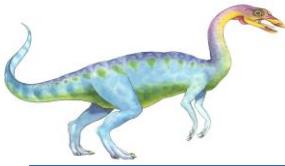
CS

//before return,
//set turn=j

remainder section

P_i entry is regulated by P_j & vice versa





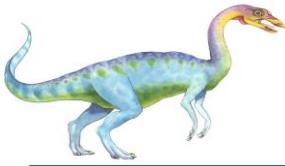
RECAP QUESTION 1

1. What do you understand by Race condition?

A situation where several processes access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which the access takes place, is called a race condition.

2. Busy wait means?





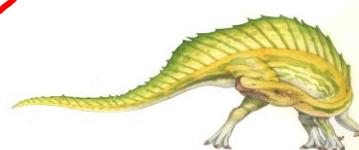
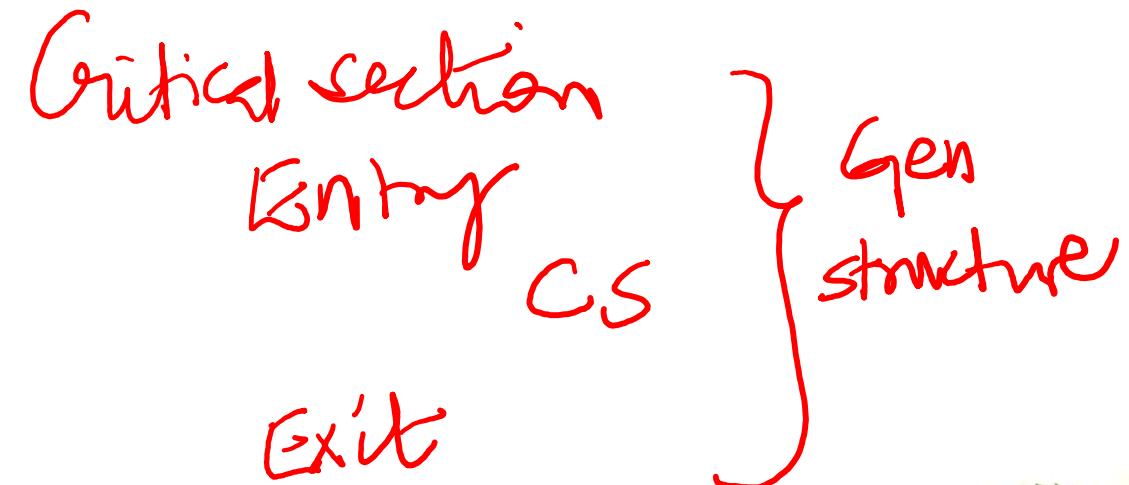
RECAP QUESTION 2 & 3

3. What is the solution to race condition?

Process synchronization

4. Critical section problem must satisfy which three conditions?

- a. Mutual exclusion
- b. Progress
- c. Bounded waiting





flag[i]=T

$P_i \rightarrow CS$

Entry:
Section

① $\text{flag}[i] = T;$ // $\text{flag}[j] = F$
 ② $\text{while } (\text{flag}[j] == T);$
 ↳ Critical section

Exit
section

$\boxed{\text{flag}[i] = F}$

render section

flag variable → 2 process → Peterson
bargain

P_j
 $\boxed{P_j}$ Progress X

① $\text{flag}[j] = T;$ ✓
 ② $\text{while } (\underline{\text{flag}[i] == T});$

Critical section

$\text{flag}[j] = F$

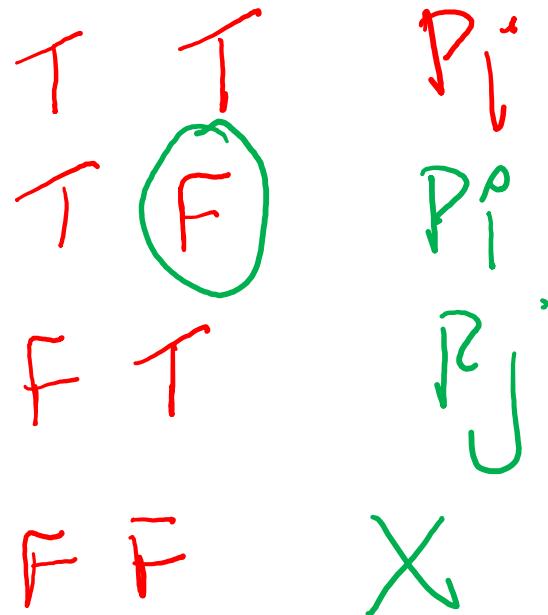
render section





Flag ←

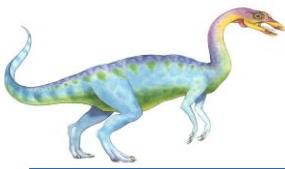
P_i P_j who enters CS?



Flag_{i,j} ≠ T
Flag_{j,i} ≠ T

P_i: ① { P_j: ① ② } ↓
P_i ↓
{ while }
↳ Deadlock!





Critical-Section Handling in OS *HOW?*

→ 2 processes → open files simultaneously?

Two approaches depending on if kernel is preemptive or non-preemptive

- **Preemptive** – allows preemption of process when running in kernel mode
- **Non-preemptive** – runs until exits kernel mode, blocks, or voluntarily yields CPU
- ▶ Essentially free of race conditions in kernel mode

Execution in Kernel mode → OS designer!





s/w

Peterson's Solution

(User process)

- Good algorithmic description of solving the problem
- **Two process** solution
- Assume that the **load** and **store** machine-language instructions are **atomic**; that is, **cannot be interrupted**
 - whose turn ?
- The two processes share two variables:
 - `int turn;` → takes only 2 values (i or j) (ME)
 - Boolean `flag[2]` → Before entering CS → process has to raise flag
- The variable `turn` indicates whose turn it is to enter the critical section
- The `flag` array is used to indicate if a process is ready to enter the critical section. `flag[i] = true` implies that process P_i is ready!





Algorithm for Process P_i

P_i wishes to enter CS & raises its flag!

```
do {  
    flag[i] = true;  
    turn = j; // false & &  
    while (flag[j] && turn == j);  
    enters critical section  
    flag[i] = false;  
    remainder section  
} while (true);
```

Initially
 $\underline{\text{flag}[i] = F}$

$\underline{\text{flag}[j] = F}$

and sets
 $\text{turn} = j$ // next process

False

checks next process

re-
sets flag = F
On exit!





Algorithm for Process P_j

P_j wants to enter CS \rightarrow raise flag $flag[j] = F$

```
do {  
    flag[j] = true;  
    turn = i;  
    while (flag[i] && turn == i); → False  
        ↳ critical section  
        flag[j] = false; ↳ Pj enters  
    remainder section  
} while (true);  
on exit, reset flag[i] = F
```





Algorithm for Process P_i

```
do {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j);  
        critical section  
    flag[i] = false;  
    remainder section  
} while (true);
```

At most one turn, wait!

$P_i \rightarrow$ wants to enter , suppose

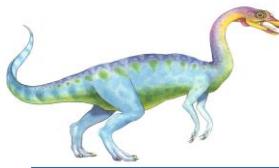
$P_j \rightarrow$ wishes
turn = i

critical section $\rightarrow P_i$ enters

Controversy \rightarrow Mutual
exclusion

Progress ✓





Peterson's Solution (Cont.)

□ Provable that the three CS requirement are met:

1. Mutual exclusion is preserved

P_i enters CS only if:

either `flag[j] = false` or `turn = i`

2. Progress requirement is satisfied
3. Bounded-waiting requirement is met





Synchronization Hardware

(Solution)

- Many systems provide hardware support for implementing the critical section code.
- All solutions below based on idea of **locking** (Process enter, lock)
 - Protecting critical regions via locks
- Uniprocessors – could disable interrupts
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems
 - ▶ Operating systems using this not broadly scalable
- Modern machines provide special atomic hardware instructions
 - ▶ **Atomic** = non-interruptible
 - Either test memory word and set value
 - Or swap contents of two memory words





Solution to Critical-section Problem Using Locks

```
do {  
    → acquire lock → entry section  
    ↘ critical section  
    → release lock → exit section  
    ↙ remainder section  
} while (TRUE);
```

Atomic





test_and_set Instruction

rv=0 ✓

Definition:

```
boolean test_and_set (boolean *target)
```

```
{
```

long vs
short vms

```
    boolean rv = *target;
```

```
    *target = TRUE;
```

```
    return rv;
```

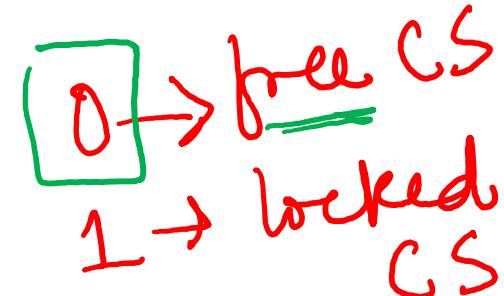
```
}
```

→ address of a
mem loc.

→ passing mem
loc.

original value

1. Executed atomically
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to “TRUE”.





Solution using test_and_set()

- Shared Boolean variable lock, initialized to FALSE
- Solution:

```
do {  
    while (test_and_set(&lock)) /*while true*/  
        ; /* do nothing */  
        /* critical section */  
    lock = false;  
        /* remainder section */  
} while (true);
```





compare_and_swap Instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
    return temp;  
}
```

1. Executed atomically
2. Returns the original value of passed parameter “value”
3. Set the variable “value” the value of the passed parameter “new_value” but only if “value” ==“expected”. That is, the swap takes place only under this condition.





Solution using compare_and_swap

- Shared integer “lock” initialized to 0;
- Solution:

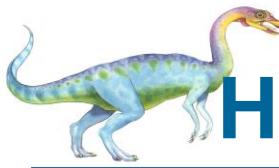
```
do {  
    while (compare_and_swap(&lock, 0, 1) !=  
0)  
        ; /* do nothing */  
    /* critical section */  
    lock = 0; ——————→ GLOBAL VARIABLE  
    /* remainder section */  
    INITIALIZED TO ZERO  
} while (true);
```





BREAK.

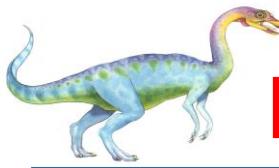




How to satisfy all critical section requirements?

- Hardware
- Software





Bounded-waiting Mutual Exclusion with test_and_set

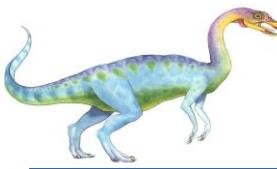
```
do {  
    waiting[i] = true;           ✓  
    key = true;                 ✓  
    while (waiting[i] && key)  
        key = test_and_set(&lock);  
    waiting[i] = false;  
    /* critical section */  
    j = (i + 1) % n;           // 2  
    {  
        while ((j != i) && !waiting[j])  
            j = (j + 1) % n;
```

$\pi = \{1, 2, 5\}$
 π_p

(n - 1)

```
if (j == i)  
    lock false;  
else  
    waiting[j] = false;  
    /* remainder section */  
} while (true);
```





Mutex Locks (MUTUAL EXCLUSION)

- Previous **hardware solutions** are complicated and generally inaccessible to **application programmers**
- **OS designers** build software tools to solve critical section problem
- **Simplest is mutex lock**
 - Protect a critical section by first **acquire()** a lock then **release()** the lock
 - Prevents what? **RACE CONDITION**
 - Boolean variable **(available)** indicating if lock is available or not





Mutex Locks

- Calls to **acquire()** and **release()** must be atomic
 - Usually implemented via hardware atomic instructions
- But this solution **requires busy waiting**
 - This lock therefore is called a **spinlock**





acquire() and release()

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
  
    if available → acquire succeeds  
    & available = false;  
}  
  
release() {  
    available = true;  
}
```

if lock is available / not ?

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

Process is blocked
if unavailable !





Semaphore (*SPW solution*)

- Synchronization tool that provides more sophisticated ways (not Mutex locks) for process to synchronize activities -- > Semaphore **S** – integer variable

- Can only be accessed via two indivisible (atomic) operations

- **wait()** and **signal()**

- Definition of the **wait()** operation

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

↓ Semaphore

Entry → wait

- Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```

↑ Semaphore





Semaphore Usage

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
 - Same as a **mutex lock**
- **Can solve** various synchronization problems
- Consider P_1 and P_2 that require S_1 to happen before S_2
Create a semaphore “**syn**” initialized to 0

P1:

```
S1;  
signal(syn);
```

P2:

```
wait(syn);  
S2;
```

- Can implement a counting semaphore **S** as a binary semaphore





Semaphore Implementation

- Must guarantee that no two processes can execute the `wait()` and `signal()` on the same semaphore at the same time
- Thus, the implementation becomes the critical section problem where the `wait` and `signal` code are placed in the critical section
 - Could now have **busy waiting** in critical section implementation
 - ▶ But implementation code is short
 - ▶ Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution

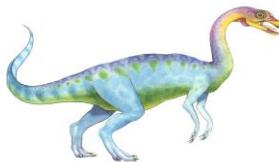




Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue
 - Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list
 - Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
 - **typedef struct{**
 int value;
 struct process *list;
} semaphore;
- (Alternative)
- To overcome Busy waiting





Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
    signal(semaphore *S) {  
        S->value++;  
        if (S->value <= 0) {  
            remove a process P from S->list;  
            wakeup(P);  
        }  
    }  
}
```





Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let s and Q be two semaphores initialized to 1

P_0

```
wait(s) ;  
wait(Q) ;  
...  
signal(S) ;  
signal(Q) ;
```

P_1

```
wait(Q) ;  
wait(S) ;  
...  
signal(Q) ;  
signal(S) ;
```

Sloboh

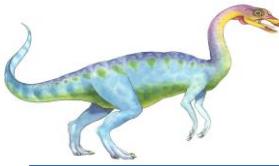
- **Starvation – indefinite blocking**
 - A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
 - Solved via **priority-inheritance protocol**



Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
 - Bounded-Buffer Problem
 - Readers and Writers Problem
 - Dining-Philosophers Problem





Bounded-Buffer Problem

- n buffers, each can hold one item
- Semaphore mutex initialized to the value 1
- Semaphore full initialized to the value 0
- Semaphore empty initialized to the value n





Bounded Buffer Problem (Cont.)

- The structure of the producer process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```



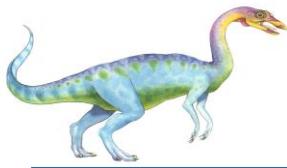


Bounded Buffer Problem (Cont.)

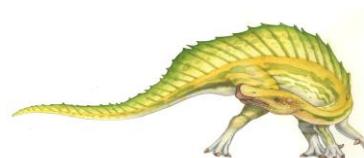
- Structure of the consumer process

```
do {  
    wait(full);  
    wait(mutex);  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    /* consume the item in next_consumed */  
    ...  
} while (true);
```





BREAK!

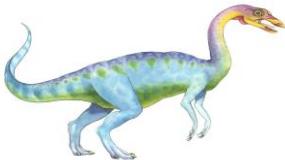




Readers-Writers Problem

- A **data set** is **shared** among a **number of concurrent processes**
 - **Readers:** only read data set; they do *not* perform any updates
 - **Writers:** can both read and write
- Problem: allow multiple readers to read at the same time
 - Only one single writer can access shared data at the same time
- Several variations of how readers and writers are considered – all involve some form of priorities
- Shared Data
 - Data set
 - Semaphore **`rw_mutex`** initialized to 1
 - Semaphore **`mutex`** initialized to 1
 - Integer **`read_count`** initialized to 0





Readers-Writers Problem - Solution

- The structure of a **writer** process

```
do  {  
    wait(rw_mutex);           rw=1; rw<=0 false  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);        1  
} while (true);
```





Readers-Writers Problem - Solution

Writer 1

```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```

$rW < 0$
No
 $rW = 0$

Writer 2

```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```

$rW = -$
 $rW < 0$
busy wait





Readers-Writers Problem (Cont.) Writer writing

- Structure of a reader process:

```
do {  
    wait(mutex); // 0, continues  
    read_count++;  
    if (read_count == 1) // 1  
        wait(rw_mutex); // first reader?  
    signal(mutex); // rw_mutex = -1 ✗  
    mutex=0 ...  
    /* reading is performed */  
    ...  
    wait(mutex); → mutex-- 1  
    read_count--; → 0 true  
    if (read_count == 0) // last reader  
        signal(rw_mutex); // last reader  
        signal(mutex); // 1 ✓  
} while (true);
```

2 Semaphores

$$rw_mutex = 1$$

$$mutex = 1 \leftarrow$$

$$rw_mutex = 0 \}$$

* Writer completes

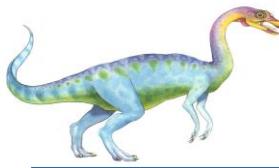


Readers-Writers Problem Variations

- First variation – no reader kept waiting unless writer has permission to use shared object
- Second variation – once writer is ready, it performs the write ASAP
- Both may have starvation leading to even more variations
- Problem is solved on some systems by kernel providing reader-writer locks

Check for reader1 & reader2?





Dining-Philosophers Problem

- Philosophers spend their lives alternating thinking and eating
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
 - Need both to eat, then release both when done
- In the case of 5 philosophers
 - Shared data
 - Bowl of rice (data set)
 - Semaphore chopstick[5] initialized to 1



makeagif.com



Dining-Philosophers Problem Algorithm

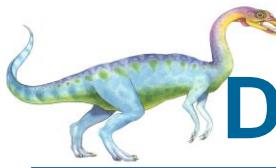
□ The structure of Philosopher *i*:

```
do {  
    wait (chopstick[i] );  
    wait (chopStick[ (i + 1) % 5] );  
        // eat  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
        // think  
} while (TRUE);
```



What is the problem with this algorithm?





Dining-Philosophers Problem Algorithm (Cont.)

□ Deadlock handling

- Allow at most 4 philosophers to be sitting simultaneously at the table.
- Allow a philosopher to pick up the forks only if both are available (picking must be done in a critical section)
- Use an asymmetric solution -- an odd-numbered philosopher picks up first the left chopstick and then the right chopstick. Even-numbered philosopher picks up first the right chopstick and then the left chopstick.





Problems with Semaphores

- Incorrect use of semaphore operations:
 - signal (mutex) wait (mutex)
 - wait (mutex) ... wait (mutex)
 - Omitting of wait (mutex) or signal (mutex) (or both)
- Deadlock and starvation are possible.



End of Chapter 5

