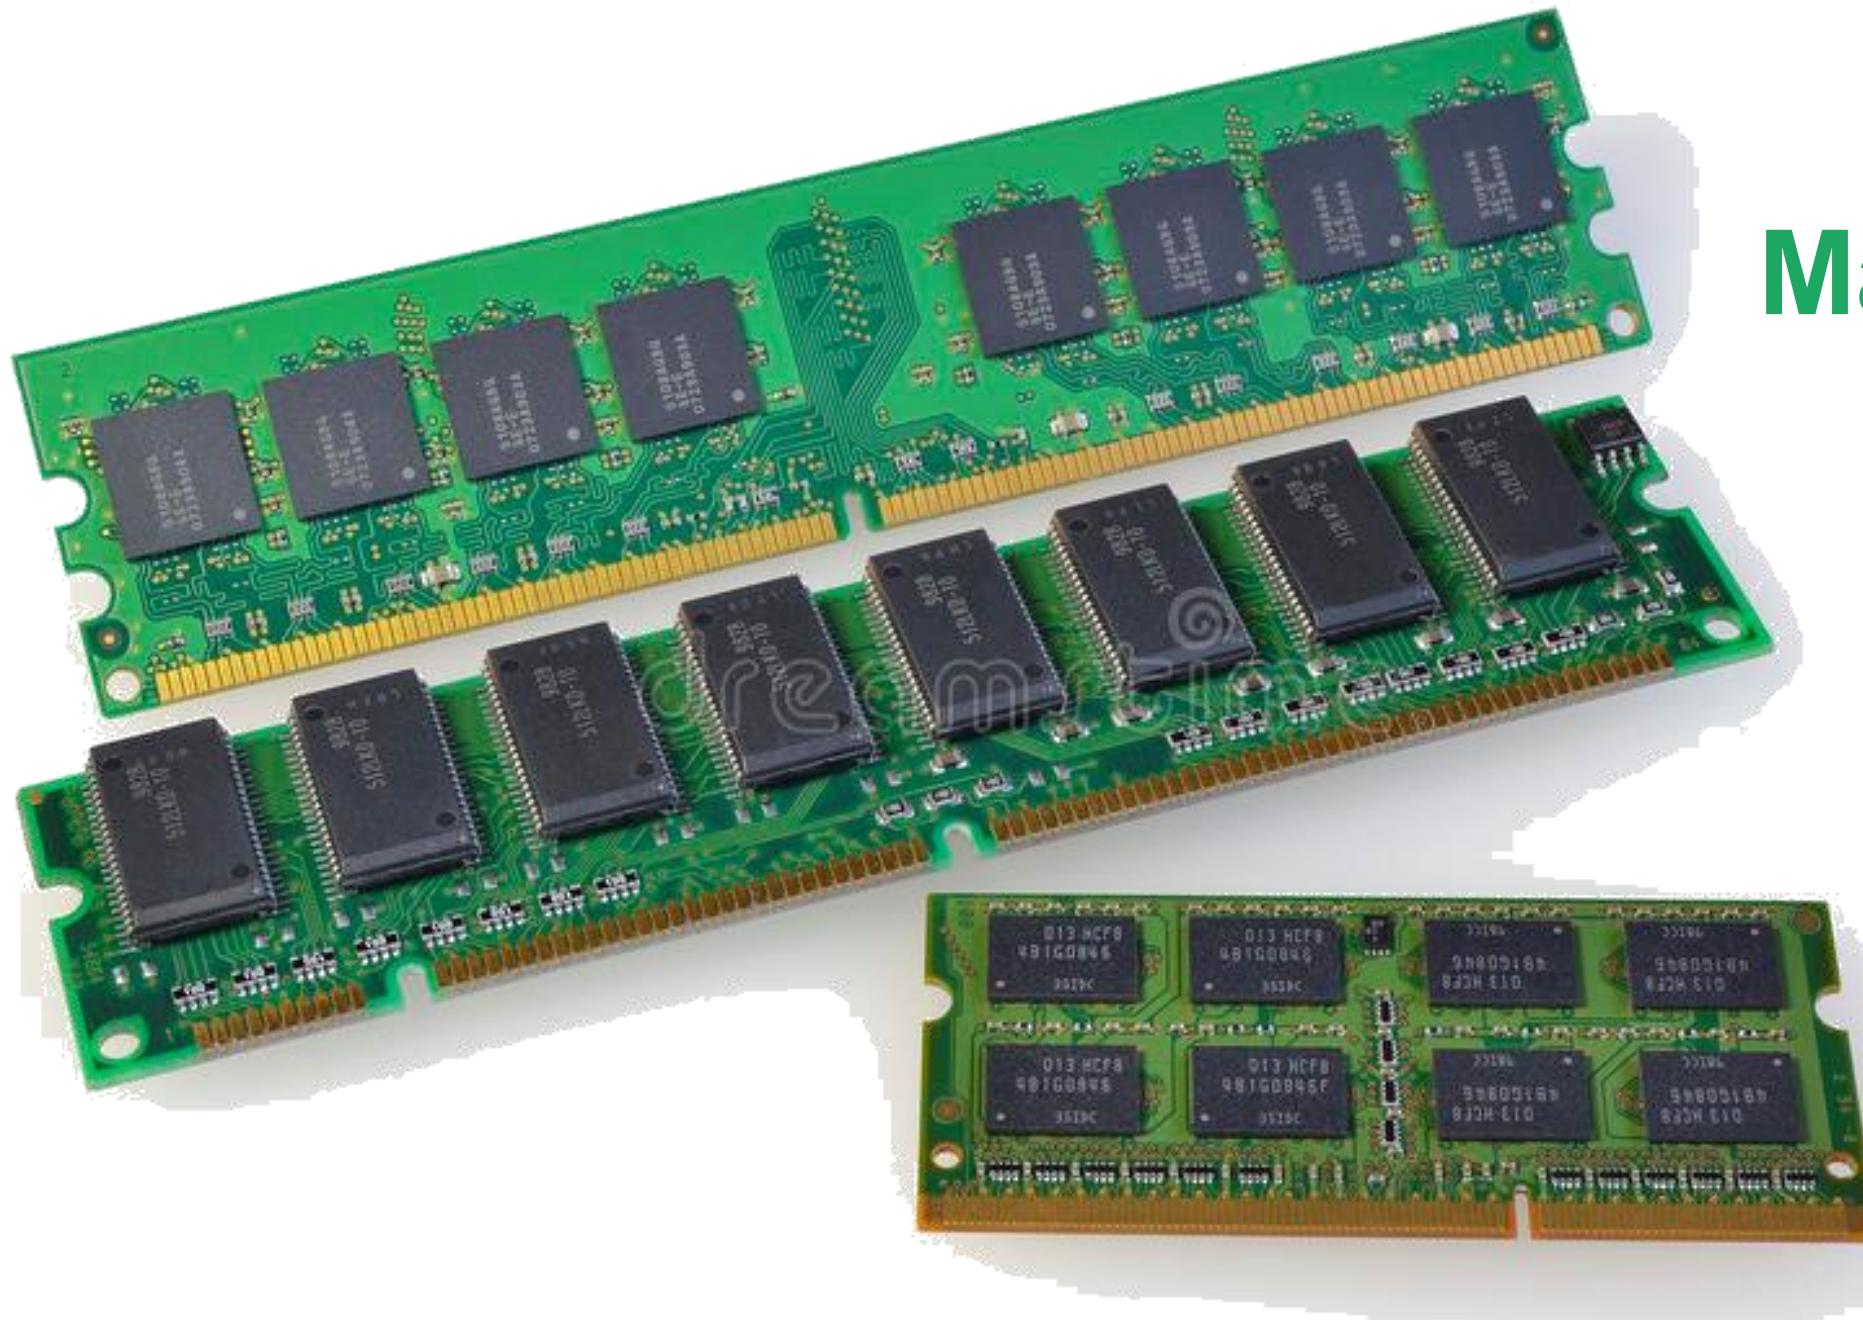


# Chapter 8

## Main Memory



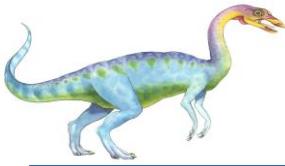


# Chapter 8: Memory Management

---

- Background
- Swapping
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table



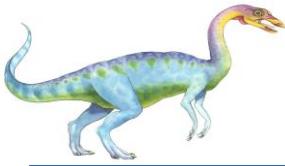


# Objectives

---

- To provide a detailed description of various ways of organizing memory hardware
- To discuss various memory-management techniques, including paging and segmentation

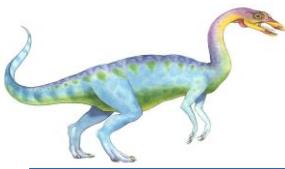




# Background

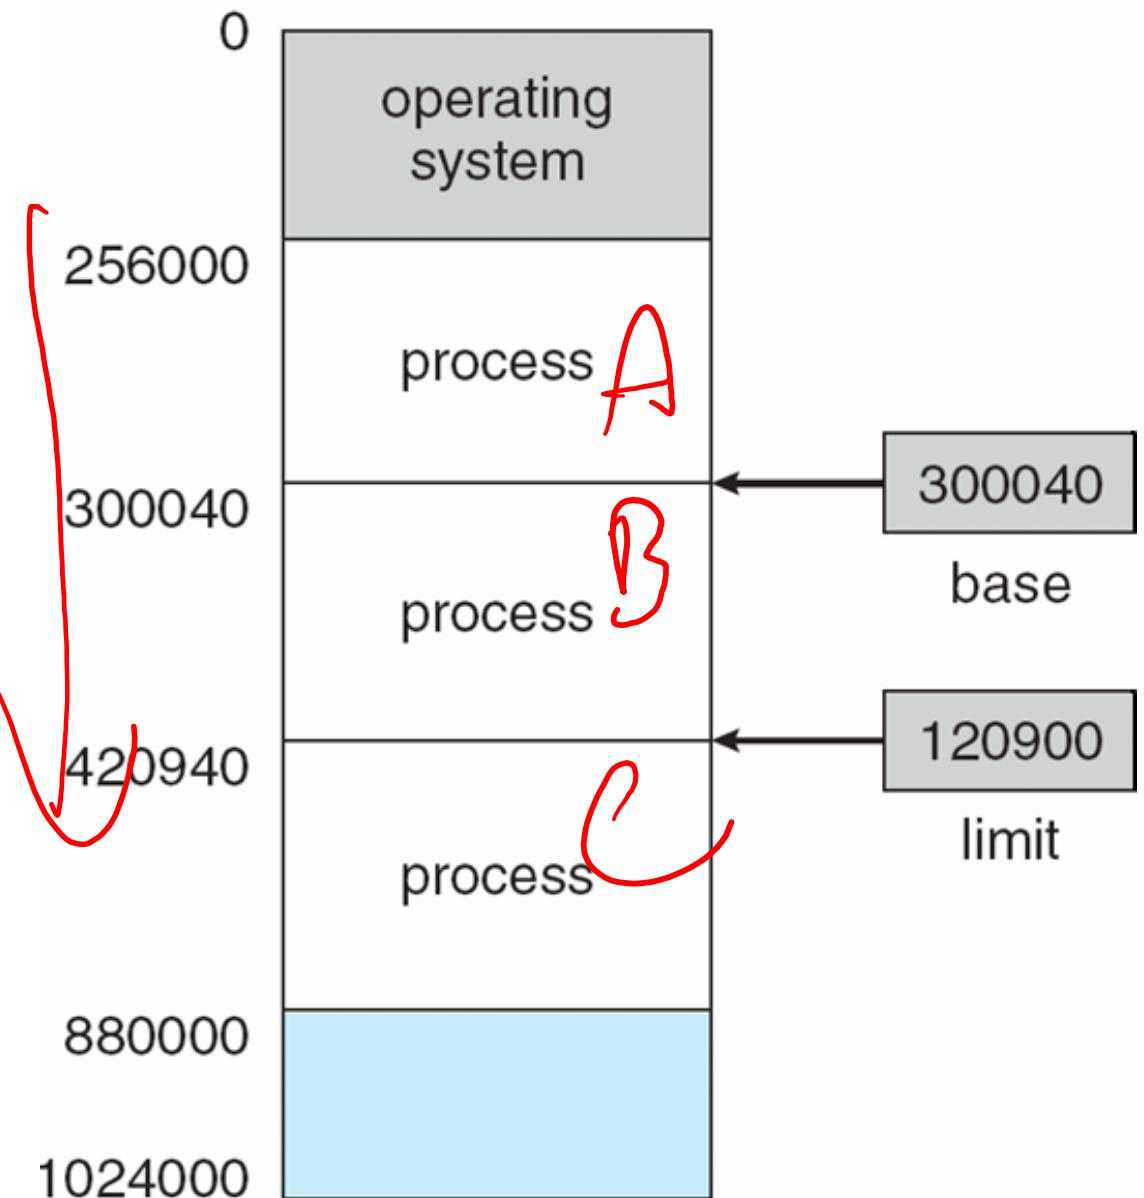
- Program must be brought (from disk) into memory and placed within a process for it to be run → *Execute prgm* —→ *Memory Needed*
- **Main memory and registers** are only storage CPU can access directly
- Memory unit only sees a **stream of addresses** + read requests, or address + data and write requests
- **Register** access in **one CPU clock** (or less)
- Main memory **can take many cycles**, causing a **stall**
- **Cache** sits between main memory and CPU registers
- Protection of memory required to ensure correct operation

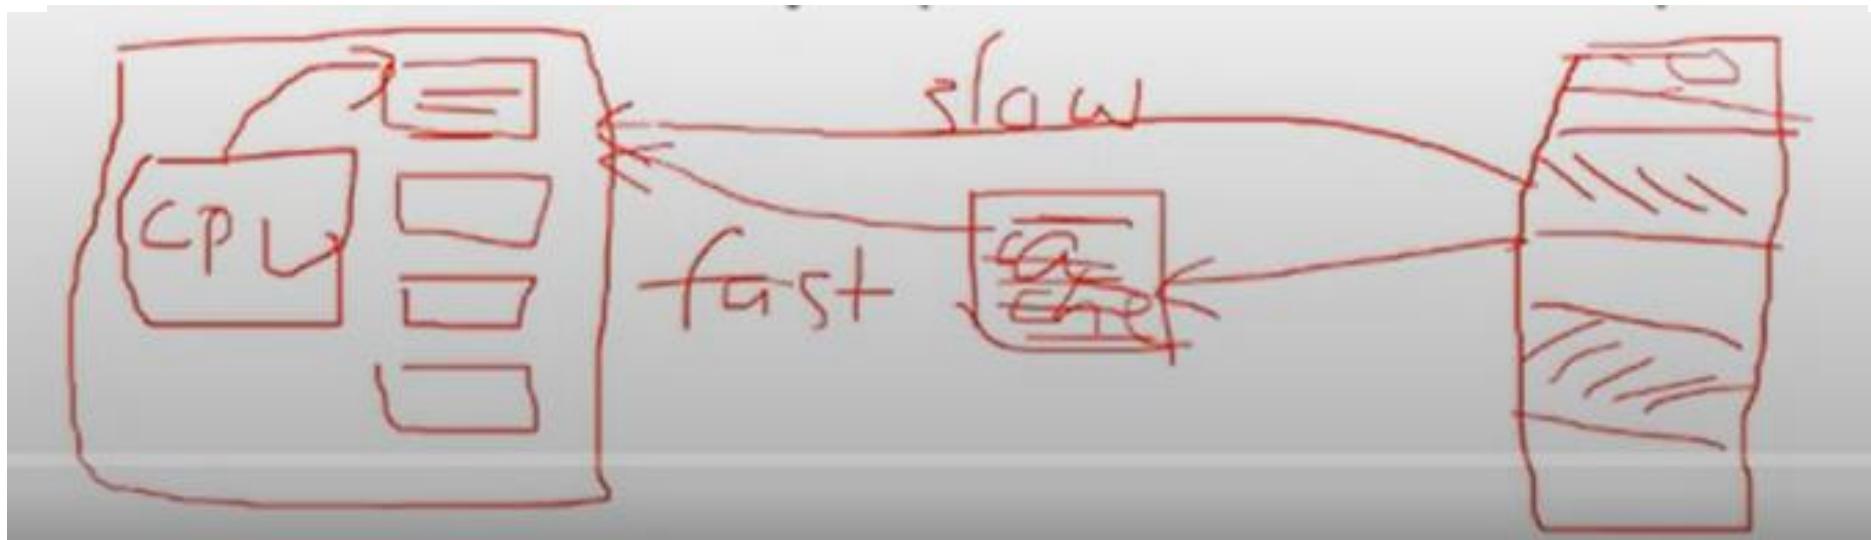




# Base and Limit Registers

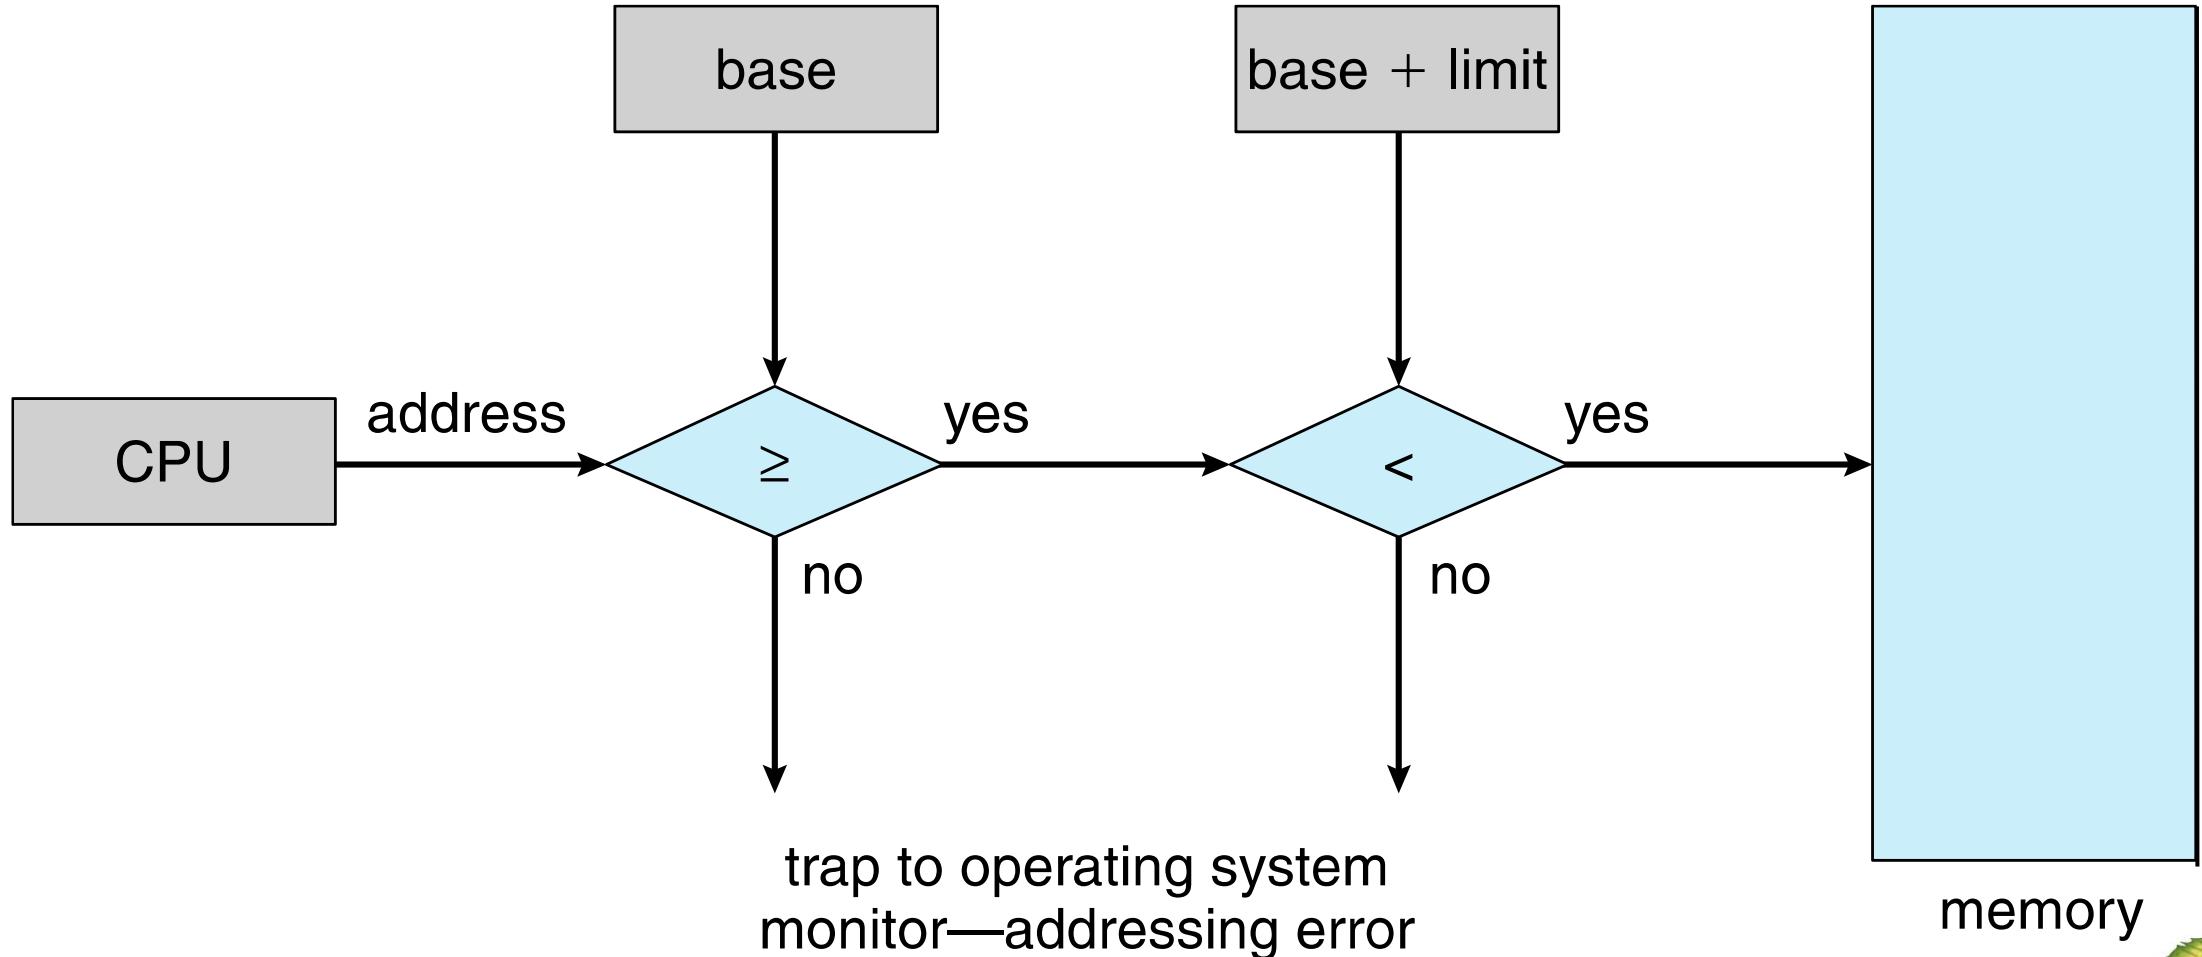
- A pair of **base** and **limit registers** define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user

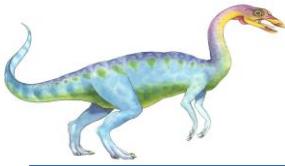






# Hardware Address Protection





# Address Binding

- Programs on disk, ready to be brought into memory to execute from an **input queue**
  - Without support, must be loaded into address 0000
- Inconvenient to have first user process physical address always at 0000
  - How can it not be?
- Further, addresses represented in different ways at different stages of a program's life
  - Source code addresses usually symbolic
  - Compiled code addresses **bind** to relocatable addresses
    - ▶ i.e. “14 bytes from beginning of this module”
  - Linker or loader will bind relocatable addresses to absolute addresses
    - ▶ i.e. 74014
  - Each binding maps one address space to another





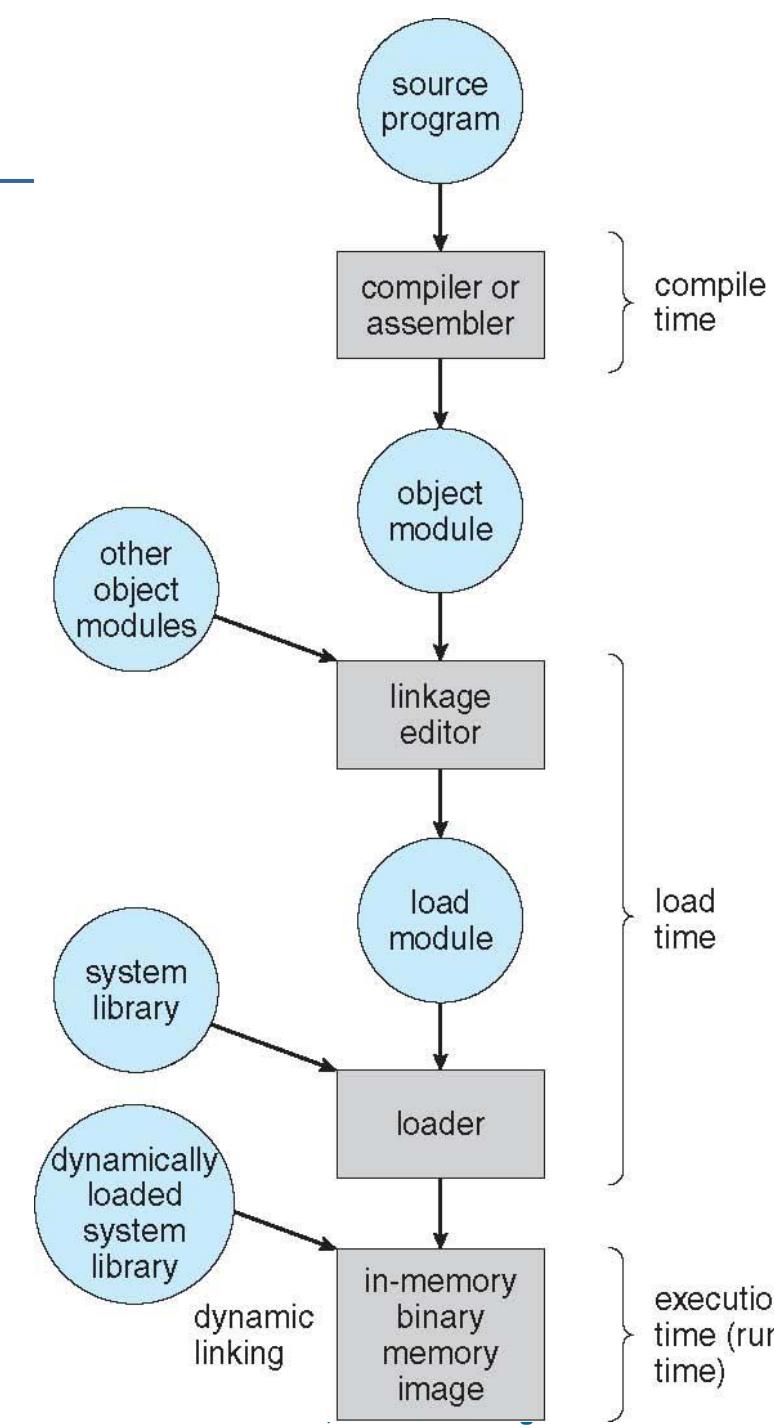
## Binding of Instructions and Data to Memory

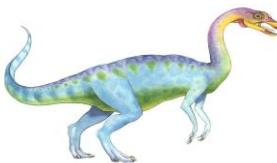
- Address binding of instructions and data to memory addresses can happen at three different stages
  - **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
  - **Load time:** Must generate **relocatable code** if memory location is not known at compile time
  - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
    - ▶ Need hardware support for address maps (e.g., base and limit registers)





# Multistep Processing of a User Program





# Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
  - **Logical address** – generated by the CPU; also referred to as **virtual address**
  - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program





# Memory-Management Unit (MMU)

- Hardware device that at run time maps virtual to physical address
- Many methods possible, covered in the rest of this chapter
- To start, consider simple scheme where the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
  - Base register now called **relocation register**
  - MS-DOS on Intel 80x86 used 4 relocation registers
- The user program deals with *logical* addresses; it never sees the *real*/physical addresses
  - Execution-time binding occurs when reference is made to location in memory
  - Logical address bound to physical addresses



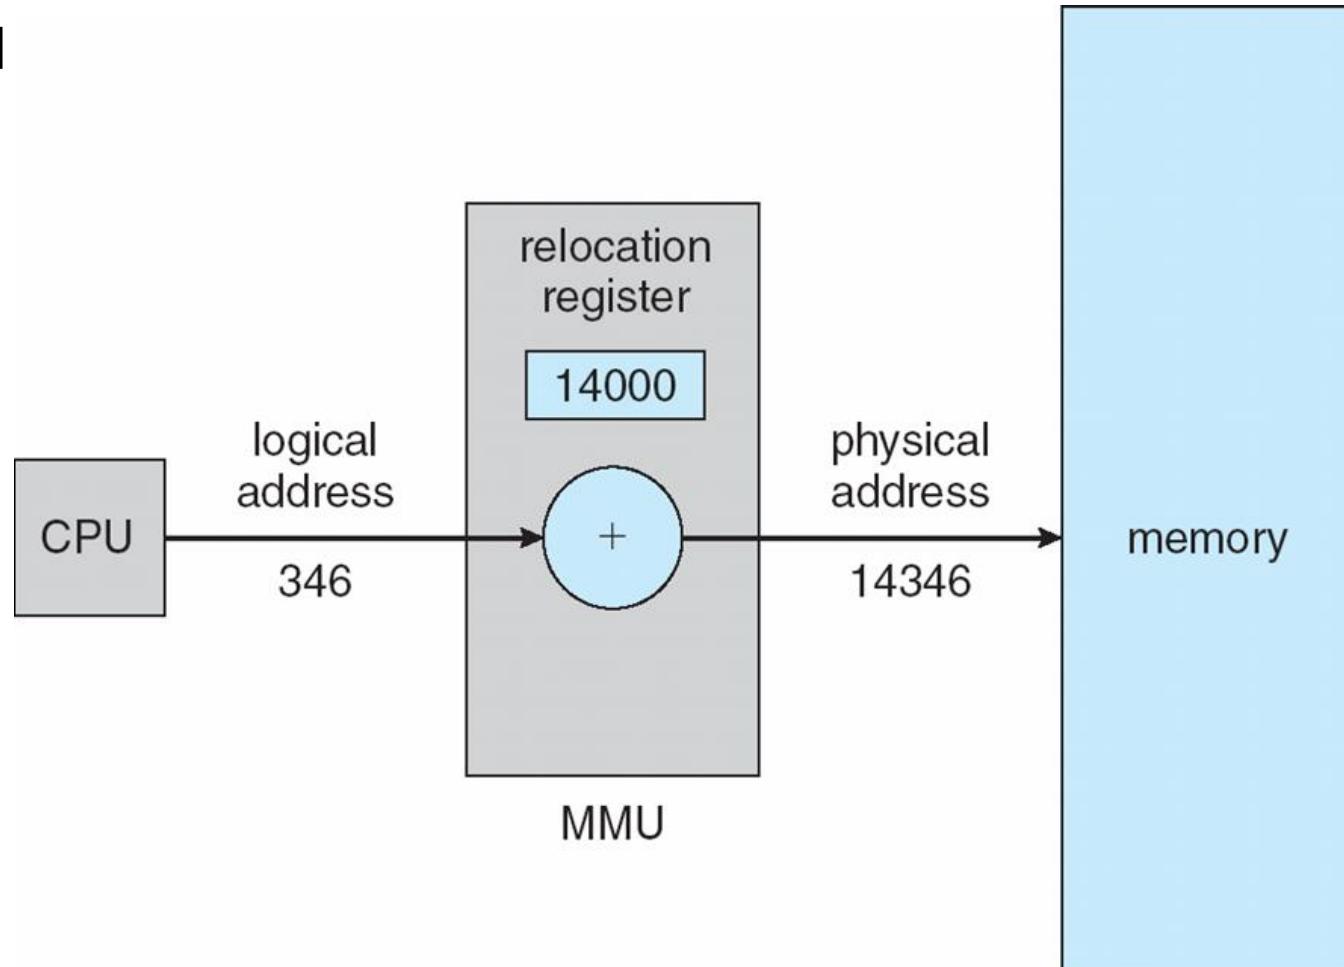


BREAK  
TIME



## Dynamic relocation using a relocation register

- Routine is not loaded until it is called
- Better memory-space utilization; unused routine is never loaded
- All routines kept on disk in relocatable load format
- Useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required
  - Implemented through program design
  - OS can help by providing libraries to implement dynamic loading





# Dynamic Linking

- **Static linking** – system libraries and program code combined by the loader into the binary program image
  - Dynamic linking –linking postponed until execution time
  - Small piece of code, **stub**, used to locate the appropriate memory-resident library routine
  - Stub replaces itself with the address of the routine, and executes the routine
  - Operating system checks if routine is in processes' memory address
    - If not in address space, add to address space
  - Dynamic linking is particularly useful for libraries
  - System also known as **shared libraries**
  - Consider applicability to patching system libraries
    - Versioning may be needed
- ← Big executables*





## Out of memory SWAPPING

Less RAM

- Where does program reside? Secondary memory
- Transfer into main memory from **secondary memory**, why?  
CPU has direct access to main memory
- Size of main memory limited? Yes!
- Degree of multiprogramming to be maintained, how?
  - A process can be **swapped** temporarily out of memory to a backing store (*secondary memory*) and then brought back into memory for continued execution
- What does **swapping achieve**?
  - Total physical memory space of processes can exceed physical memory

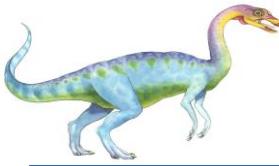




# Swapping

- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- **Major part of swap time is transfer time**
- **Total transfer time** is directly proportional to the amount of memory swapped
- System maintains a **ready queue** of ready-to-run processes which have memory images on disk





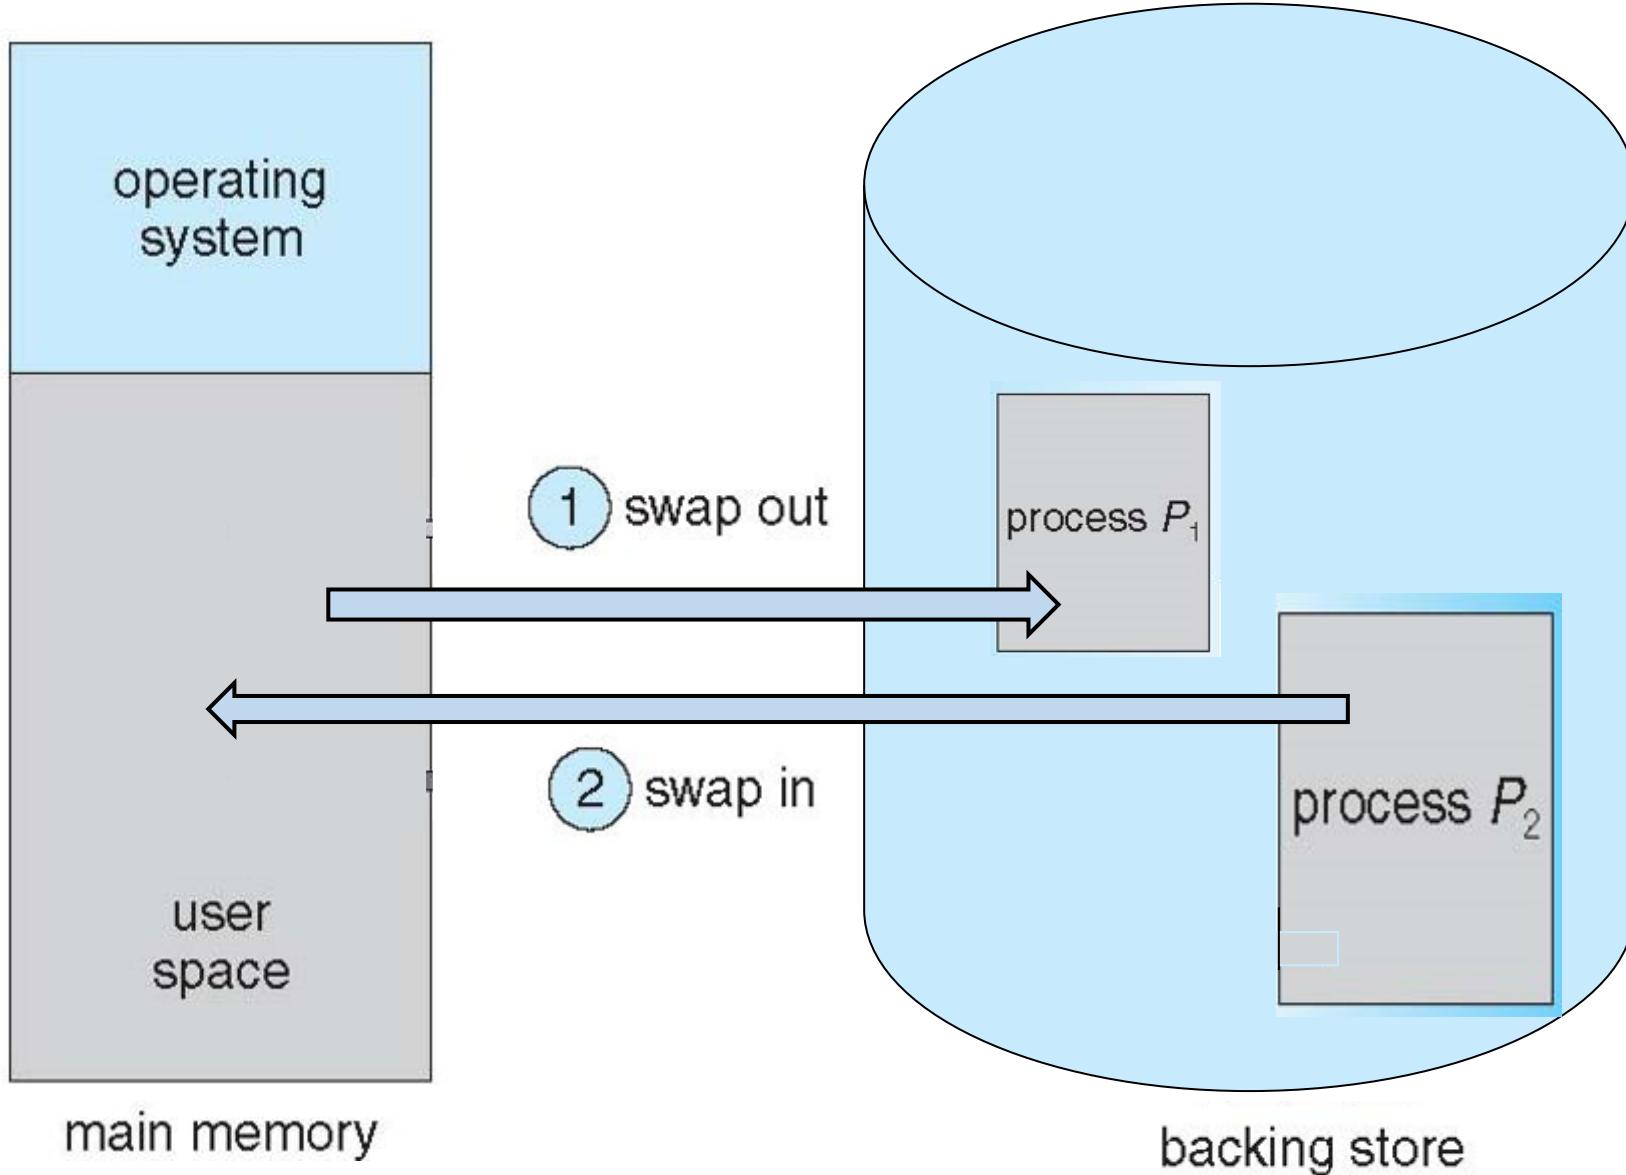
## Swapping (Cont.)

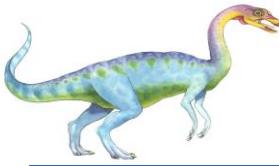
- Does the swapped out process need to swap back in to same physical addresses?
- Depends on address binding method
  - Plus consider pending I/O to / from process memory space
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
  - Swapping normally disabled
  - Started if more than threshold amount of memory allocated
  - Disabled again once memory demand reduced below threshold





# SWAPPING

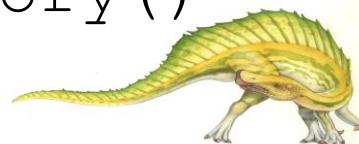




## Context Switch Time including Swapping

---

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be very high
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
  - Swap out time of 2000 ms
  - Plus swap in of same sized process
  - Total context switch swapping component time of 4000ms (4 seconds)
- Can reduce if reduce size of memory swapped – by knowing how much memory really being used
  - System calls to inform OS of memory use via `request_memory()` and `release_memory()`

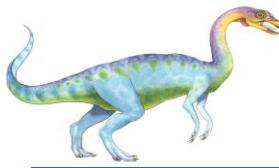




## Context Switch Time and Swapping (Cont.)

- Other constraints as well on swapping
  - Pending I/O – can't swap out as I/O would occur to wrong process
  - Or always transfer I/O to kernel space, then to I/O device
    - ▶ Known as **double buffering**, adds overhead
- Standard swapping not used in modern operating systems
  - But modified version common
    - ▶ Swap only when free memory extremely low





# Memory Allocation Techniques

Contiguous  
allocation

Non-contiguous  
memory allocation



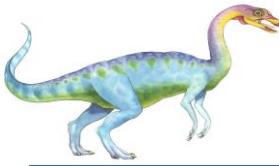


# Contiguous Allocation

---

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method
- Main memory usually into two **partitions**:
  - Resident operating system, usually held in low memory with interrupt vector
  - User processes then held in high memory
  - Each process contained in single contiguous section of memory

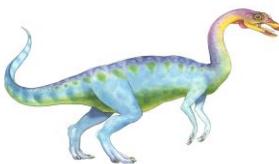




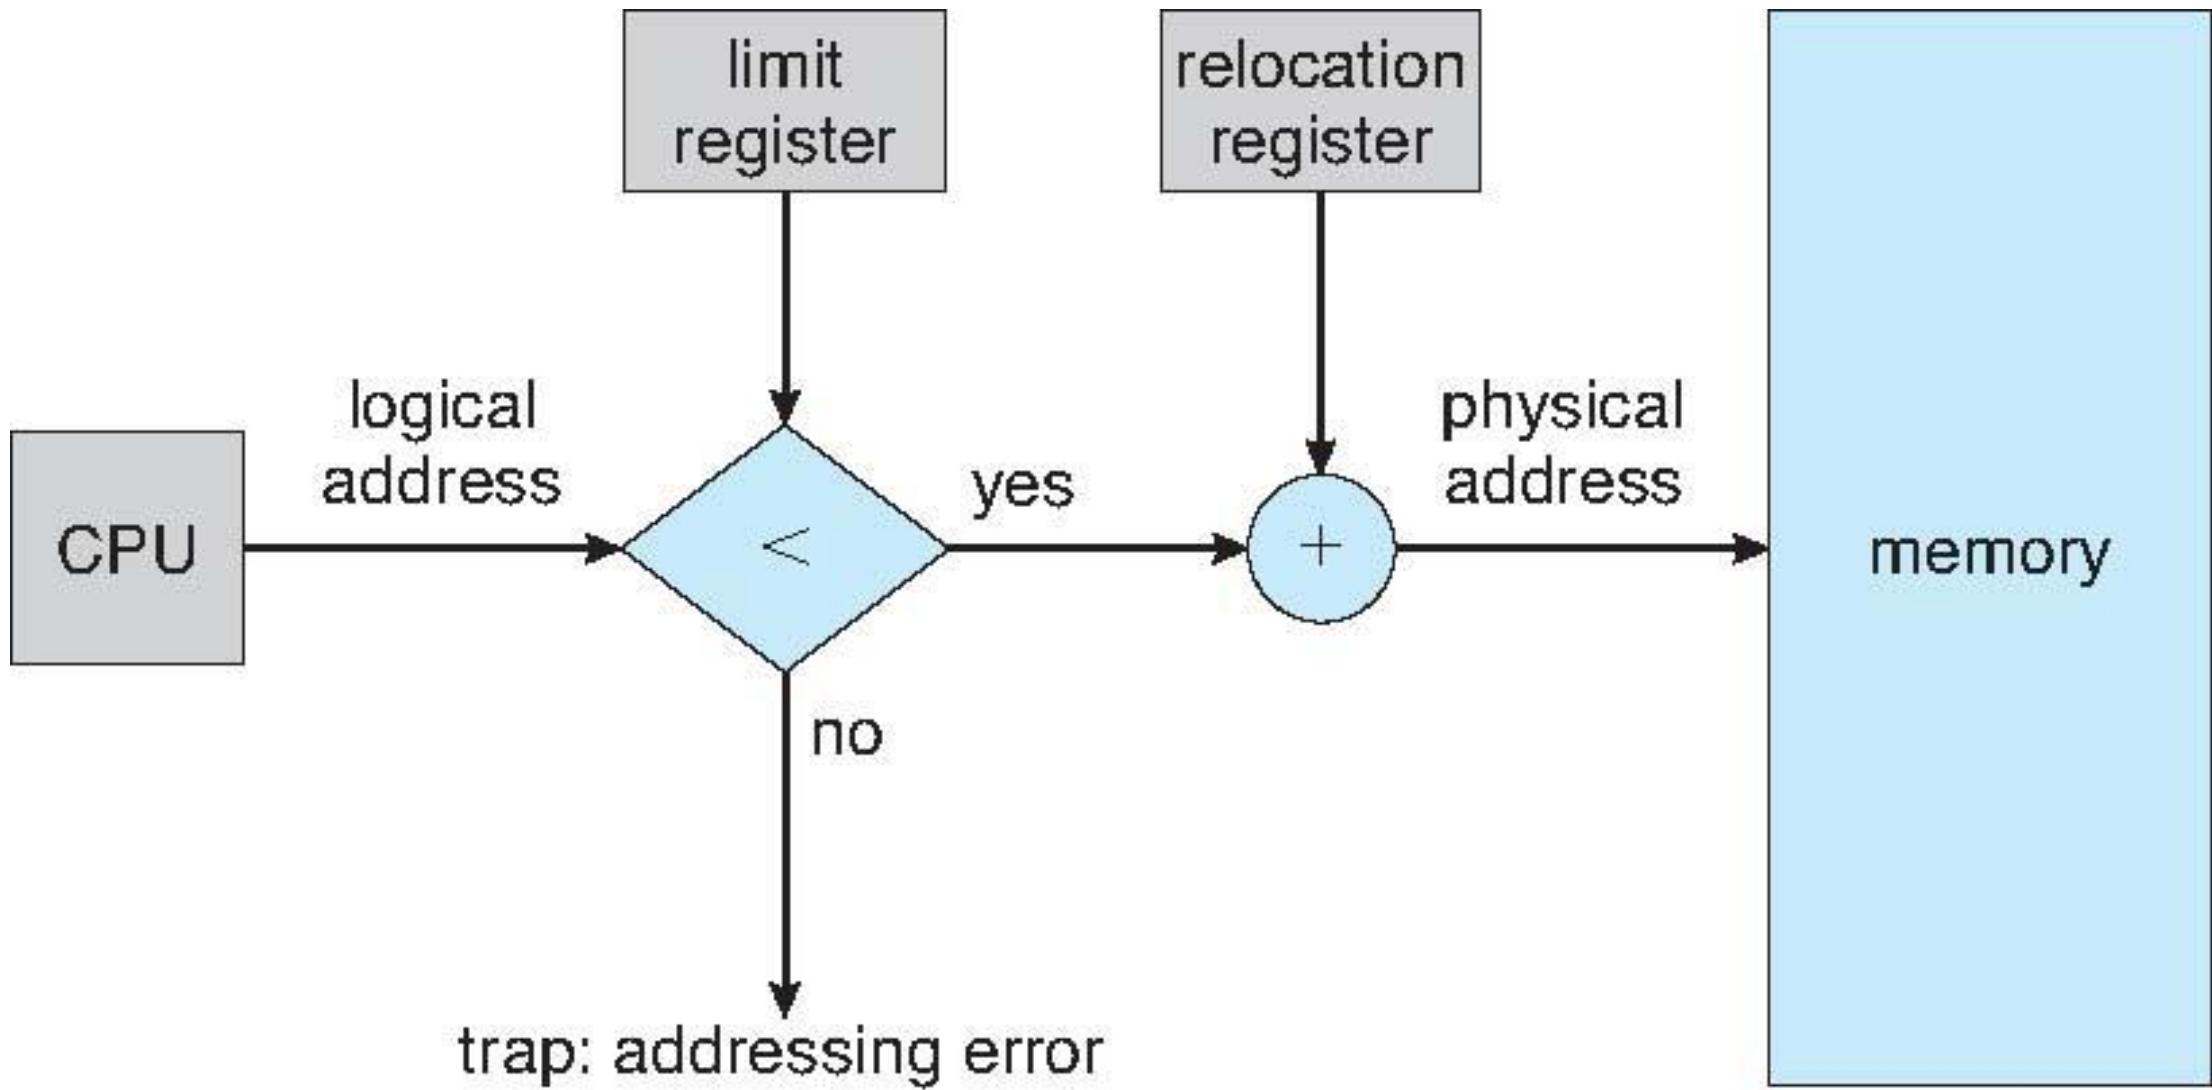
## Contiguous Allocation (Cont.)

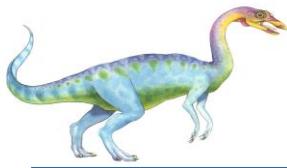
- Relocation registers used to protect user processes from each other, and from changing operating-system code and data
  - Base register contains value of smallest physical address
  - Limit register contains range of logical addresses – each logical address must be less than the limit register
  - MMU maps logical address *dynamically*
  - Can then allow actions such as kernel code being **transient** and kernel changing size



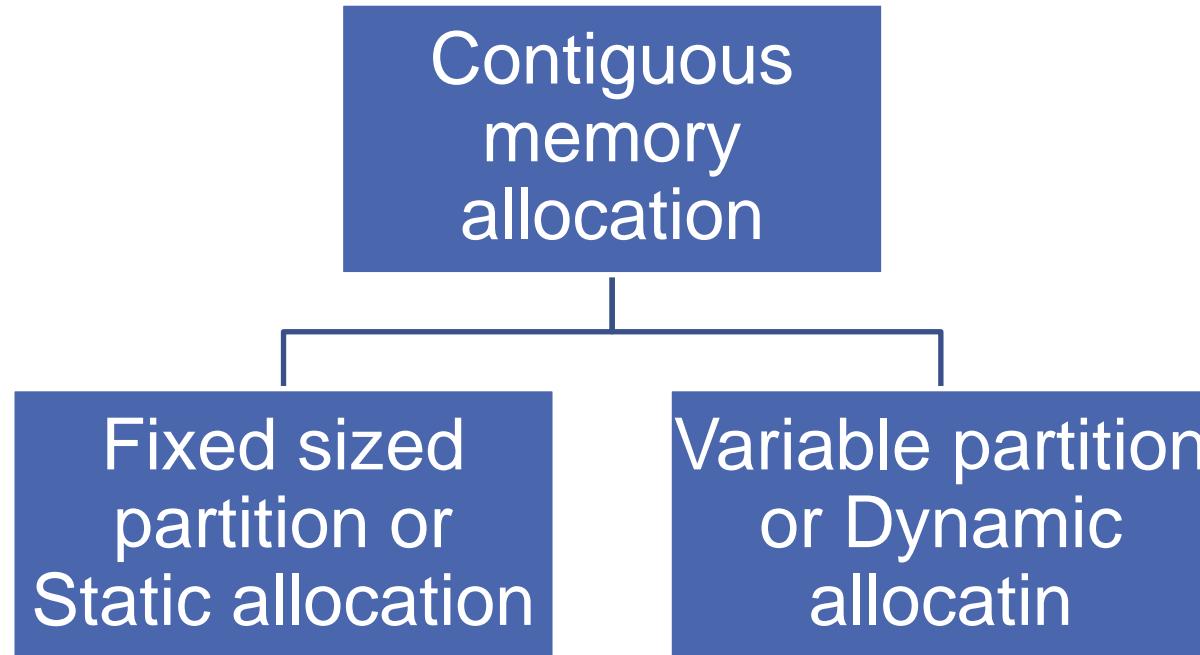


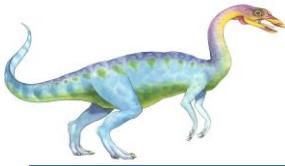
# Hardware Support for Relocation and Limit Registers





# Contiguous Memory Allocation





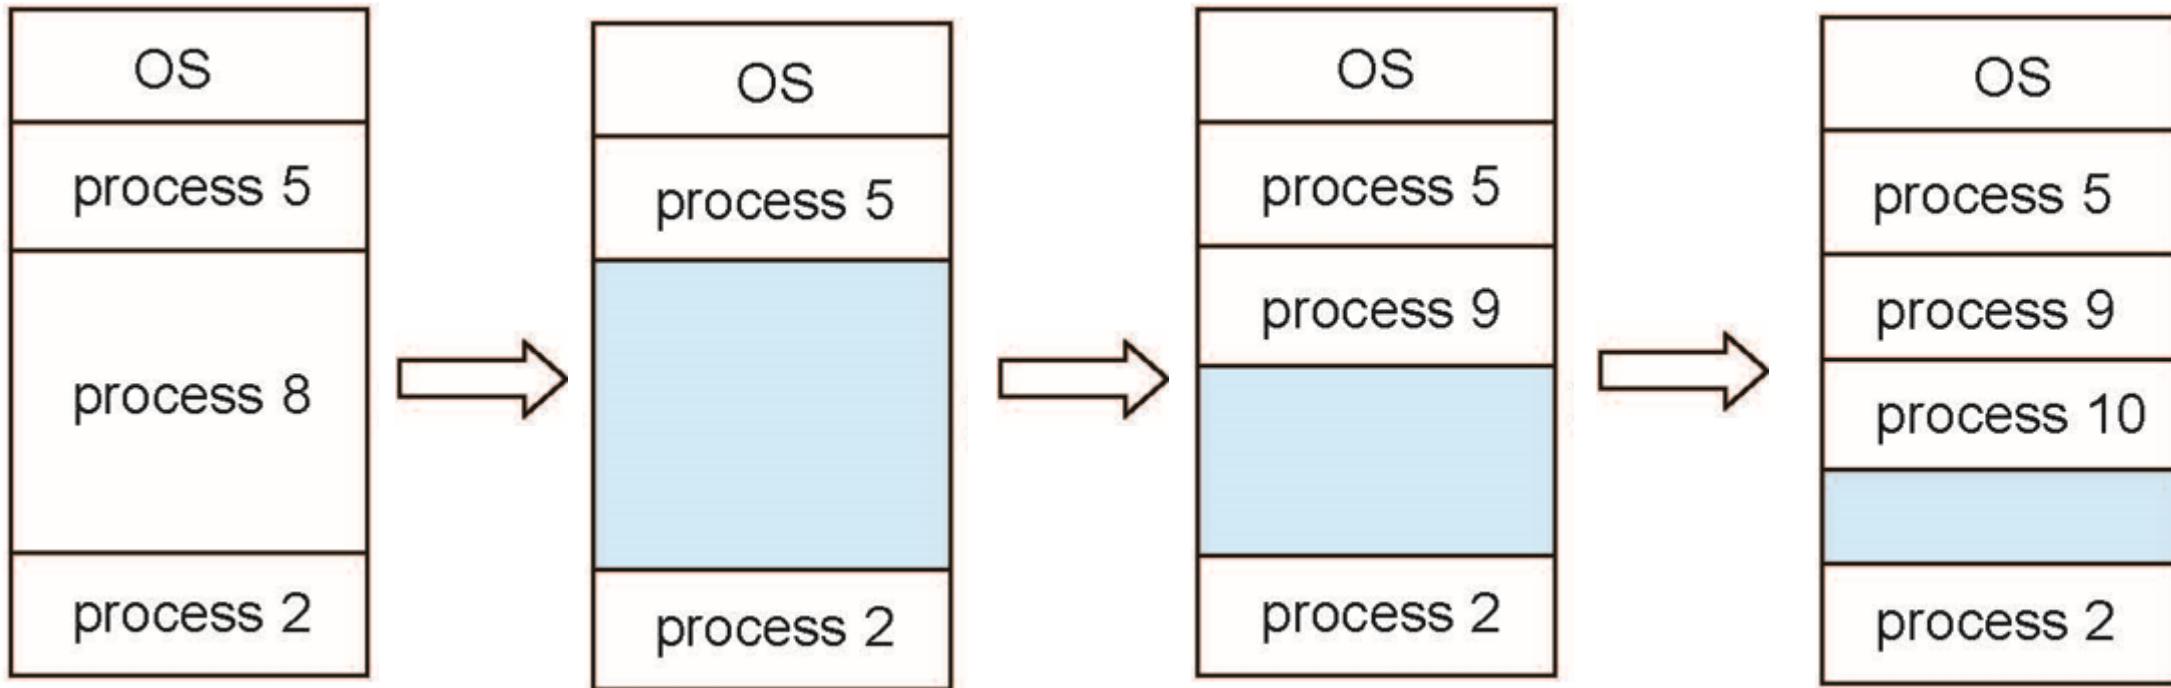
# Multiple-partition allocation

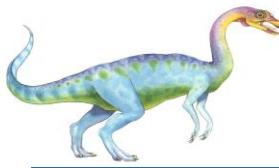
- Multiple-partition allocation
  - Degree of multiprogramming limited by number of partitions
  - **Variable-partition** sizes for efficiency (sized to a given process' needs)
  - **Hole** – block of available memory; holes of various size are scattered throughout memory
  - When a process arrives, it is allocated memory from a hole large enough to accommodate it
  - Process exiting frees its partition, adjacent free partitions combined
  - Operating system maintains information about:
    - a) allocated partitions
    - b) free partitions (hole)





# Multiple-partition allocation





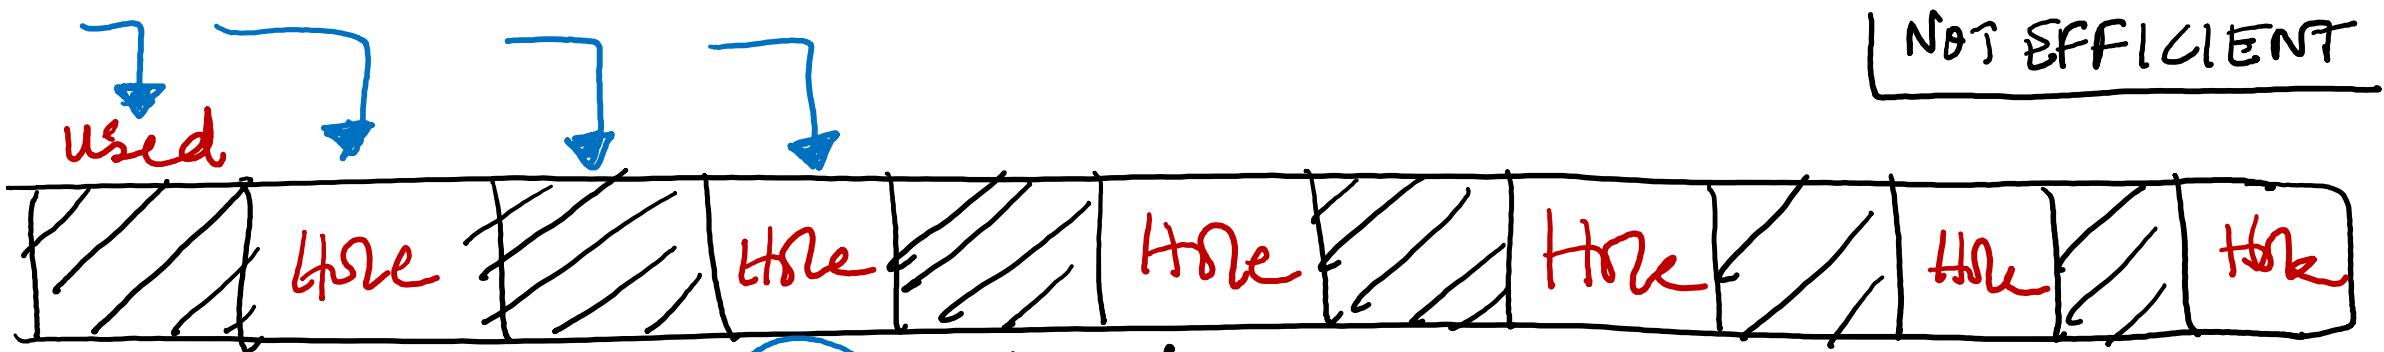
# Dynamic Storage-Allocation Problem

How to satisfy a request of size  $n$  from a list of free holes?

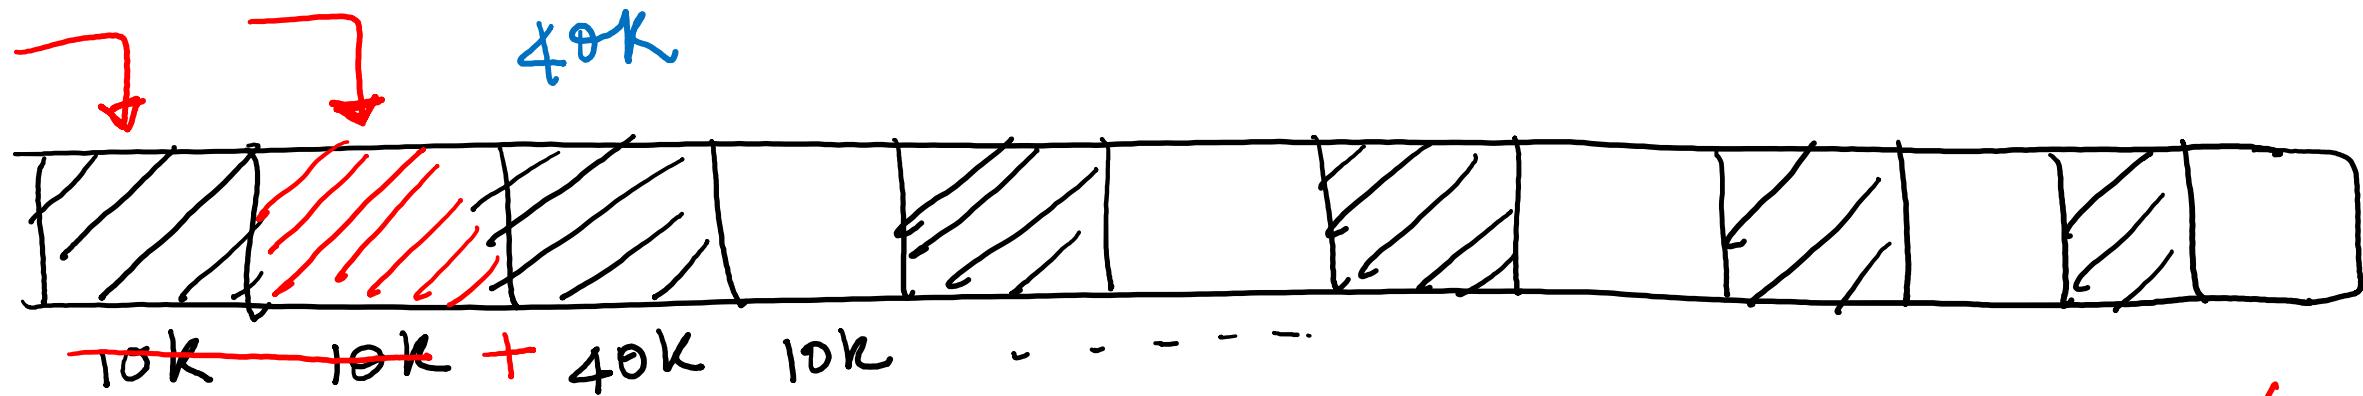
- **First-fit**: Allocate the *first* hole that is big enough
- **Best-fit**: Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size
  - Produces the smallest leftover hole
- **Worst-fit**: Allocate the *largest* hole; must also search entire list
  - Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization





$10K$     $10K$     $20K$     $30K$     $+ 20K$   
 $\underline{10K}$   
 $40K$

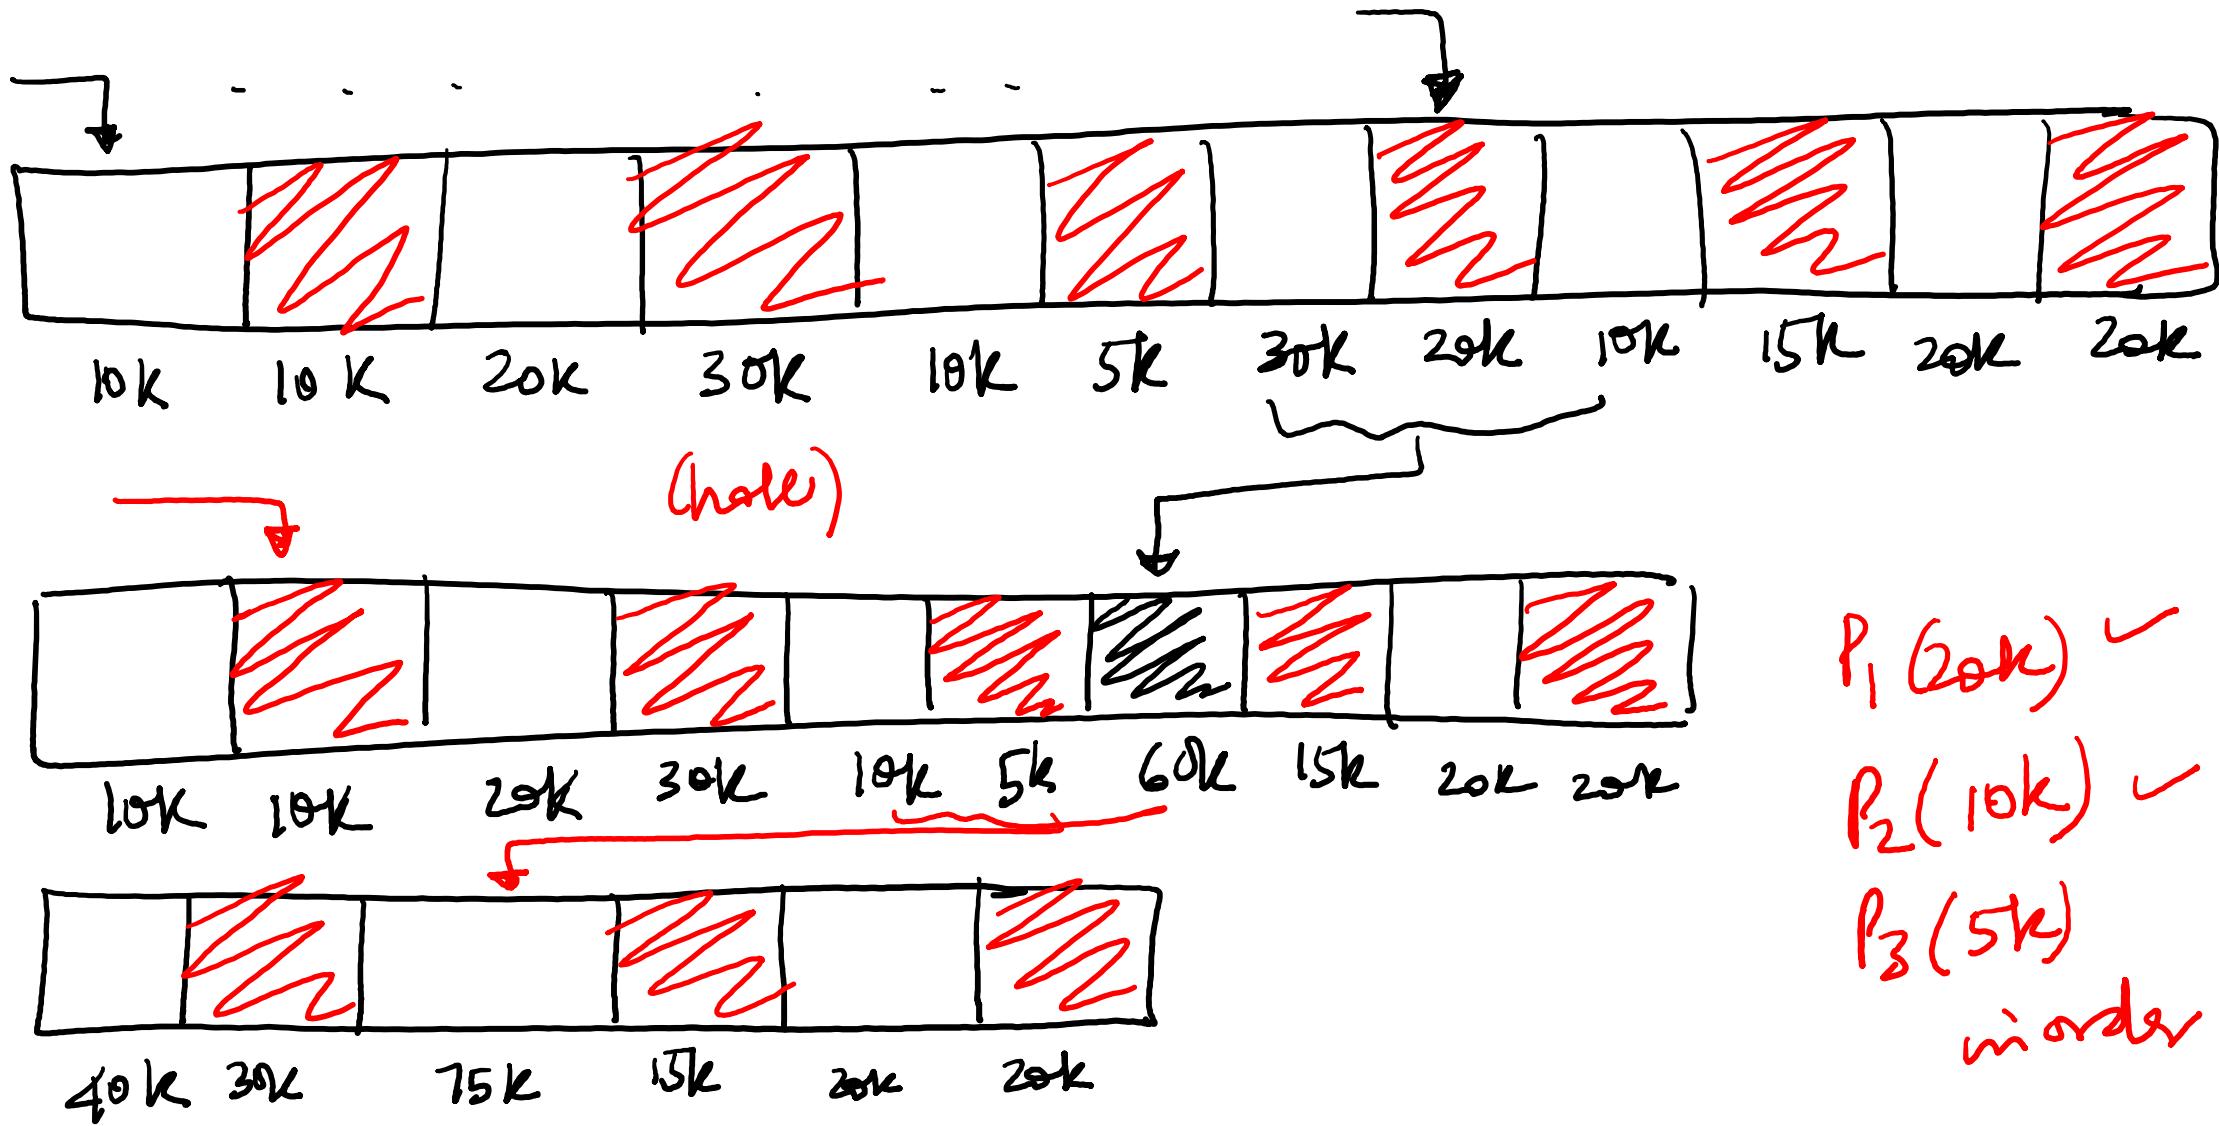


$65K$     $5K$     $\dots \dots \dots$

Request from P<sub>1</sub> (20K)  
 ✓ Request P<sub>2</sub> (10K)  
 Request P<sub>3</sub> (5K)

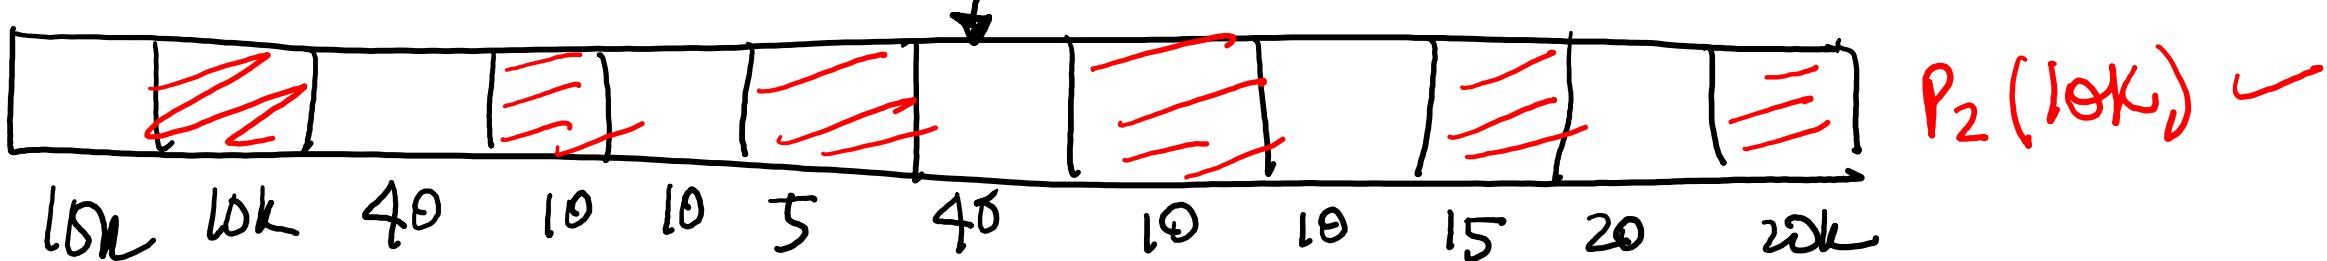
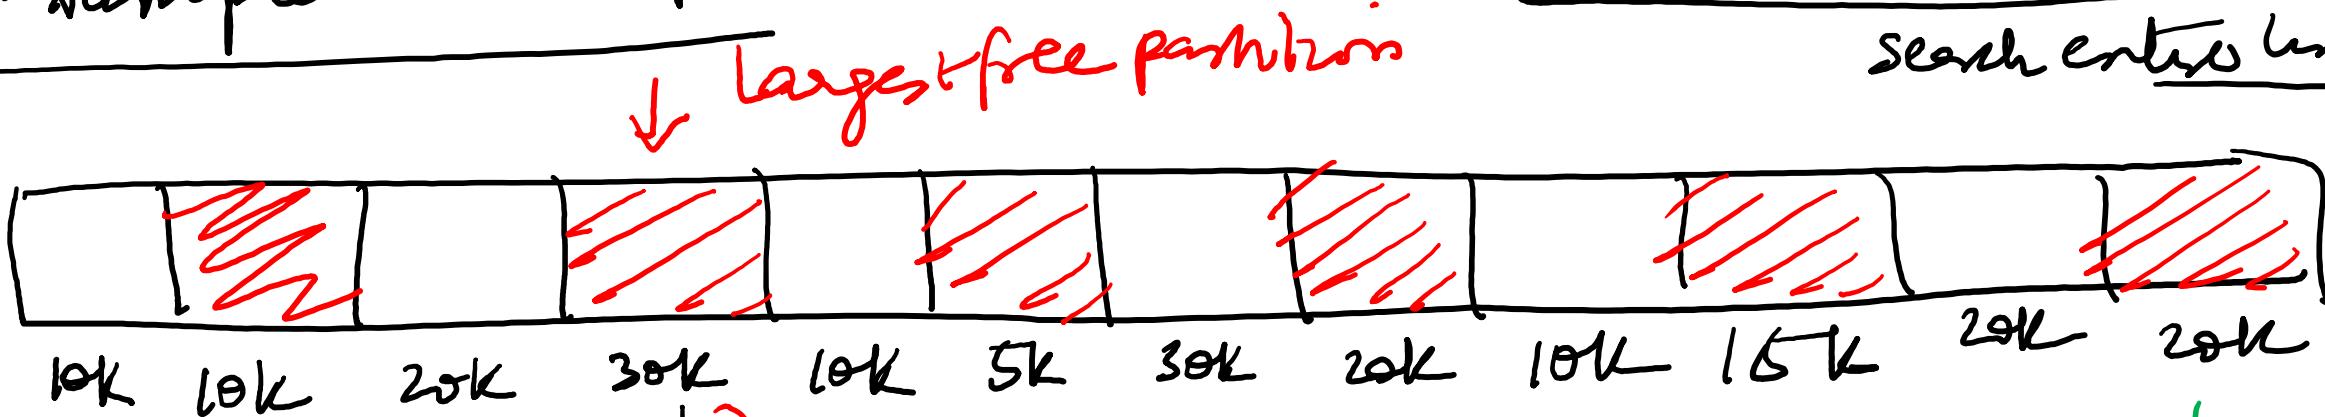
## Example 2 : Best fit

Smallest hole  
big enough



### Example 3: Worst fit

Allocate largest free  
Search entire list





# TAKE A BREAK

---





# Fragmentation

~~185K~~ 25

- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous (*Variable size*)
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used (*fixed size*)
- First fit analysis reveals that given  $N$  blocks allocated,  $0.5N$  blocks lost to fragmentation
  - 1/3 may be unusable -> **50-percent rule**





## Fragmentation (Cont.)

---

- Reduce external fragmentation by **compaction**
  - Shuffle memory contents to place all free memory together in one large block
  - Compaction is possible *only* if relocation is dynamic, and is done at execution time
  - I/O problem
    - ▶ Latch job in memory while it is involved in I/O
    - ▶ Do I/O only into OS buffers
- Now consider that backing store has same fragmentation problems





# Segmentation



- Memory-management scheme that supports user view of memory
- A program is a collection of segments
  - A segment is a logical unit such as:

main program

procedure

function

method

object

local variables, global variables

common block

stack

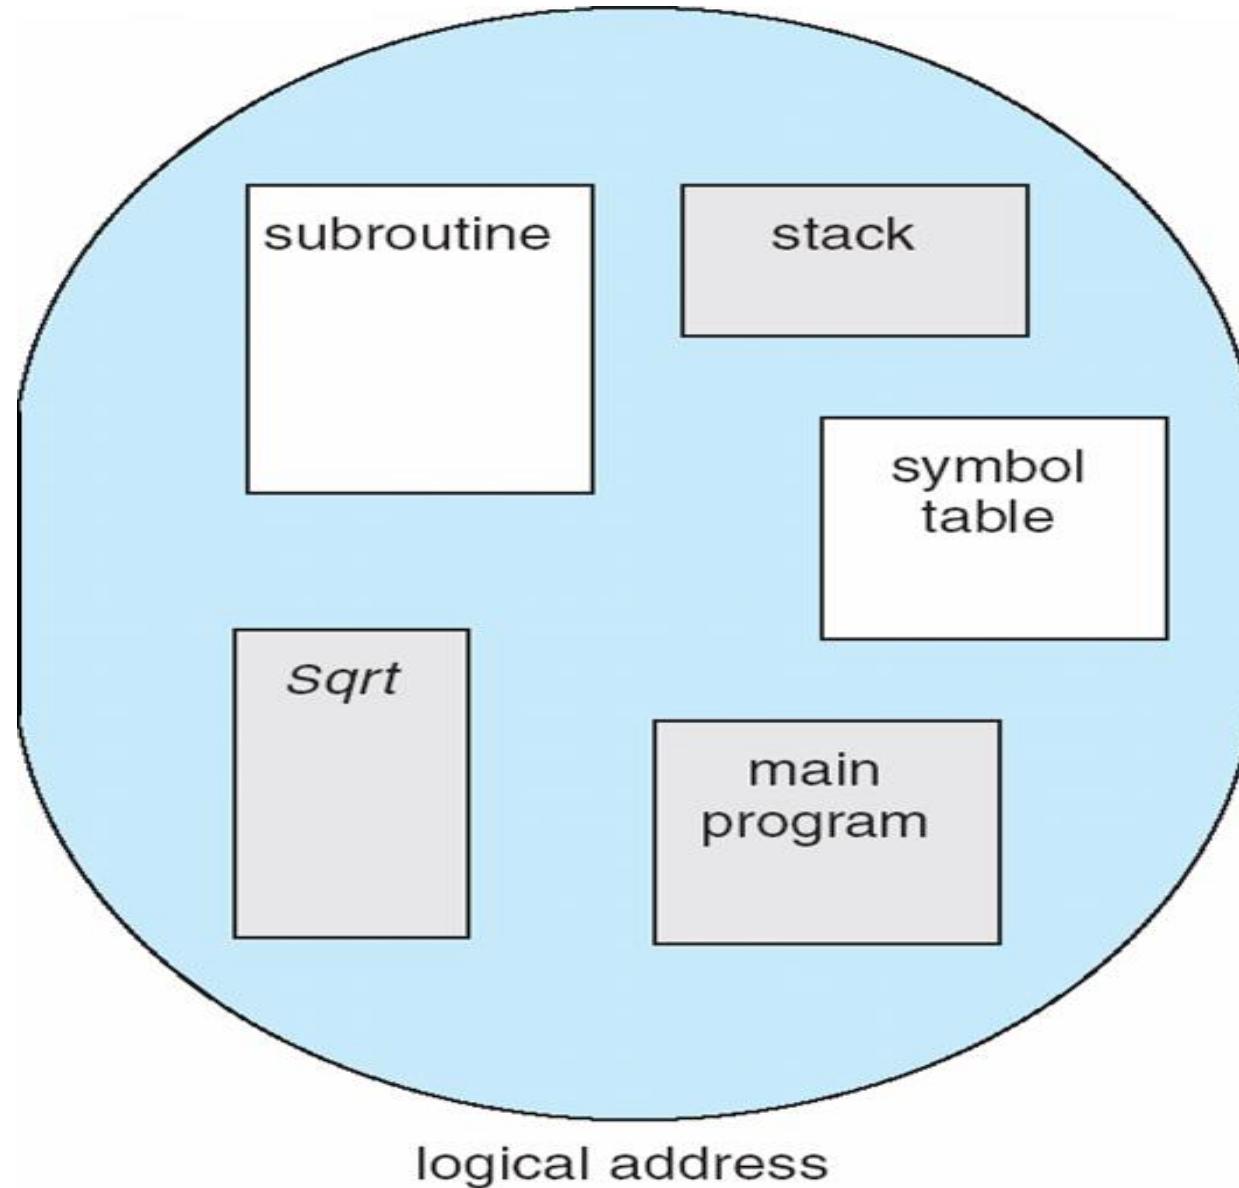
symbol table

arrays



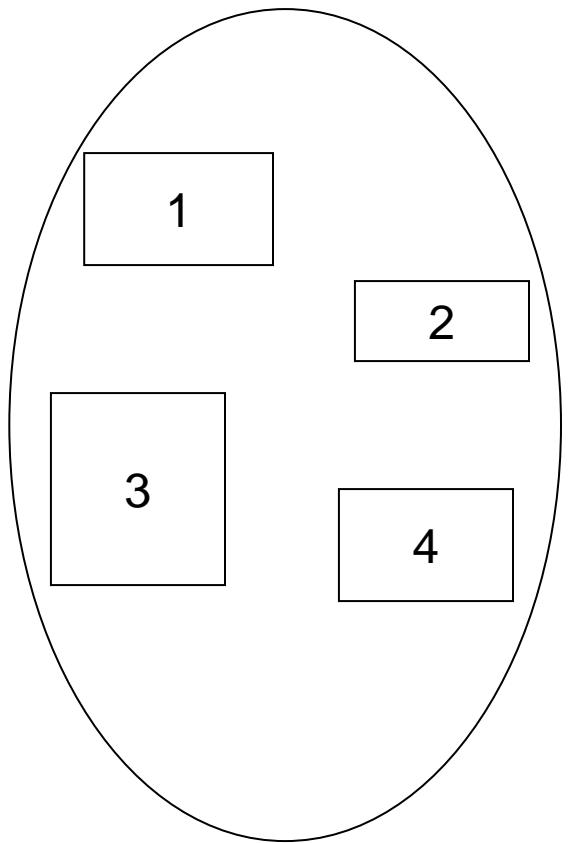


# User's View of a Program

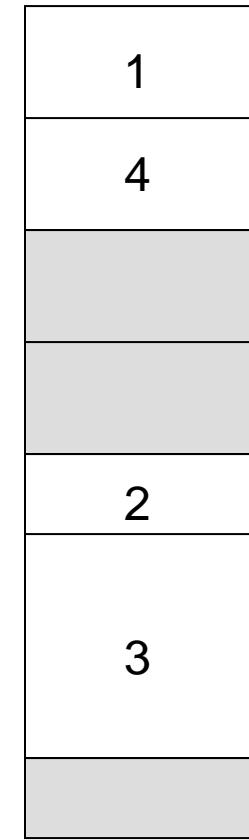




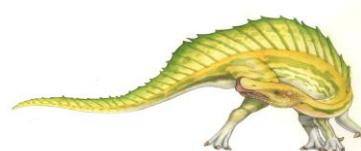
# Logical View of Segmentation



user space



physical memory space





# Segmentation Architecture

- Logical address consists of a two tuple: <segment-number, offset>
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
  - **base** – contains the starting physical address where the segments reside in memory
  - **limit** – specifies the length of the segment
- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program;  
segment number **s** is legal if **s < STLR**





# Segmentation Architecture (Cont.)

---

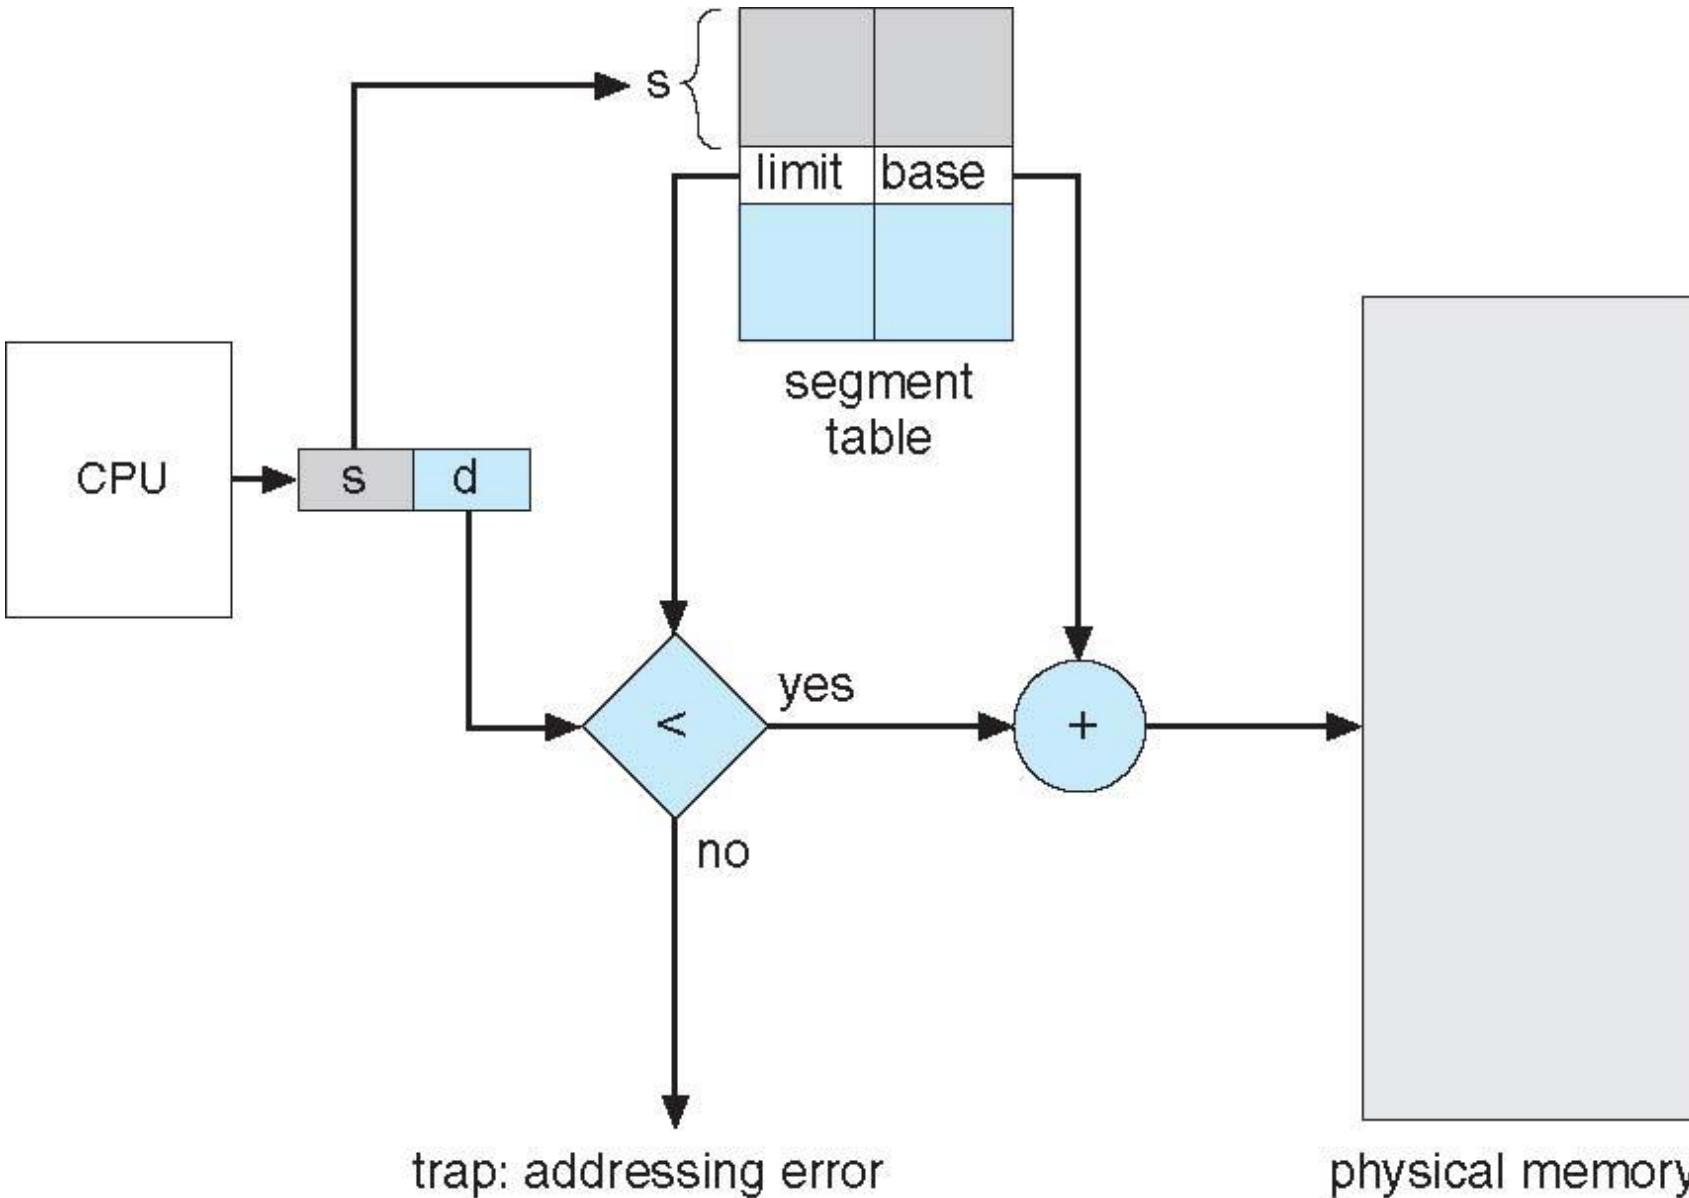
## □ Protection

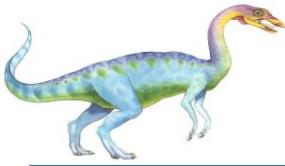
- With each entry in segment table associate:
  - ▶ validation bit = 0  $\Rightarrow$  illegal segment
  - ▶ read/write/execute privileges
- Protection bits associated with segments; code sharing occurs at segment level
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem
- A segmentation example is shown in the diagram (*shown in next slide*)





# Segmentation Hardware





# Paging

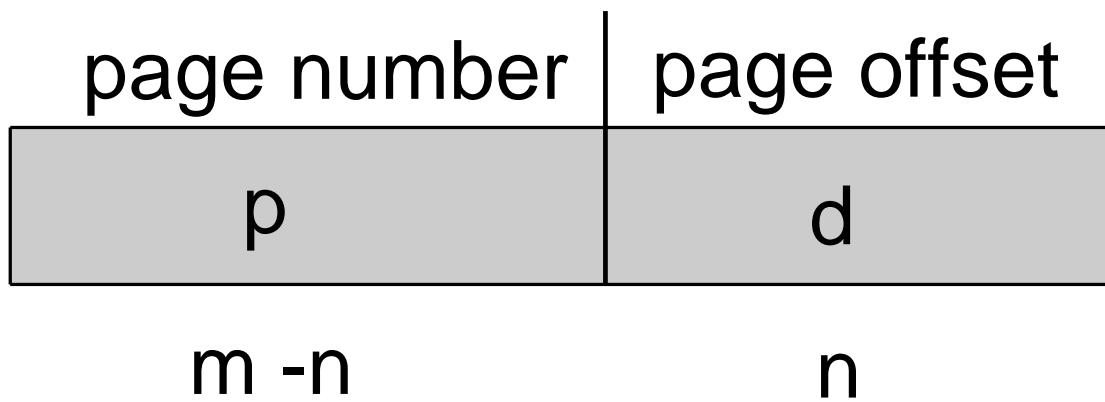
- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
  - Avoids external fragmentation
  - Avoids problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called **frames**
  - Size is power of 2, between 512 bytes and 16 Mbytes
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size  $N$  pages, need to find  $N$  free frames and load program
- Set up a **page table** to translate logical to physical addresses
- Backing store likewise split into pages
- Still have Internal fragmentation





# Address Translation Scheme

- Address generated by CPU is divided into:
    - **Page number** ( $p$ ) – used as an index into a **page table** which contains base address of each page in physical memory
    - **Page offset** ( $d$ ) – combined with base address to define the physical memory address that is sent to the memory unit

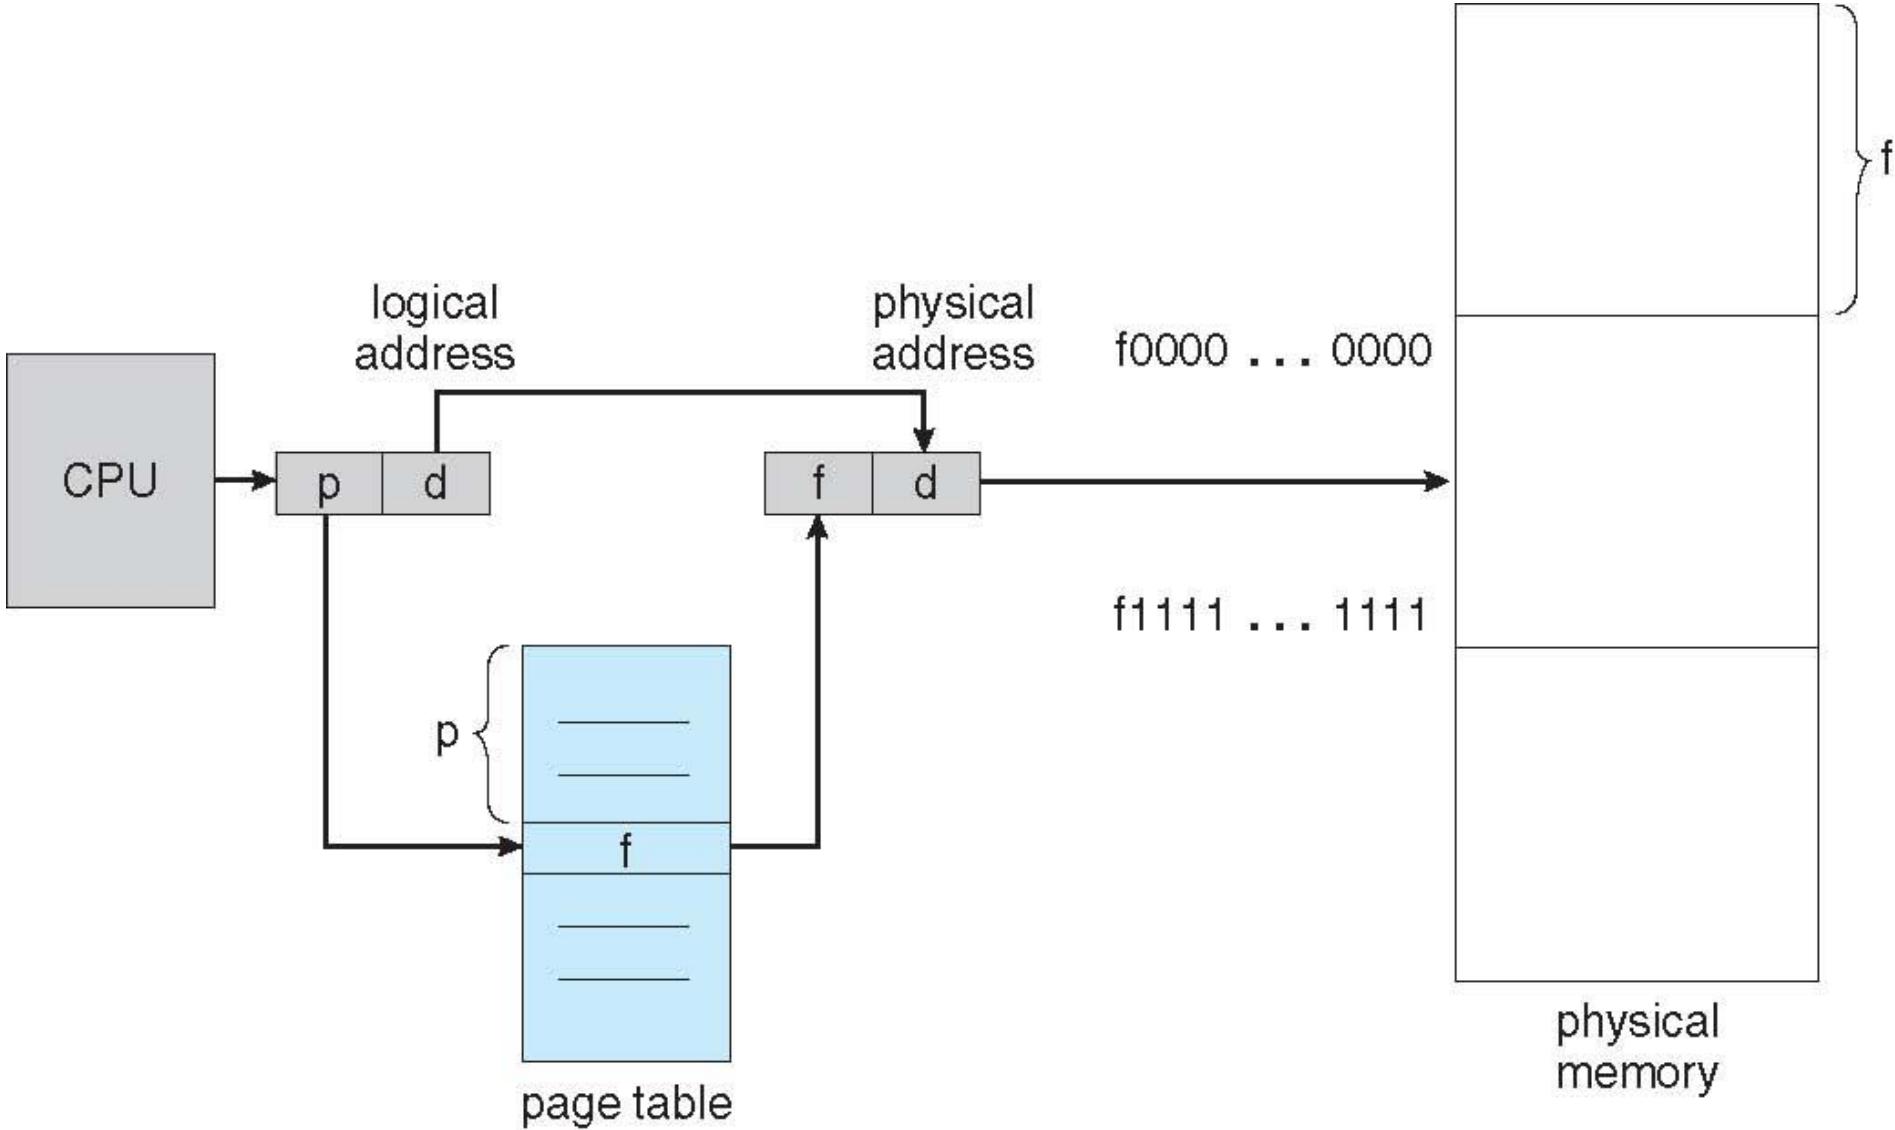


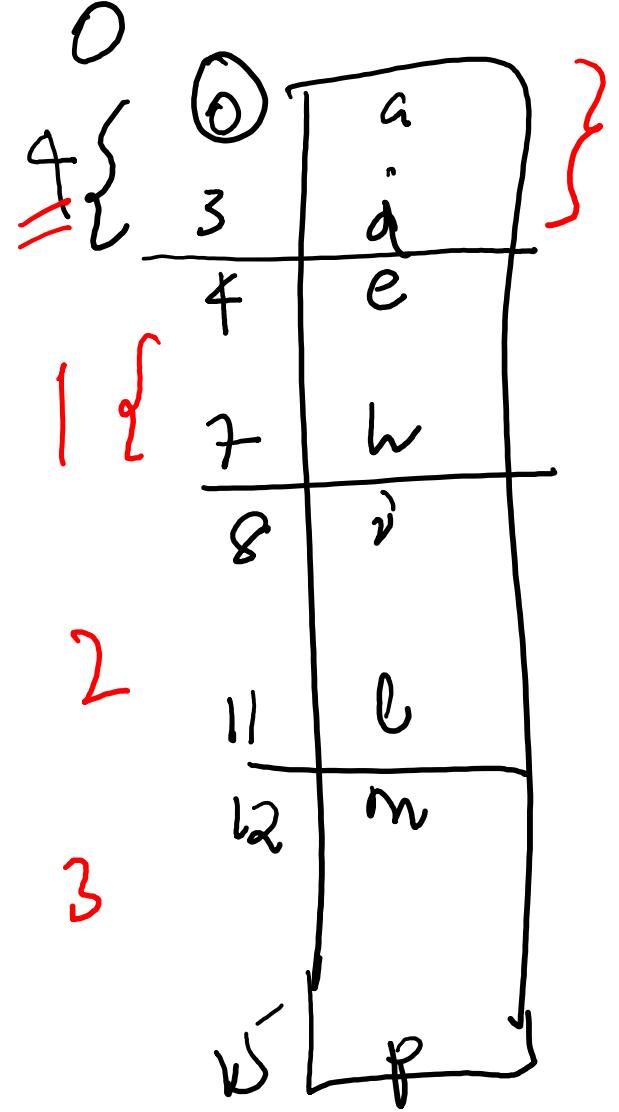
- For given logical address space  $2^m$  and page size  $2^n$





# Paging Hardware



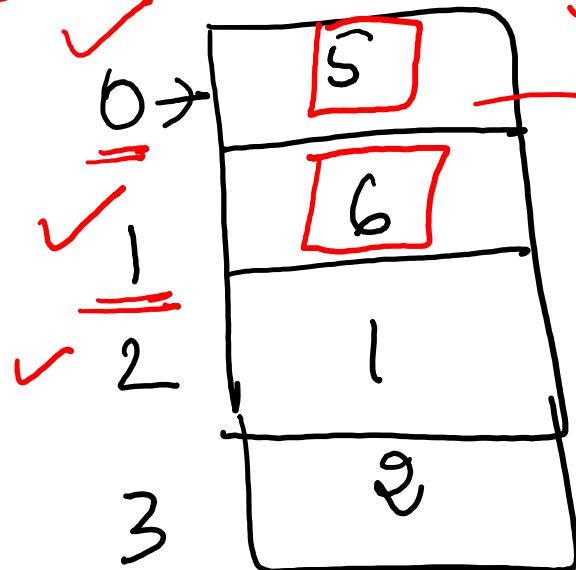


$$\underline{m=4} \quad \underline{n=2}$$

Pagesize = 4 bytes

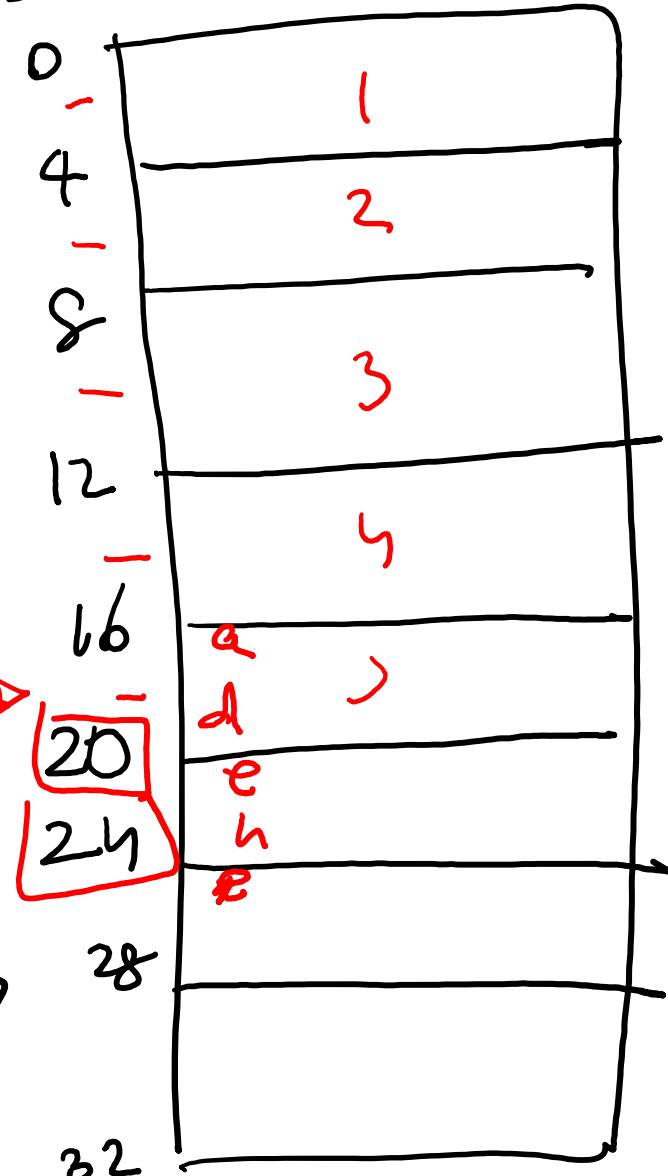
$$20 = (5 \times 4) + 0$$

$$= (5 \times 4) + 1$$



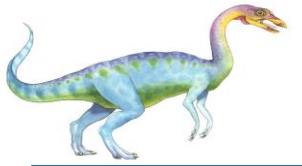
8 pages

Physical memory 32 bytes

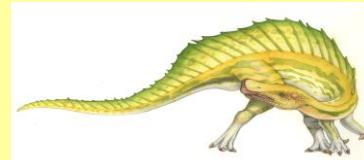




BREAK  
TIME



# RECAP



1. Segmentation is a memory-management scheme that supports programmer view of memory

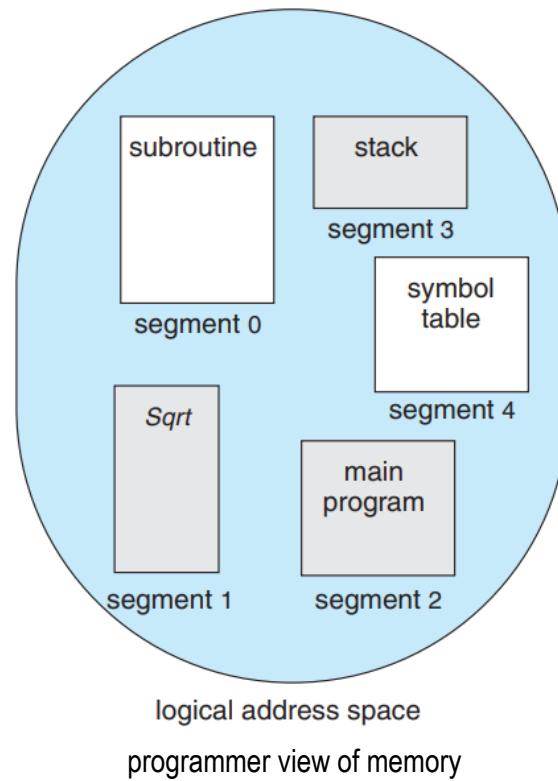
**True**

2. In segmentation, the logical address consists of which two parts?

**Segment-number**  
**offset**

# RECAP QUESTIONS (1)

3.



	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

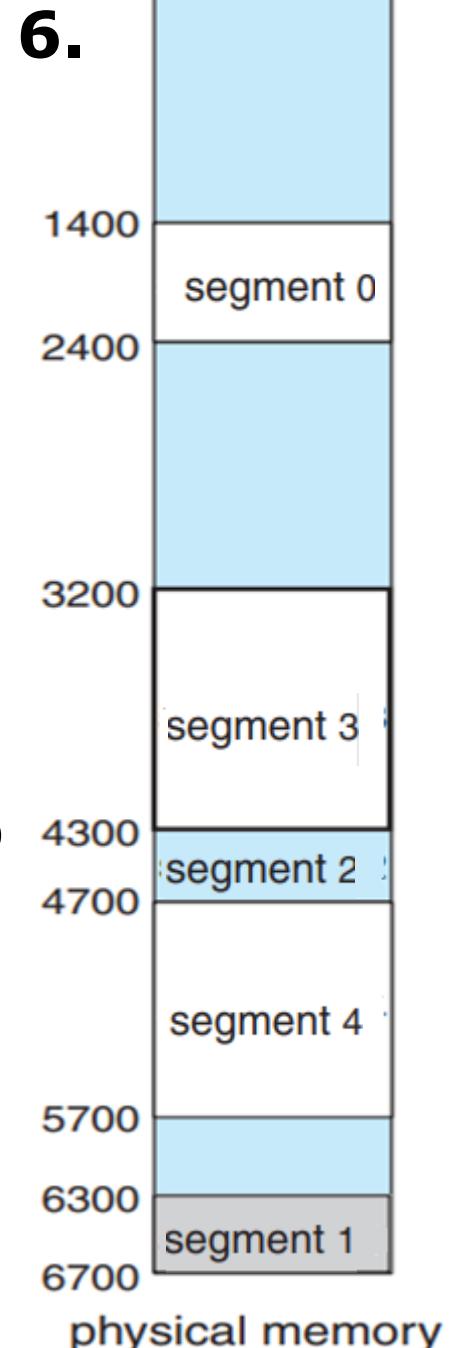
segment table

4. Reference to segment 3, **byte 852**, is mapped to which location?

**4052**

5. Reference to segment 0, **byte 1222**, is mapped to which location?

**Error**





## RECAP QUESTIONS (2)

7. For given logical address space  $2^m$  and page size  $2^n$  and  $m = 4$  and  $n = 2$ , what is the page size and total logical address space?

4 bytes

16 bytes

8. Also, given that physical memory space is 32 bytes. Calculate the number of **associated number pages** for this physical memory?

8 pages





## RECAP QUESTIONS (2)

9. Paging is a memory management scheme that permits physical memory to be non-contiguous? True
10. Paging differs from segmentation as it avoids external fragmentation and need for compaction? True
11. Paging involves breaking **logical memory** into blocks of the same size called Frames? **Pages** False
12. In paging, the logical address consists of which two parts?

**Page number**

**Page offset**

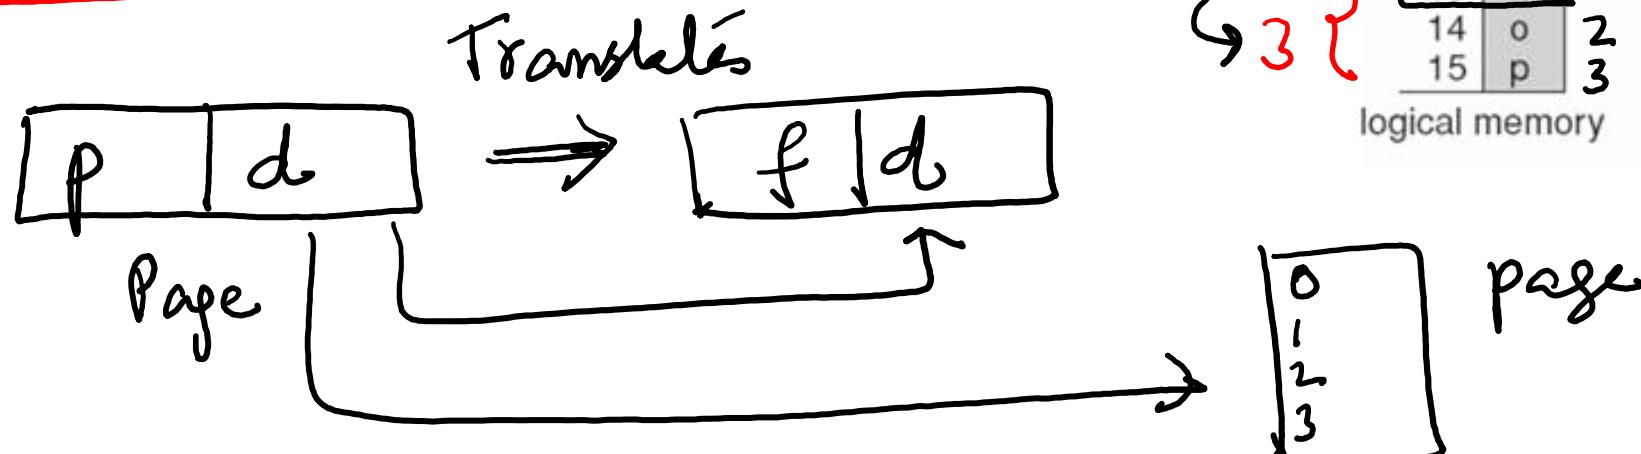


**Q.13.** How can the following logical memory space be mapped into the given physical memory space?

Where does logical address 13 map into physical address ?

9

Page Size = Frame Size



page no.

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

1 {  
2 {  
3 {

0 1 2 3

logical memory

frame no.  
↓

0	5
1	6
2	1
3	2

page table

0	
4	
8	1 2 3
12	
16	
20	a b c d
24	e f g h
28	

physical memory

**Q.13.** How can the following logical memory space be mapped into the given physical memory space?

PN = 3

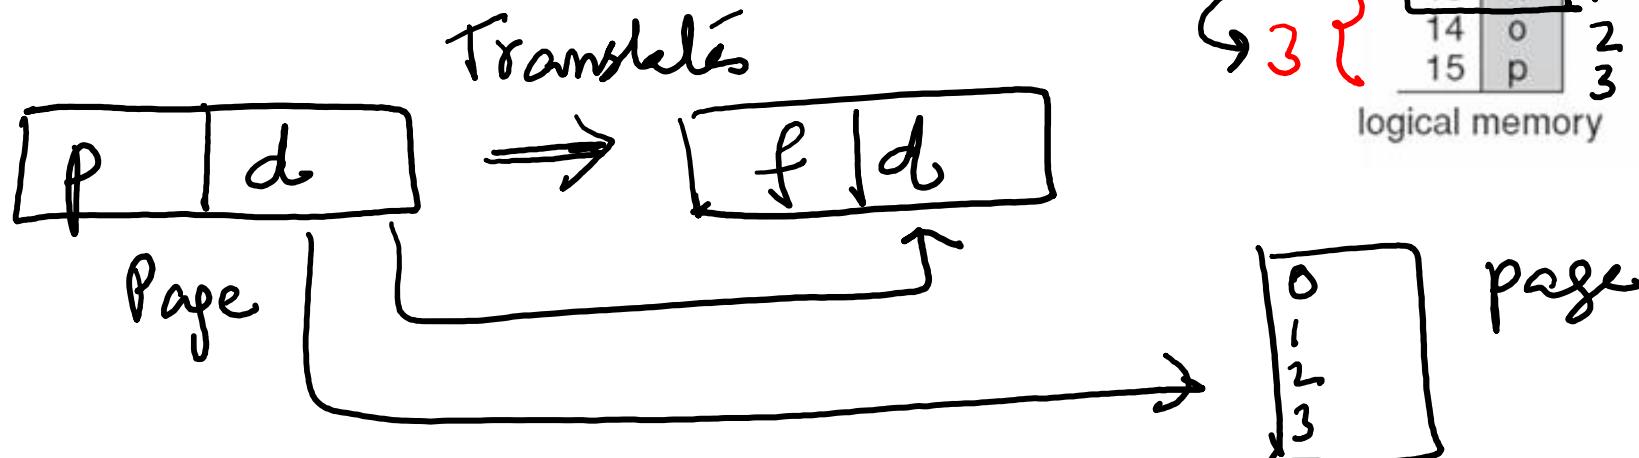
FN = 2

offset = 1

Where does logical address 13 map into physical address ?

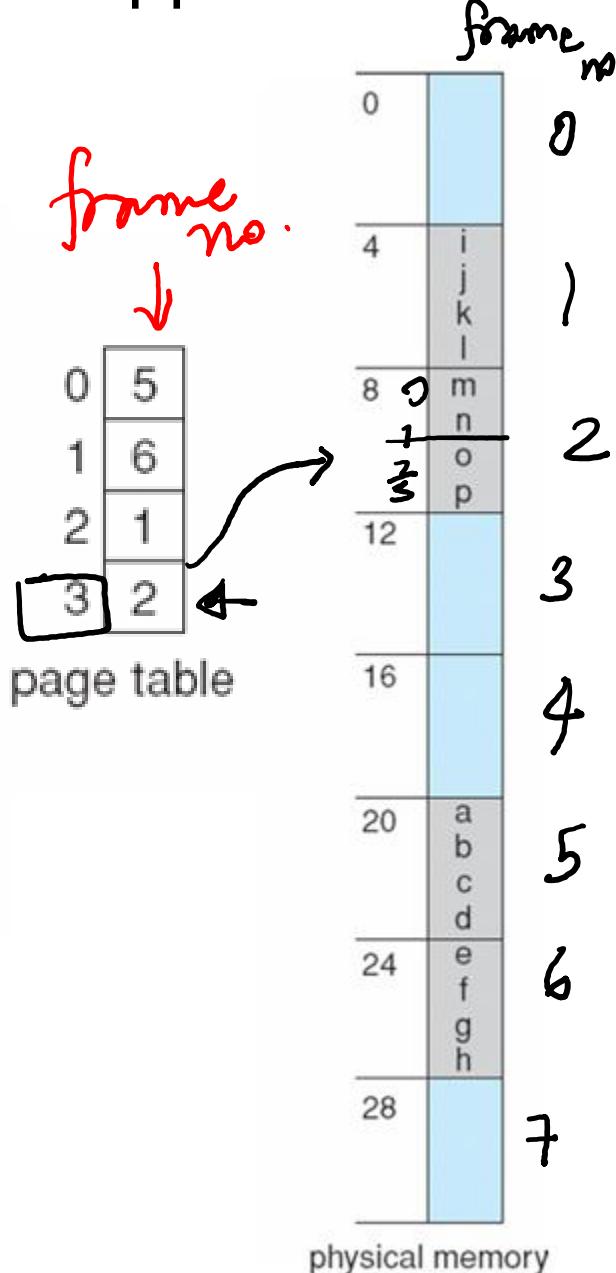
9

$$\begin{aligned}\text{Physical Address} &= \text{FN} * \text{framesize} + \\ &= 2 * 4 + 1 \quad \text{offset} \\ &= 9\end{aligned}$$



page no.

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p



**Q.14.** We have a 32-bit machine with 16GB RAM using 4Kbyte page size. How many memory frames are there in the machine?

We have,

32-bit machine  $\Rightarrow$  each address  $\Rightarrow$  32 bits

Frame size = page size

$$= 4 \text{ KByte} = 2^{10} \times 4 \quad \text{Why? } 1 \text{ KByte} = ?$$

$$\boxed{1 \text{ GB} = 2^{30}}$$

$$= 2^{10} \times 2^2$$

$$= 2^{12}$$

$$\text{Total RAM} = 16 \times 1 \text{ GB}$$

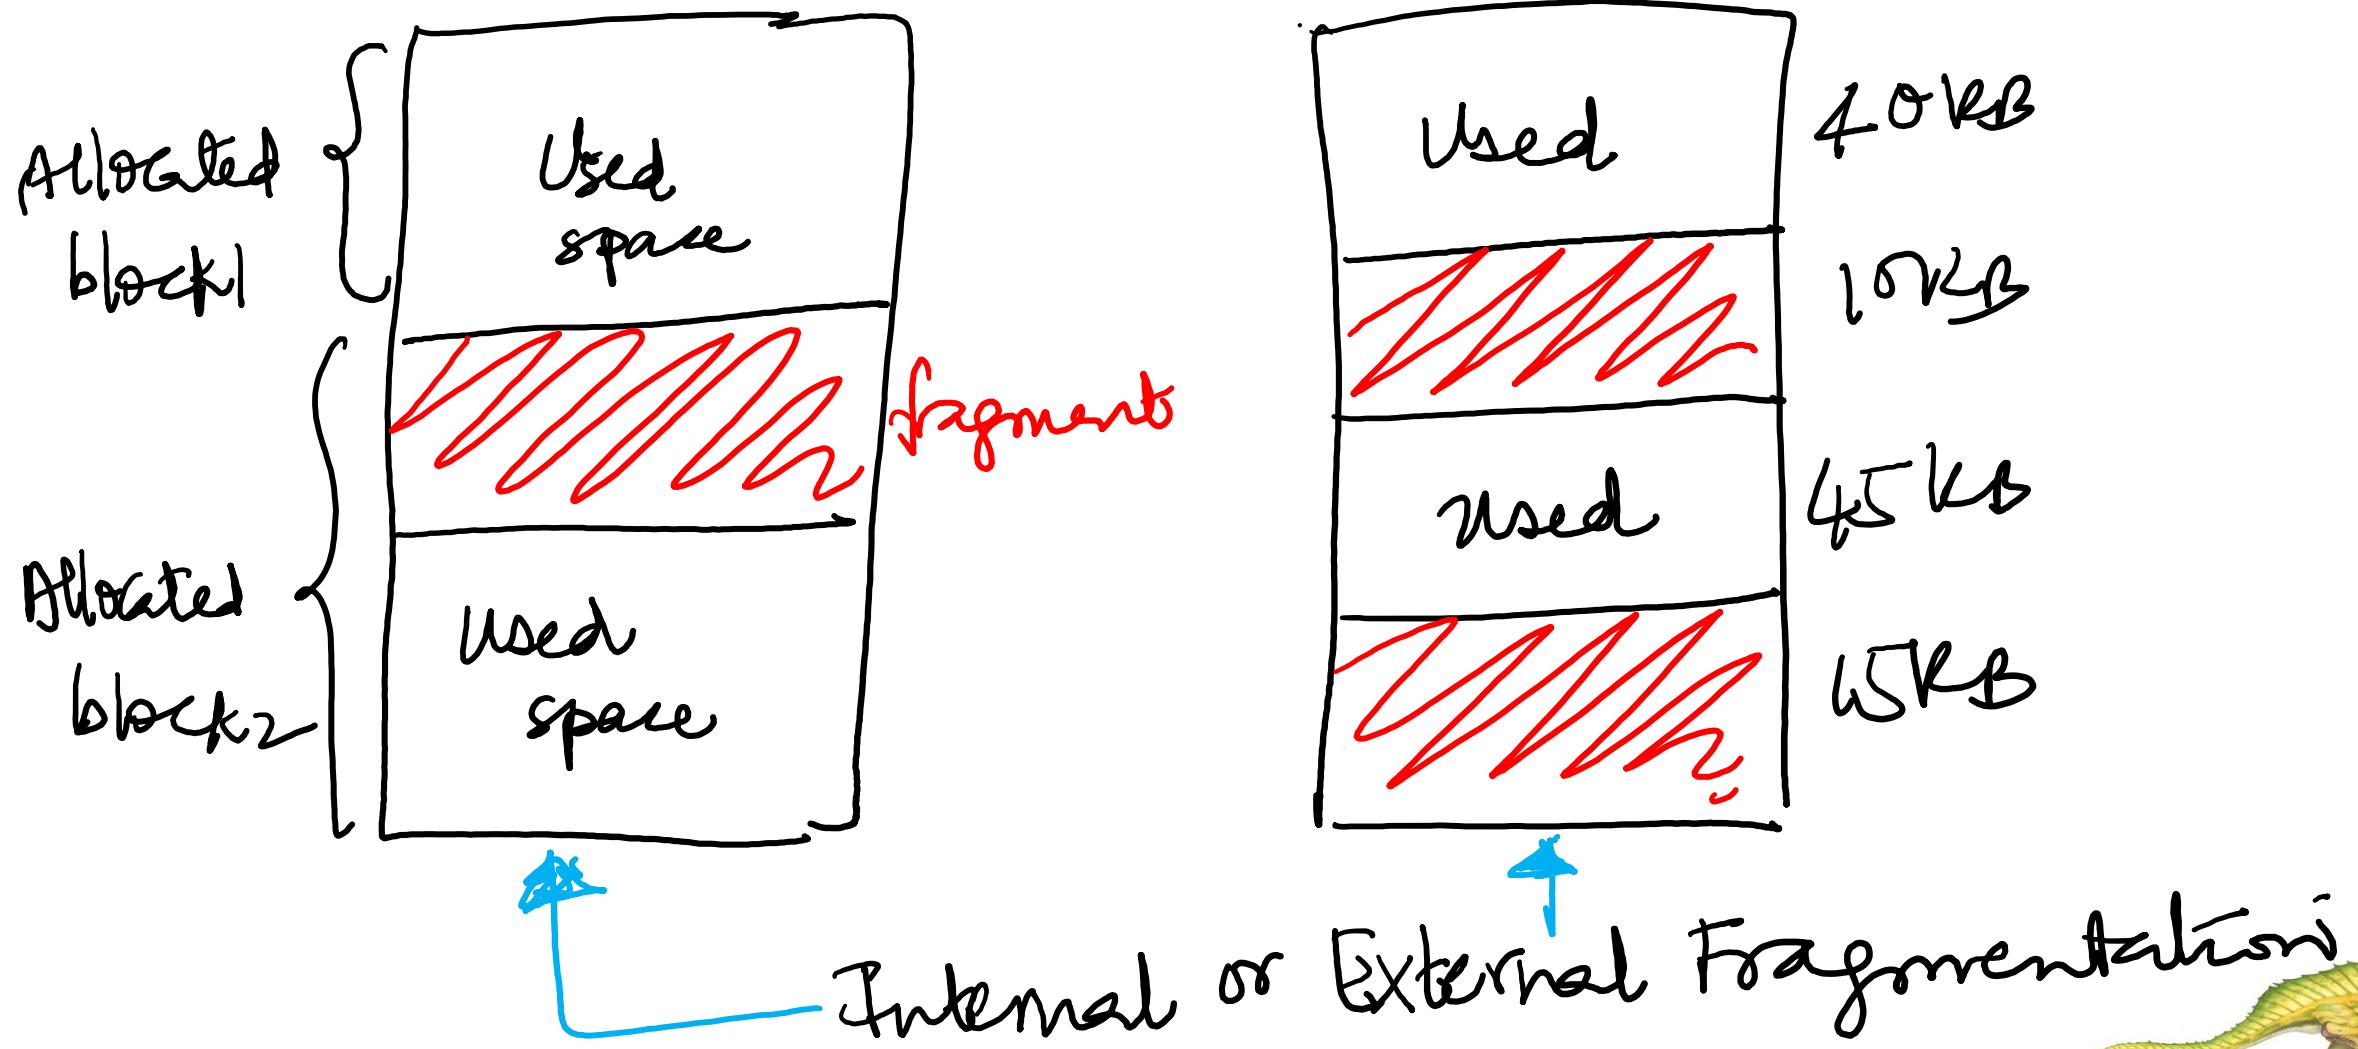
$$= 2^4 \times 2^{30}$$

$$= 2^{34}$$

$$\therefore \text{Number of frames} = \frac{\text{Total RAM}}{\text{Frame Size}}$$
$$= \frac{2^{34}}{2^{12}} = 2^{22} \text{ frames}$$



# One last RECAP!





## Paging (Cont.)

- Takes care of external fragmentation but not internal fragmentation.
- How do we calculate internal fragmentation?
- Suppose
  - Page size = 2,048 bytes
  - Process size = 72,766 bytes
- How many pages would be needed?
  - 35 pages + 1,086 bytes
- Therefore, internal fragmentation would be calculated as:
  - $2,048 - 1,086 = 962$  bytes
- Worst case fragmentation = 1 frame – 1 byte
  - A process would need n pages + 1 byte.
  - *It would be allocated n + 1 frames, resulting in internal fragmentation of almost an entire frame.*





## Paging (Cont.)

---

- On average fragmentation =  $1 / 2$  frame size
- So small frame sizes desirable?
  - Each page table entry takes memory to track
  - Page sizes growing over time
- Today, pages typically are between 4 KB and 8 KB in size, and some systems support even larger page sizes.
- Some CPUs and kernels even support multiple page sizes.
  - Solaris uses page sizes of 8 KB and 4 MB, depending on the data stored by the pages.





# Free Frames

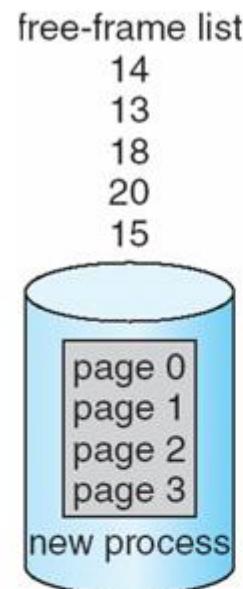
- A system with **32-bit page-table** entries may address less physical memory than the possible maximum.
- A **32-bit CPU** uses **32-bit addresses**, meaning that a given process space can only be **2<sup>32</sup> bytes (4 TB)**.
- Therefore, ***paging lets us use physical memory*** that is larger than what can be addressed by the CPU's address pointer length





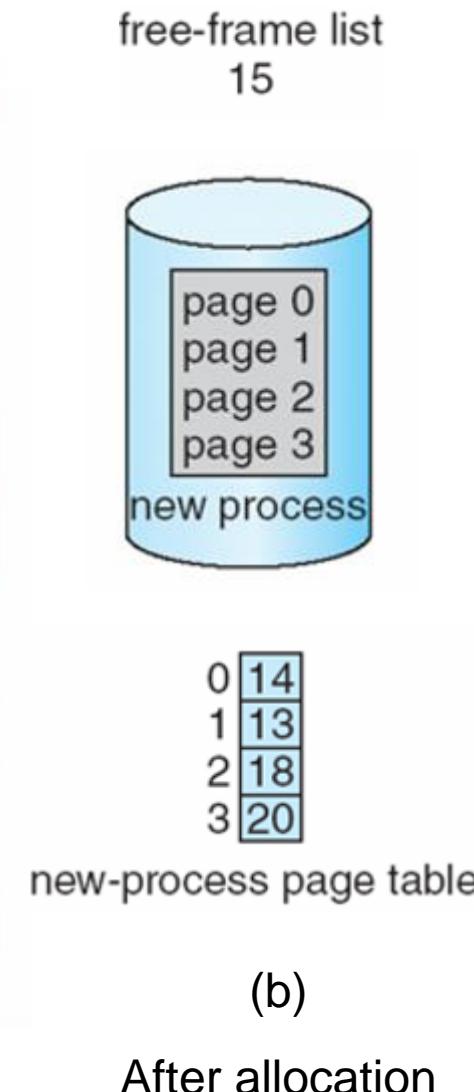
# Free Frames

- Process arrives in system to be executed, **size in page** is examined.
- Each page of the process needs one frame.
- **Process of n pages needs atleast n frames and must be available in memory.**
- If **n frames** are available, allocated to arriving process.
- First page of the process is loaded into one of the allocated frames, and the frame number is put in the page table for this process.
- The next page is loaded into another frame, its frame number is put into the page table, and so on



(a)

Before allocation





# Implementation of Page Table

---

- Page table is kept in main memory
- **Page-table base register (PTBR)** points to the **page table**
- **Page-table length register (PTLR)** indicates **size of the page table**
- In this scheme **every data/instruction access** requires **two memory accesses**
  - One for the page table and one for the data / instruction
  - The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**





## Implementation of Page Table (Cont.)

- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
  - Otherwise need to flush at every context switch
- **TLBs typically small (64 to 1,024 entries)**
- On a TLB miss, value is loaded into the TLB for faster access next time
  - Replacement policies must be considered
  - Some entries can be **wired down** for permanent fast access





# Associative Memory

- Associative memory – parallel search

Page #	Frame #

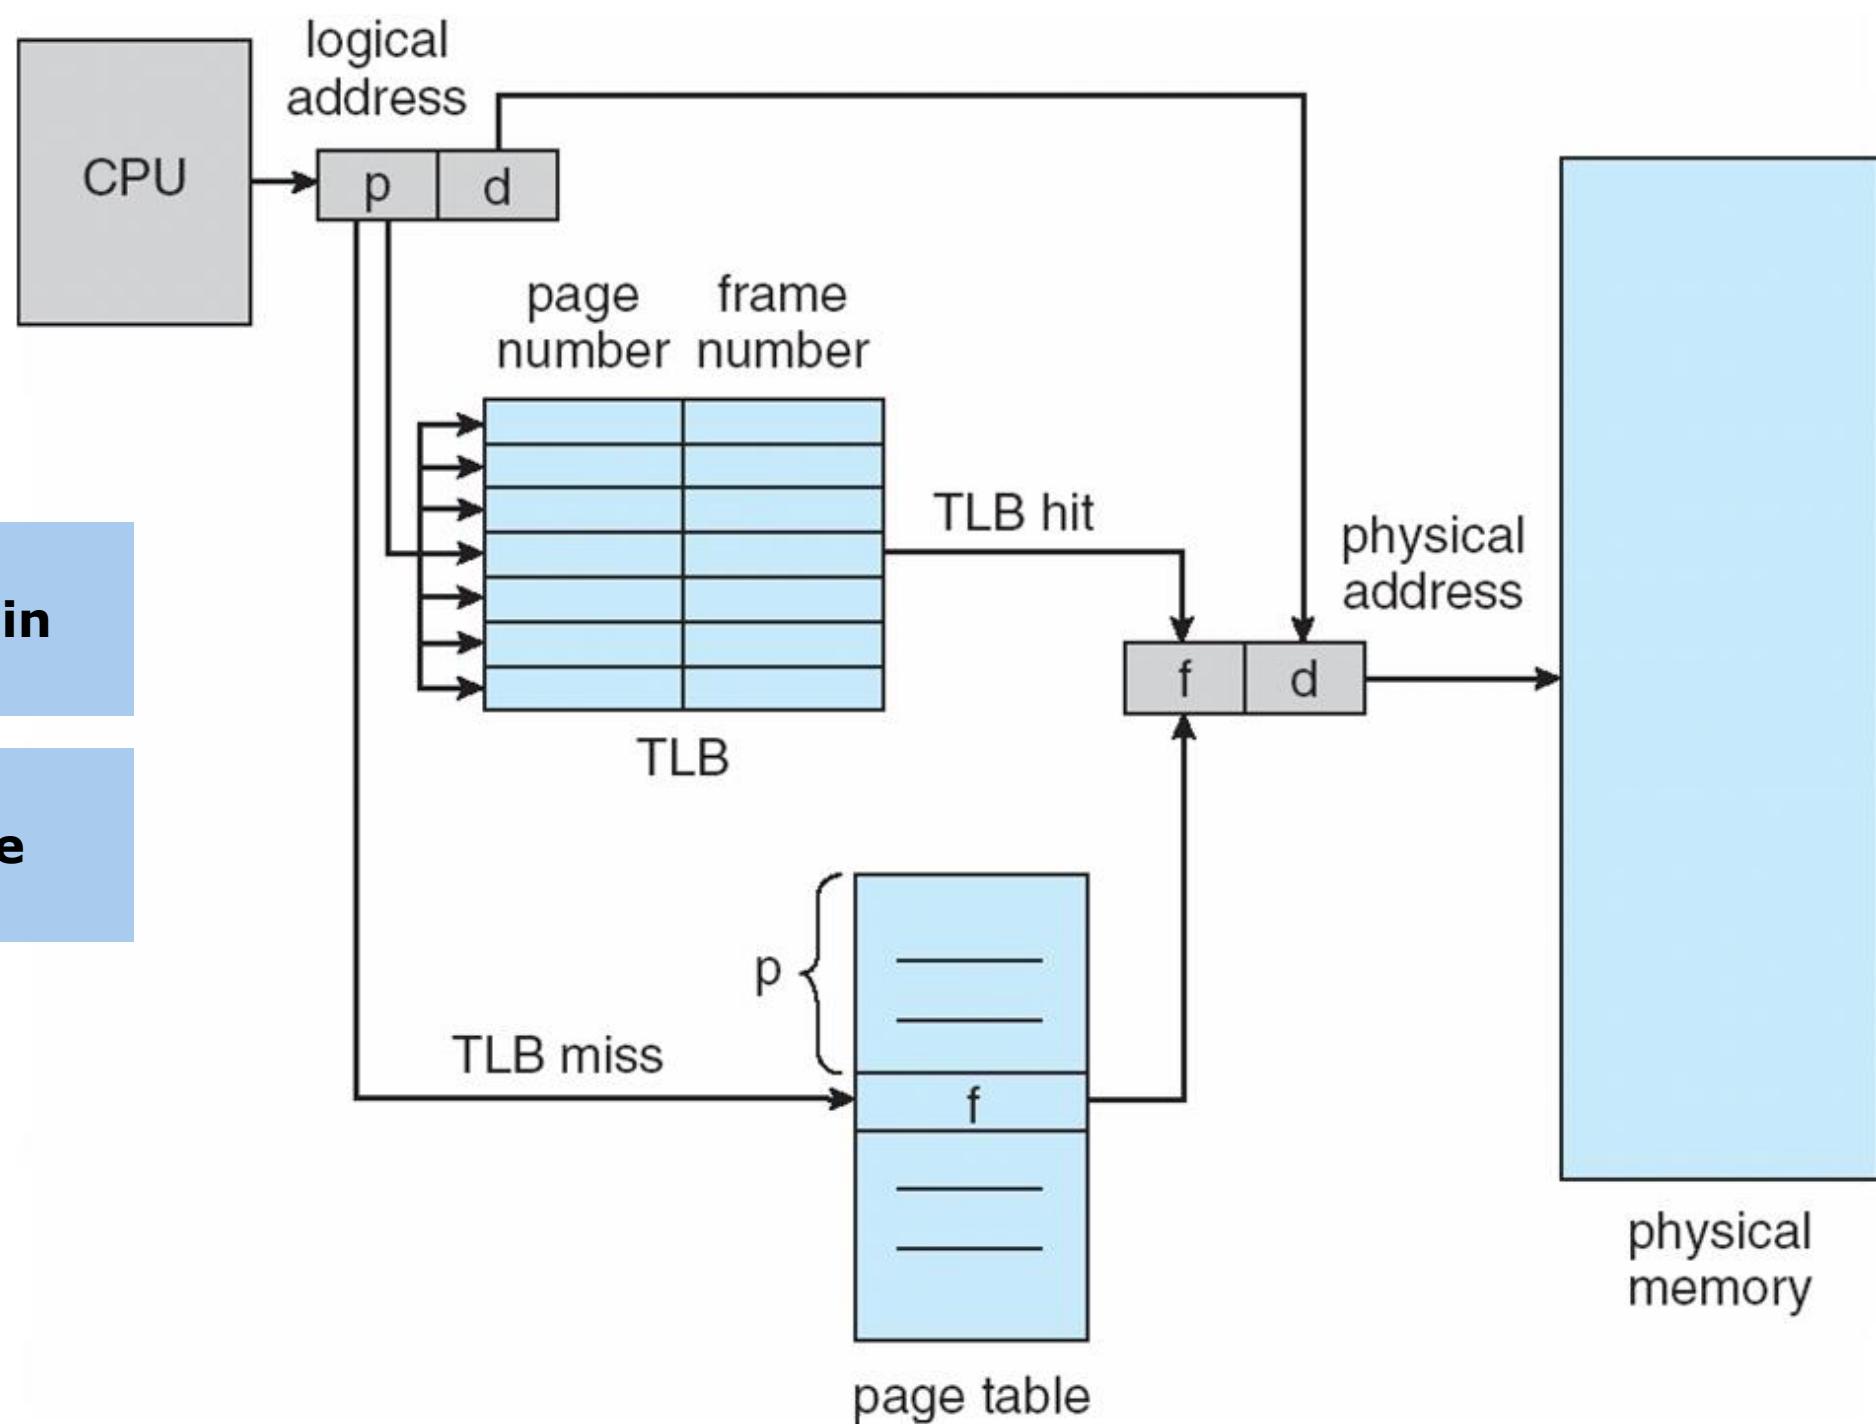
- Address translation (p, d)
  - If p is in associative register, get frame # out
  - Otherwise get frame # from page table in memory



# Paging Hardware With TLB

**p** is compared with  
every page number in  
TLB

If **p** found,  
corresponding frame  
number fetched





# Effective Access Time

- Associative Lookup =  $\varepsilon$  time unit
  - Can be < 10% of memory access time
- Hit ratio =  $\alpha$ 
  - Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- Consider  $\alpha = 80\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
- **Effective Access Time (EAT)**
$$\begin{aligned} \text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$
- Consider  $\alpha = 80\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access
  - $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- Consider more realistic hit ratio ->  $\alpha = 99\%$ ,  $\varepsilon = 20\text{ns}$  for TLB search,  $100\text{ns}$  for memory access     $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$



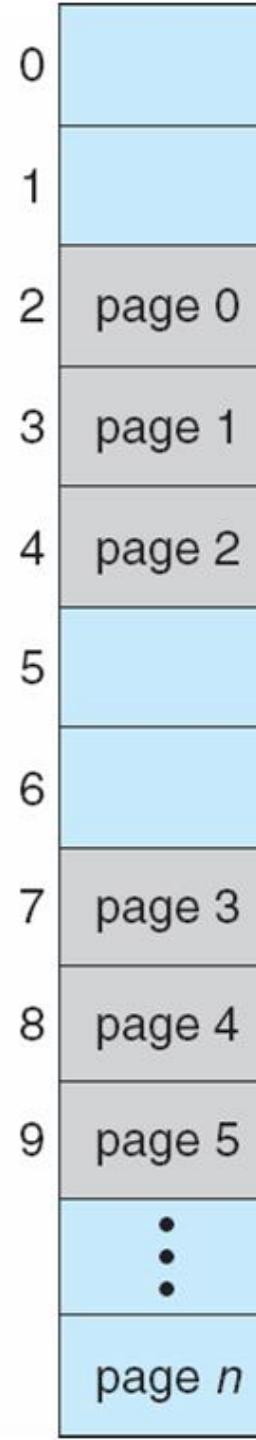
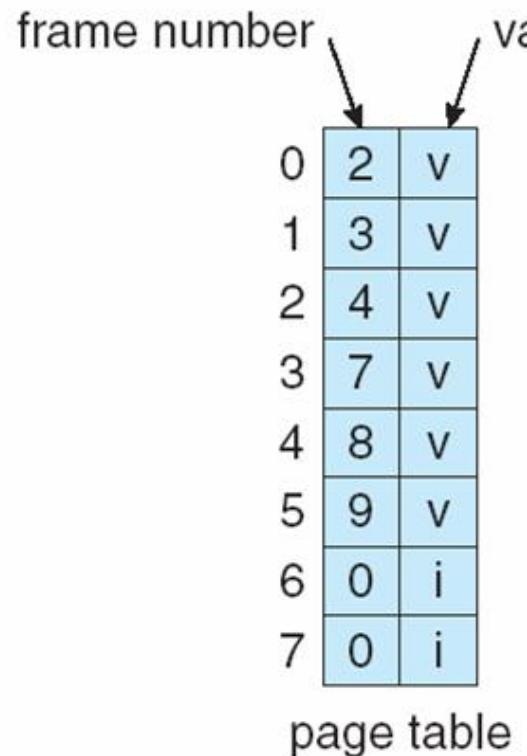
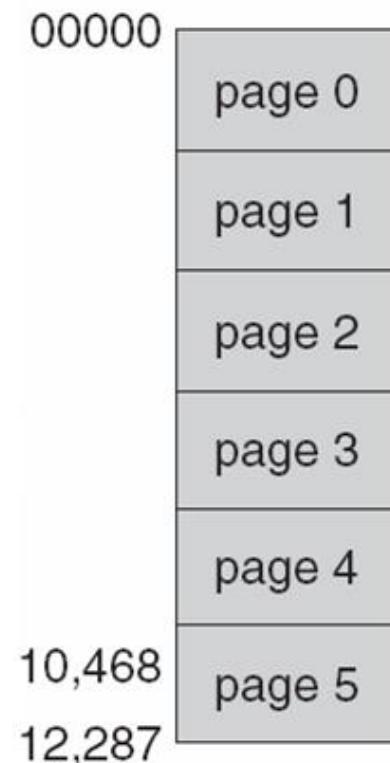


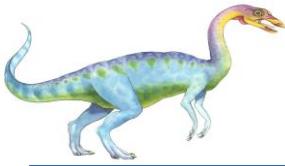
# Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
  - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
  - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
  - “invalid” indicates that the page is not in the process’ logical address space
  - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel



# Valid (v) or Invalid (i) Bit In A Page Table





# Shared Pages

## □ Shared code

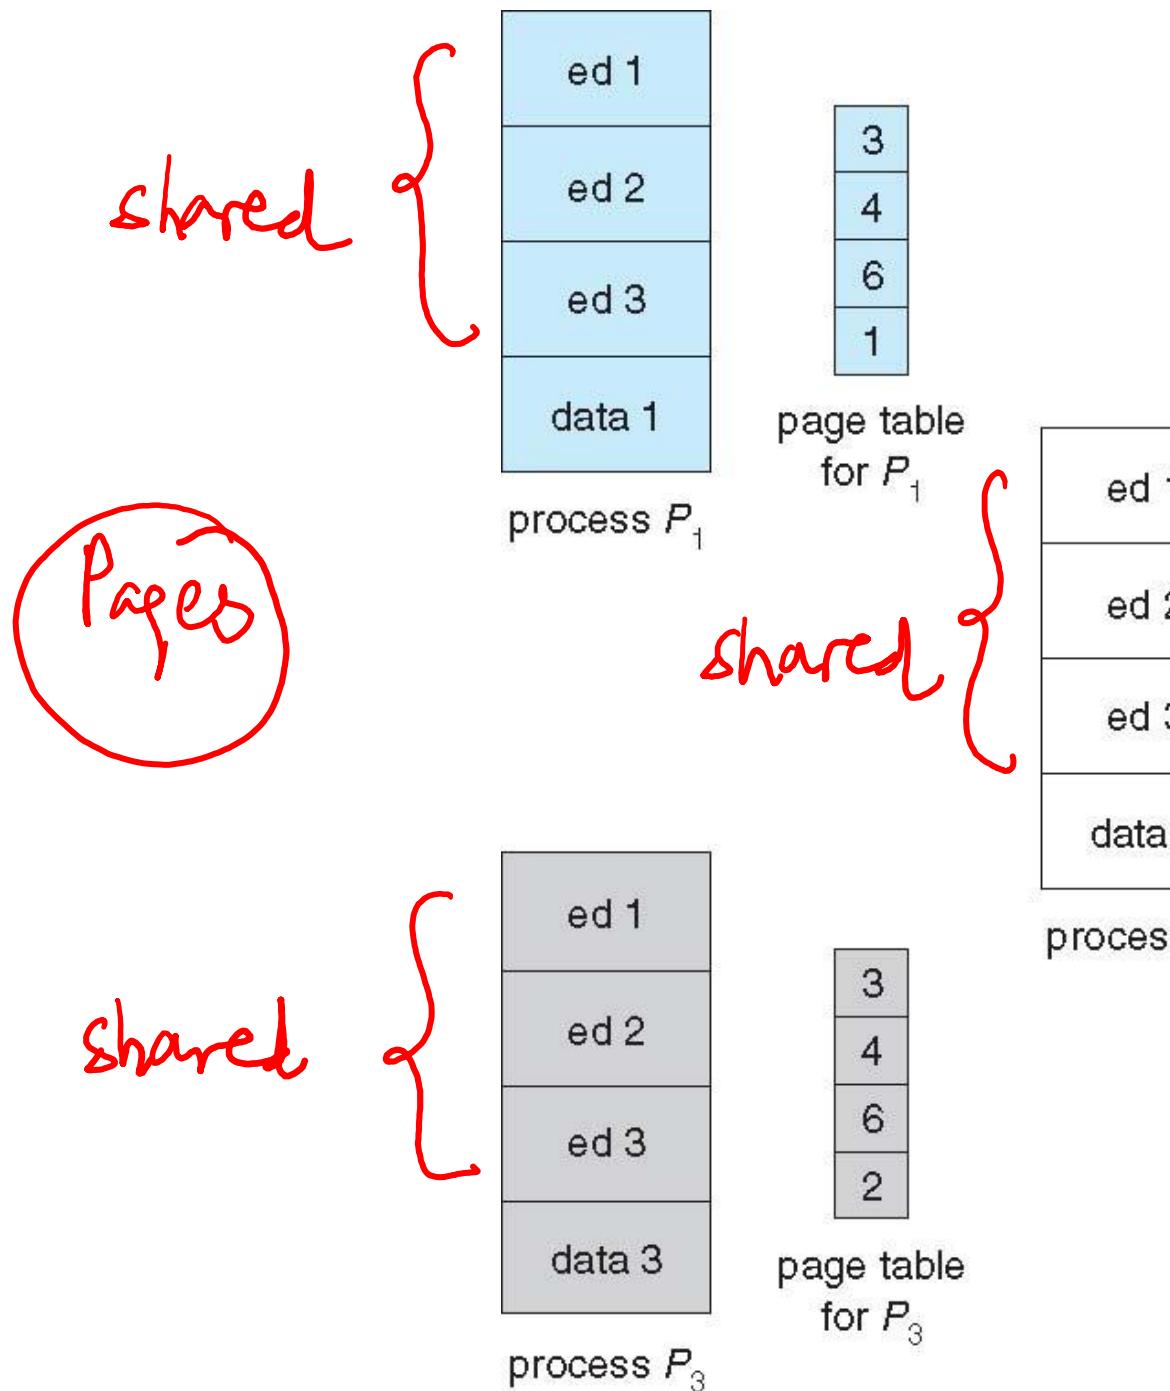
- One copy of read-only (**reentrant**) code shared among processes (i.e., **text editors**, compilers, window systems)
- Similar to multiple threads sharing the same process space
- Also useful for interprocess communication if sharing of read-write pages is allowed

## □ Private code and data

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space



# Shared Pages Example



0	
1	data 1
2	data 3
3	ed 1
4	ed 2
5	
6	ed 3
7	data 2
8	
9	
10	
11	



# Structure of the Page Table

---

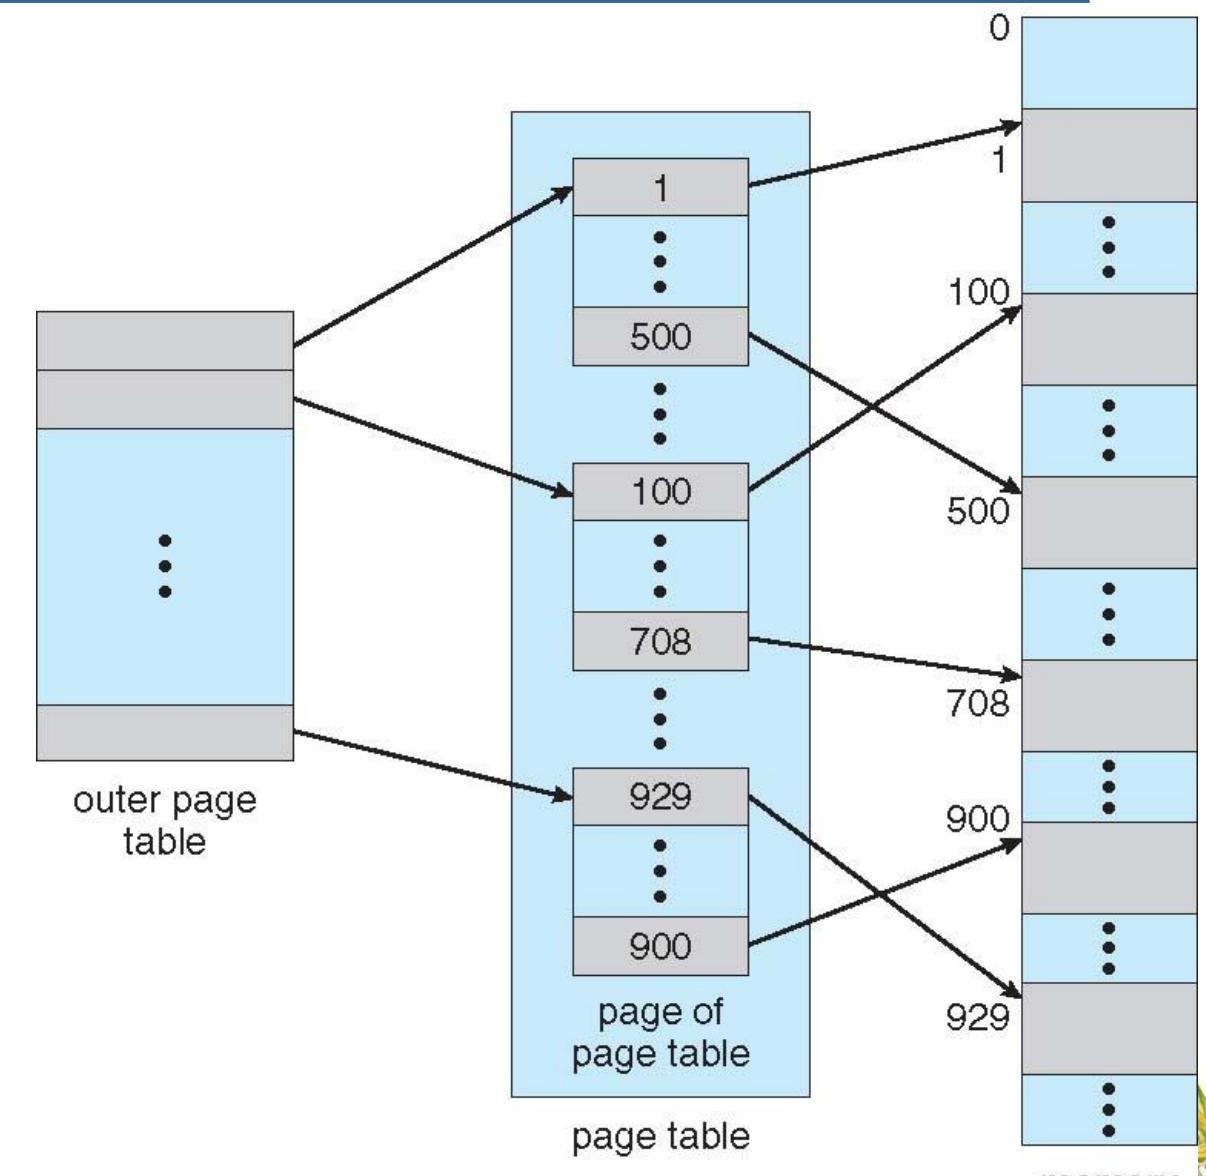
- Memory structures for paging can get huge using straight-forward methods
  - Consider a **32-bit logical address space** as on modern computers
  - Page size of **4 KB ( $2^{12}$ )**
  - Page table would have **1 million entries ( $2^{32} / 2^{12}$ )**
  - If **each entry** is 4 bytes, each process needs **4MB physical address space for page table alone**
    - ▶ Cannot allocate that contiguously in main memory
- **Solution:** Divide page table into smaller pieces
  - Hierarchical Paging
  - Hashed Page Tables
  - Inverted Page Tables





# Hierarchical Page Tables

- Break up the **logical address space** into *multiple page tables*
- A simple technique is a **two-level page table**
- We then **page** the **page table**



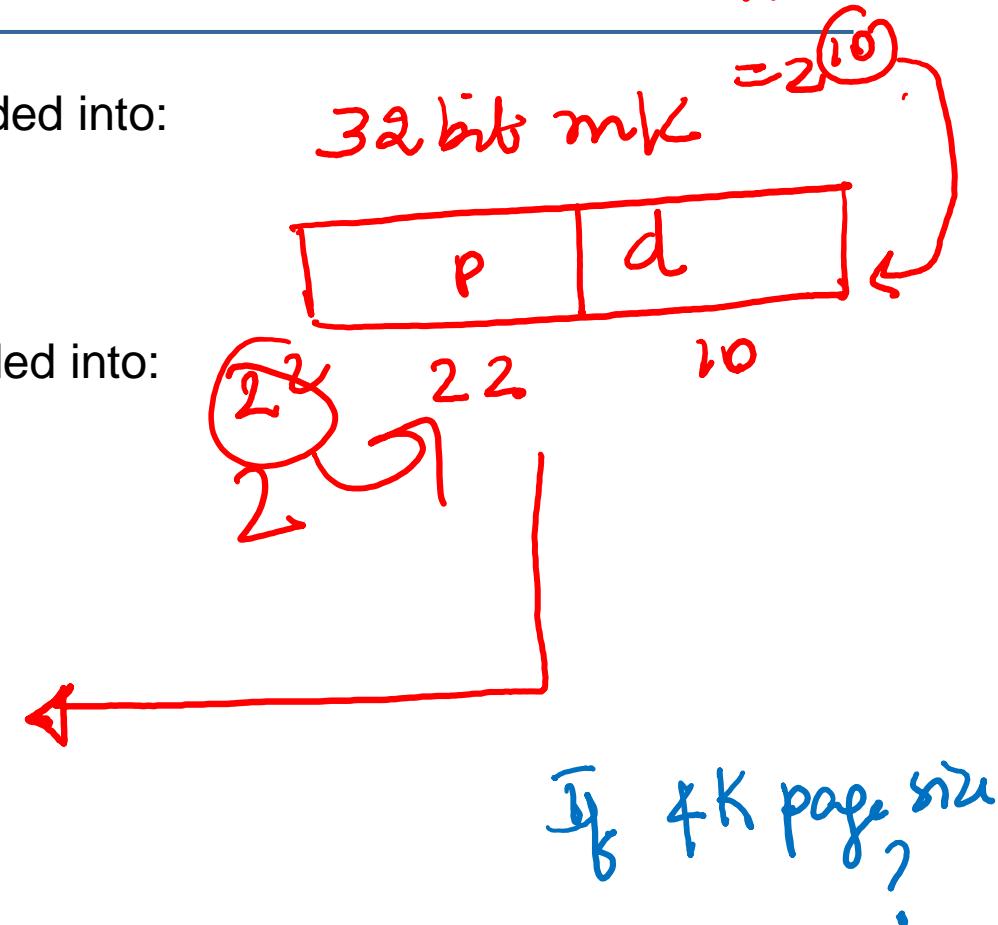


# Two-Level Paging Example

Page size  
1KB  
 $= 2^{10}$

- A logical address (on 32-bit machine with 1K page size) is divided into:
  - a page number consisting of 22 bits
  - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
  - a 12-bit page number
  - a 10-bit page offset
- Thus, a logical address is as follows:

page number	page offset
$p_1$	$p_2$
{ 12 22	10

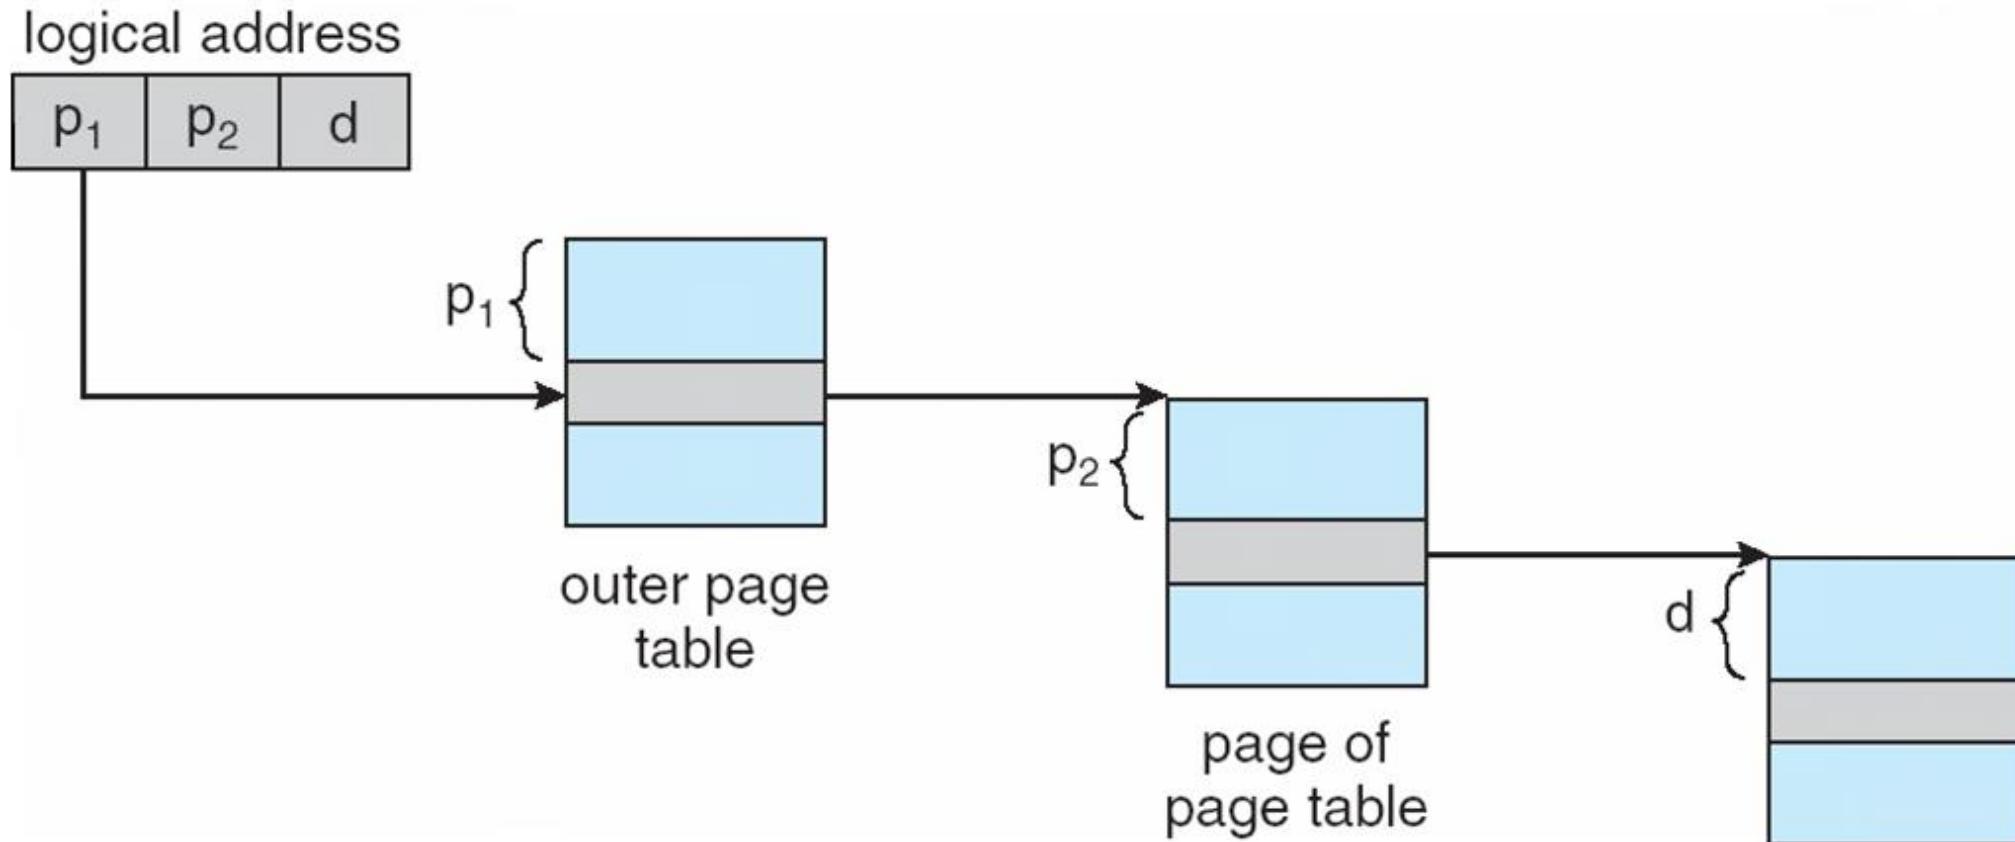


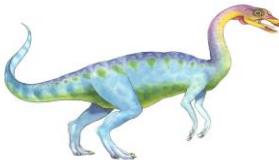
- where  $p_1$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the inner page table
- Known as **forward-mapped page table**

$d = 12$  How?



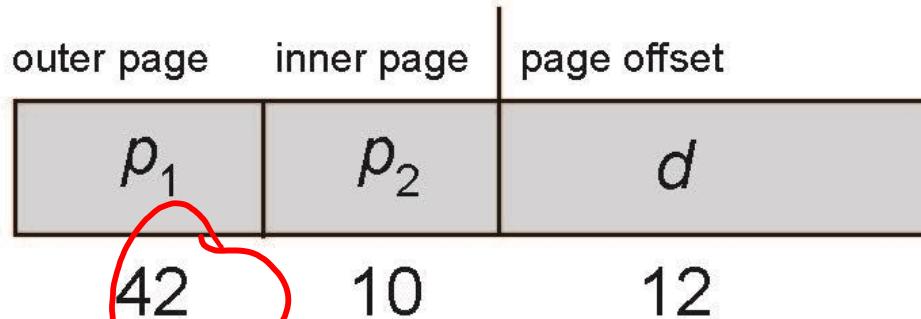
# Address-Translation Scheme





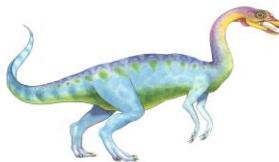
# 64-bit Logical Address Space

- Even two-level paging scheme
- If page size is 4 KB ( $2^{12}$ )
  - Then page table has  $2^{52}$  entries
  - If two level scheme, inner page tables could be  $2^{10}$  4-byte entries
  - Address would look like

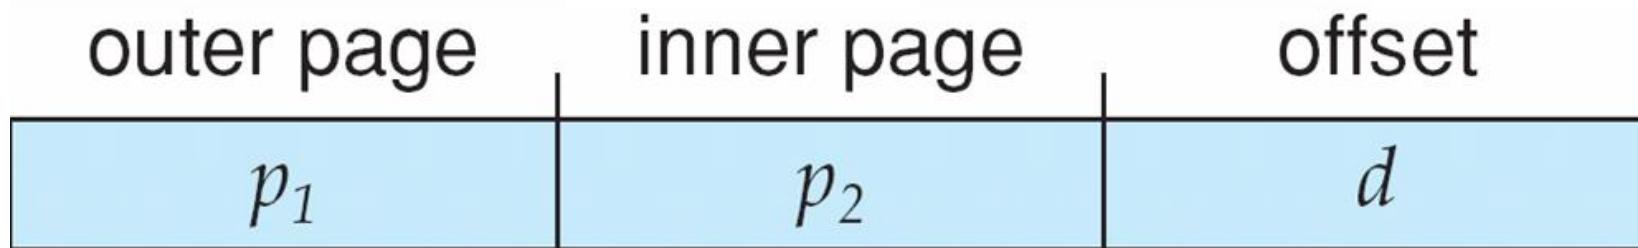
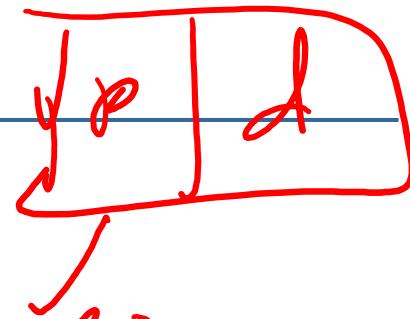


- Outer page table has  $2^{42}$  entries or  $2^{44}$  bytes
- One solution is to add a 2<sup>nd</sup> outer page table
- But in the following example the 2<sup>nd</sup> outer page table is still  $2^{34}$  bytes in size
  - ▶ And possibly 4 memory access to get to one physical memory location





# Three-level Paging Scheme

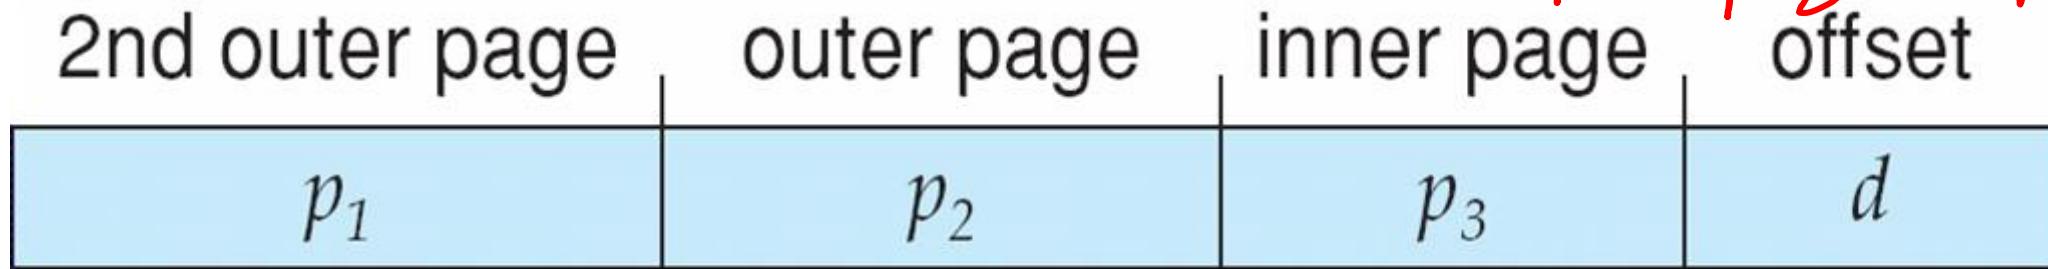


42

2nd outer page

10

12

 $p_1 \ p_2$  $p_1 \ p_2 \ p_3$ 

32

10

10

12





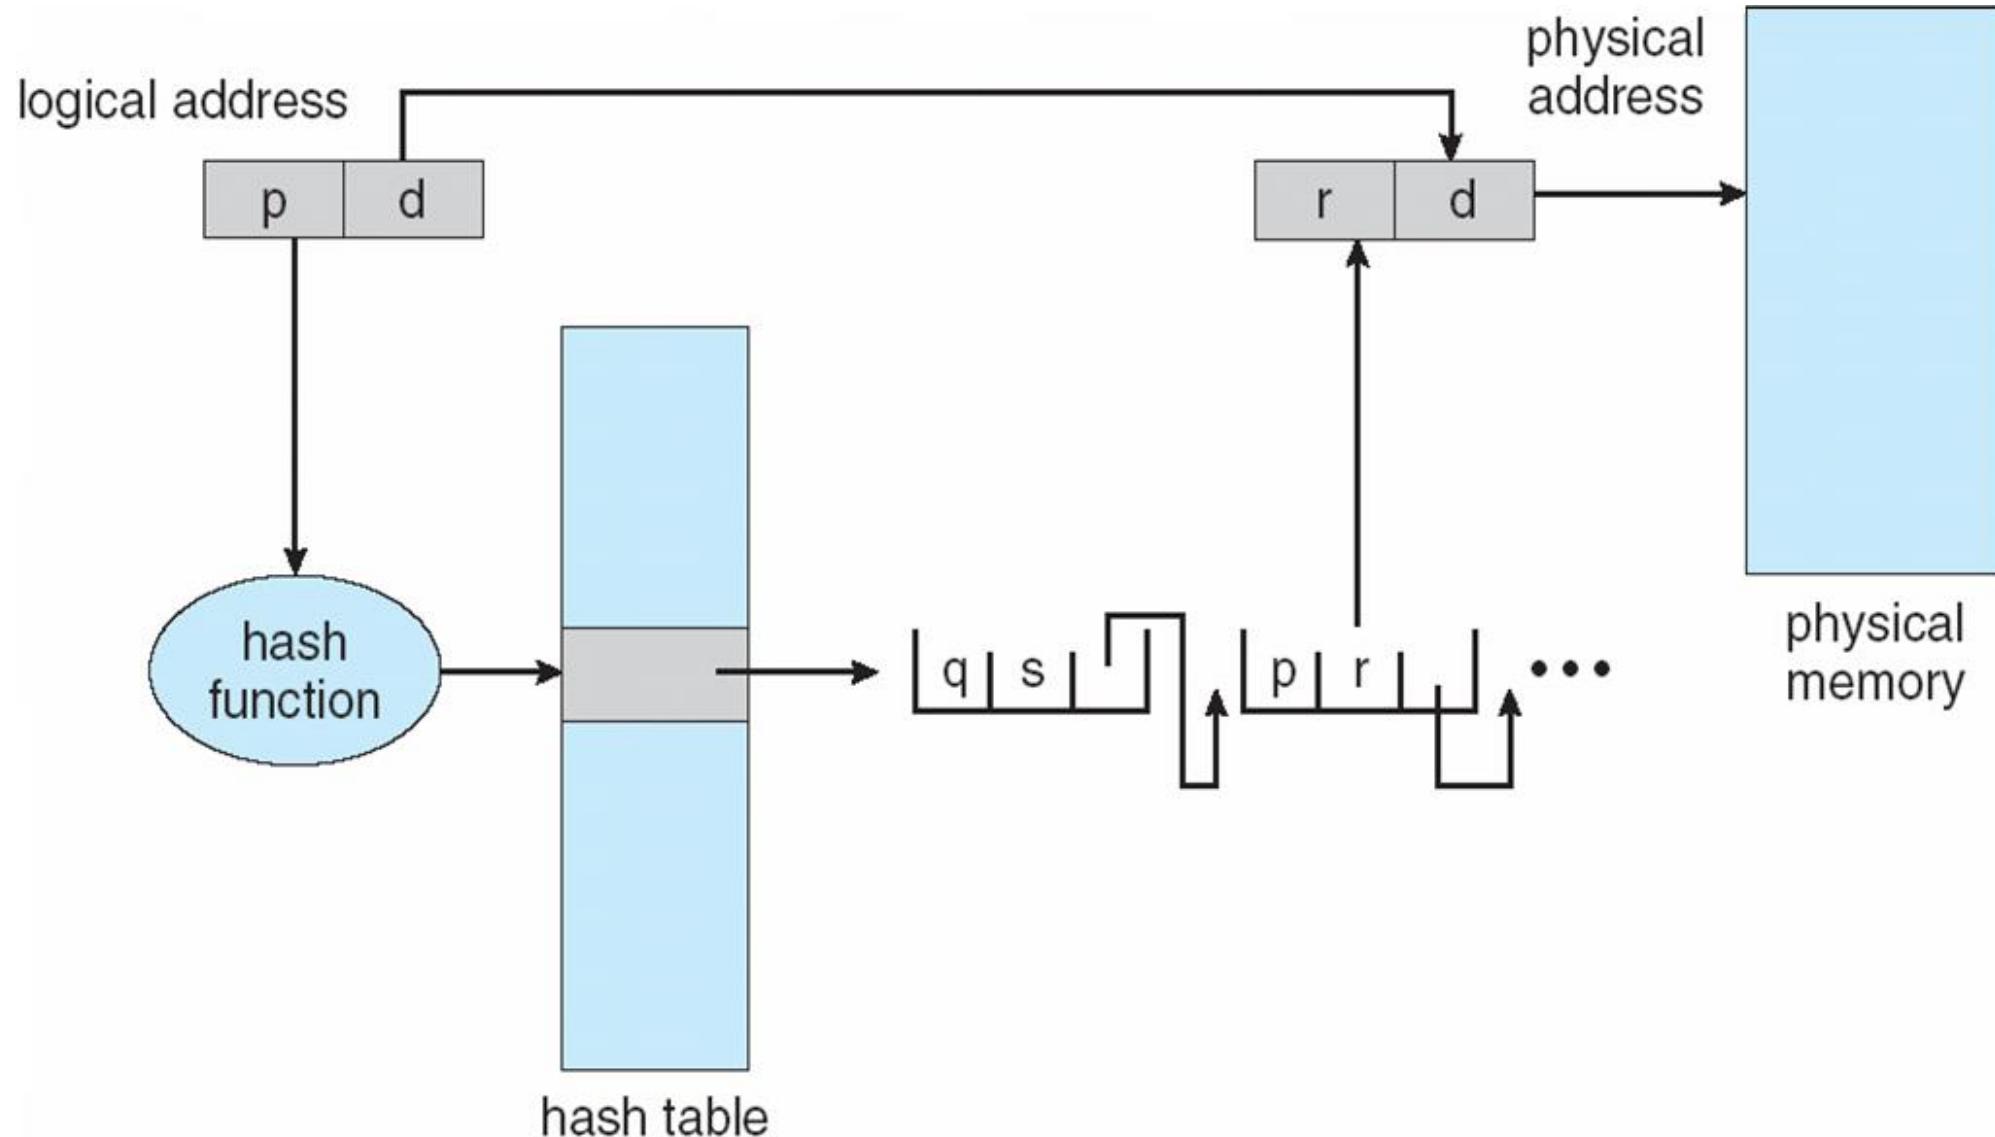
# Hashed Page Tables

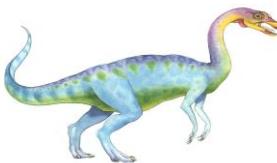
- Common approach to handle address spaces > 32 bits
- Hashed page table with **hash value (virtual page number)**
  - This page table contains a **chain of elements** hashing to the same location
- Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- Virtual page numbers are compared in this **chain searching** for a match
  - If a match is found, the corresponding physical frame is extracted
- Variation for 64-bit addresses is **clustered page tables**
  - Similar to hashed but each entry refers to several pages (such as 16) rather than 1
  - Especially useful for **sparse** address spaces (where memory references are non-contiguous and scattered)





# Hashed Page Table

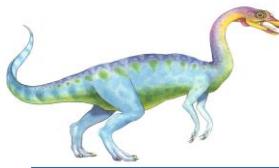




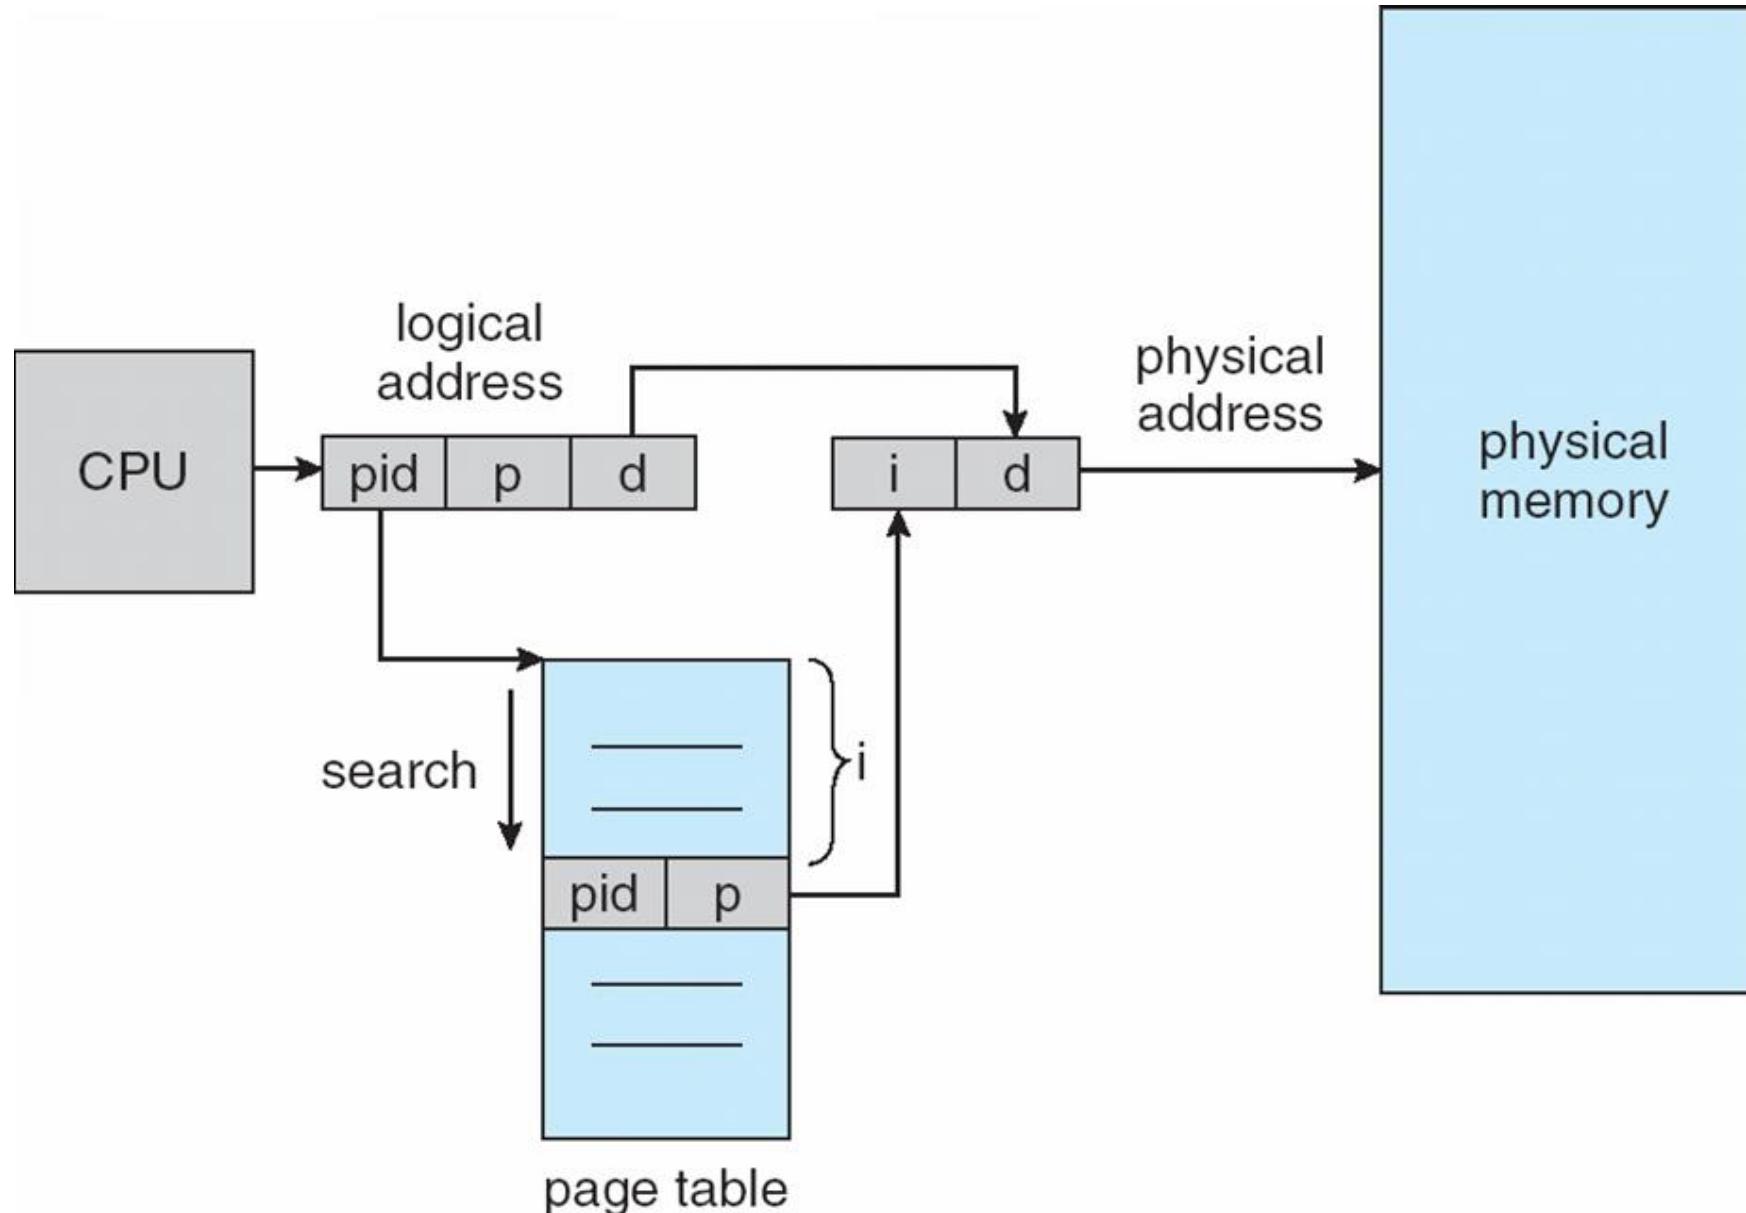
# Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- **One entry for each real page (or frame) of memory**
- Each entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- Use hash table to limit the search to one — or at most a few — page-table entries
  - TLB can accelerate access
- But how to implement shared memory?
  - One mapping of a virtual address to the shared physical address





# Inverted Page Table Architecture



# End of Chapter 8

